



```
In [79]: !pip install imblearn
# Import necessary libraries
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifie
from sklearn.svm import SVC
from sklearn.naive_bayes import GaussianNB
import requests
import io
from imblearn.over_sampling import SMOTE
```

```
Requirement already satisfied: imblearn in /usr/local/lib/python3.12/dist-pac
ges (0.0)
Requirement already satisfied: imbalanced-learn in /usr/local/lib/python3.12/di
st-packages (from imblearn) (0.14.1)
Requirement already satisfied: numpy<3,>=1.25.2 in /usr/local/lib/python3.12/di
st-packages (from imbalanced-learn->imblearn) (2.0.2)
Requirement already satisfied: scipy<2,>=1.11.4 in /usr/local/lib/python3.12/di
st-packages (from imbalanced-learn->imblearn) (1.16.3)
Requirement already satisfied: scikit-learn<2,>=1.4.2 in /usr/local/lib/python
3.12/dist-packages (from imbalanced-learn->imblearn) (1.6.1)
Requirement already satisfied: sklearn-compat<0.2,>=0.1.5 in /usr/local/lib/pyt
hon3.12/dist-packages (from imbalanced-learn->imblearn) (0.1.5)
Requirement already satisfied: joblib<2,>=1.2.0 in /usr/local/lib/python3.12/di
st-packages (from imbalanced-learn->imblearn) (1.5.3)
Requirement already satisfied: threadpoolctl<4,>=2.0.0 in /usr/local/lib/python
3.12/dist-packages (from imbalanced-learn->imblearn) (3.6.0)
```

```
In [80]: !gdown 1eiZkHD5L41lspbLXTA2BP7af0t9Z5Nw
```

```
Downloading...
From: https://drive.google.com/uc?id=1eiZkHD5L41lspbLXTA2BP7af0t9Z5Nw
To: /content/Ad_click_prediction_train (1).csv
100% 34.3M/34.3M [00:00<00:00, 76.7MB/s]
```

```
In [81]: !gdown 1eeXHpl_6WYxfTYf6Z0S7WsSxh0t3cK4W
```

```
Downloading...
From: https://drive.google.com/uc?id=1eeXHpl_6WYxfTYf6Z0S7WsSxh0t3cK4W
To: /content/Ad_Click_prediciton_test.csv
100% 9.47M/9.47M [00:00<00:00, 31.6MB/s]
```

```
In [82]: ad_click_train_data = pd.read_csv('Ad_click_prediction_train (1).csv')

ad_click_train_data
```

Out[82]:

	session_id	Datetime	user_id	product	campaign_id	webpage_id	pr
0	140690	2017-07-02 00:00	858557	C	359520	13787	
1	333291	2017-07-02 00:00	243253	C	105960	11085	
2	129781	2017-07-02 00:00	243253	C	359520	13787	
3	464848	2017-07-02 00:00	1097446	I	359520	13787	
4	90569	2017-07-02 00:01	663656	C	405490	60305	
...
463286	583588	2017-07-07 23:59	572718	H	118601	28529	
463287	198389	2017-07-07 23:59	130461	I	118601	28529	
463288	563423	2017-07-07 23:59	306241	D	118601	28529	
463289	595571	2017-07-07 23:59	306241	D	118601	28529	
463290	45023	2017-07-07 23:59	1113780	C	405490	60305	

463291 rows × 15 columns

In [83]:

```
ad_click_test_data = pd.read_csv('Ad_Click_prediciton_test.csv')  
ad_click_test_data
```

Out[83]:

	session_id	Datetime	user_id	product	campaign_id	webpage_id	pro
0	411705	2017-07-08 00:00	732573	J	404347	53587	
1	208263	2017-07-08 00:00	172910	I	118601	28529	
2	239450	2017-07-08 00:00	172910	I	118601	28529	
3	547761	2017-07-08 00:00	557318	G	118601	28529	
4	574275	2017-07-08 00:00	923896	H	118601	28529	
...
128853	215328	2017-07-09 21:29	252148	B	414149	45962	
128854	282232	2017-07-09 21:29	47955	D	98970	6970	
128855	140499	2017-07-09 21:29	314236	C	359520	13787	
128856	531038	2017-07-09 21:29	988544	E	98970	6970	
128857	349998	2017-07-09 21:29	409394	H	414149	45962	

128858 rows × 14 columns

In [84]:

```
rows, cols = ad_click_train_data.shape
print(f"The `ad_click_train_data` dataset has {rows} rows and {cols} columns.")

The `ad_click_train_data` dataset has 463291 rows and 15 columns.
```

In [85]:

```
numerical_cols = []
categorical_cols = []

for col in ad_click_train_data.columns:
    if pd.api.types.is_numeric_dtype(ad_click_train_data[col]):
        numerical_cols.append(col)
    else:
        categorical_cols.append(col)

print("Column Type Classification:")
print("-----")
for col in ad_click_train_data.columns:
    if col in numerical_cols:
        print(f"{col}: Numerical")
    else:
        print(f"{col}: Categorical")
```

```
Column Type Classification:  
-----  
session_id: Numerical  
DateTime: Categorical  
user_id: Numerical  
product: Categorical  
campaign_id: Numerical  
webpage_id: Numerical  
product_category_1: Numerical  
product_category_2: Numerical  
user_group_id: Numerical  
gender: Categorical  
age_level: Numerical  
user_depth: Numerical  
city_development_index: Numerical  
var_1: Numerical  
is_click: Numerical
```

Calculate Click-Through Rate (CTR)

```
In [86]: total_clicks = ad_click_train_data['is_click'].sum()  
total_impressions = ad_click_train_data.shape[0]  
ctr = (total_clicks / total_impressions) * 100  
  
print(f"Total Clicks: {total_clicks}")  
print(f"Total Impressions: {total_impressions}")  
print(f"Click-Through Rate (CTR): {ctr:.2f}%")
```

```
Total Clicks: 31331  
Total Impressions: 463291  
Click-Through Rate (CTR): 6.76%
```

```
In [87]: missing_values = ad_click_train_data.isnull().sum()  
missing_values = missing_values[missing_values > 0]  
  
print("Columns with Missing Values and their Null Counts:")  
print("-----")  
if not missing_values.empty:  
    print(missing_values)  
else:  
    print("No missing values found in the dataset.")
```

```
Columns with Missing Values and their Null Counts:  
-----  
product_category_2      365854  
user_group_id           18243  
gender                  18243  
age_level                18243  
user_depth                18243  
city_development_index     125129  
dtype: int64
```

Summary of `ad_click_train_data` Analysis

1. Dataset Dimensions:

- **Rows:** 463291
- **Columns:** 15

2. Column Type Classification:

- **Numerical Columns:**

- `session_id`
- `user_id`
- `campaign_id`
- `webpage_id`
- `product_category_1`
- `product_category_2`
- `user_group_id`
- `age_level`
- `user_depth`
- `city_development_index`
- `var_1`
- `is_click`

- **Categorical Columns:**

- `DateTime`
- `product`
- `gender`

3. Click-Through Rate (CTR):

- **Total Clicks:** 31331
- **Total Impressions:** 463291
- **Click-Through Rate (CTR):** 6.76%

4. Columns with Missing Values:

- `product_category_2` : 365854 null values
- `user_group_id` : 18243 null values
- `gender` : 18243 null values
- `age_level` : 18243 null values
- `user_depth` : 18243 null values
- `city_development_index` : 125129 null values

Temporal Patterns

Extract Hour from DateTime

```
In [88]: ad_click_train_data['DateTime'] = pd.to_datetime(ad_click_train_data['DateTime'])
ad_click_train_data['hour_of_day'] = ad_click_train_data['DateTime'].dt.hour
print(ad_click_train_data[['DateTime', 'hour_of_day']].head())
```

	DateTime	hour_of_day
0	2017-07-02 00:00:00	0
1	2017-07-02 00:00:00	0
2	2017-07-02 00:00:00	0
3	2017-07-02 00:00:00	0
4	2017-07-02 00:01:00	0

Calculate Click-Through Rate per Hour

```
In [89]: hourly_ctr = ad_click_train_data.groupby('hour_of_day').agg(
    total_clicks=('is_click', 'sum'),
    total_impressions=('session_id', 'count')
)
hourly_ctr['CTR'] = (hourly_ctr['total_clicks'] / hourly_ctr['total_impressions'])

print("Click-Through Rate (CTR) per Hour:")
print("-----")
print(hourly_ctr[['CTR']].sort_index())
```

Click-Through Rate (CTR) per Hour:

hour_of_day	CTR
0	6.849315
1	7.460815
2	5.273189
3	6.223734
4	5.700510
5	6.702681
6	7.282210
7	7.397823
8	7.027111
9	7.010128
10	6.777607
11	6.954296
12	6.858513
13	6.641993
14	6.530361
15	6.392927
16	6.722104
17	6.634720
18	6.622956
19	6.669113
20	6.635217
21	6.952918
22	6.651026
23	6.373626

Identify Hour with Highest CTR

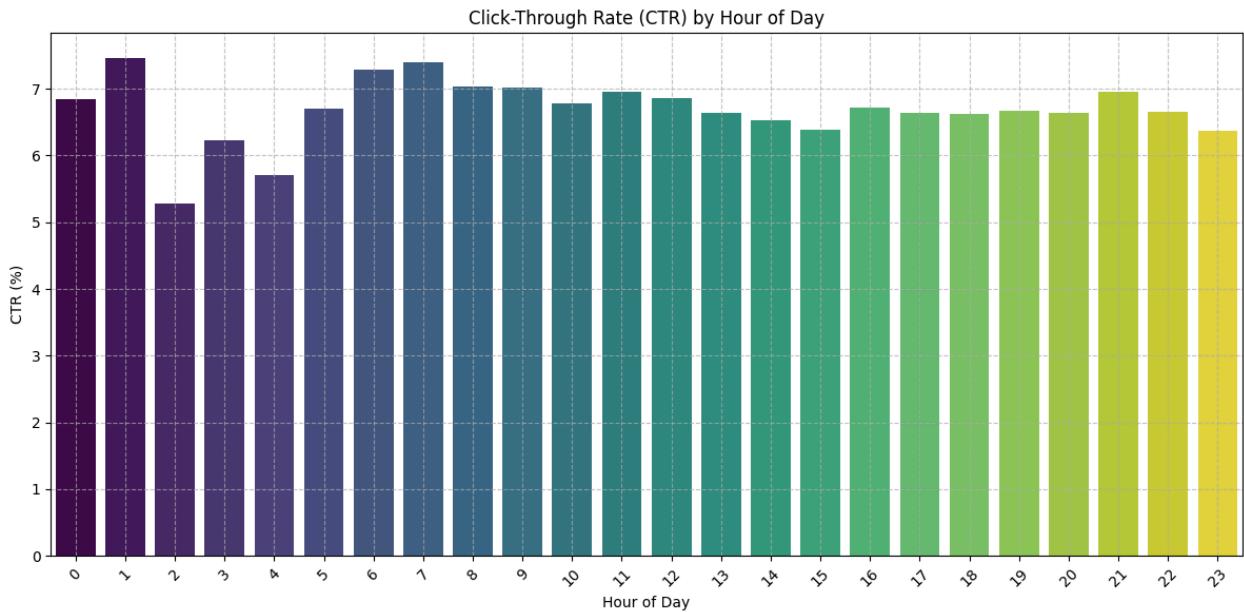
```
In [90]: max_ctr_hour = hourly_ctr['CTR'].idxmax()
max_ctr_value = hourly_ctr['CTR'].max()

print(f"The hour with the highest CTR is {max_ctr_hour} with a CTR of {max_ctr}
```

The hour with the highest CTR is 1 with a CTR of 7.46%

Visualize Hourly CTR

```
In [91]: plt.figure(figsize=(12, 6))
sns.barplot(x=hourly_ctr.index, y=hourly_ctr['CTR'], hue=hourly_ctr.index, palette='viridis')
plt.title('Click-Through Rate (CTR) by Hour of Day')
plt.xlabel('Hour of Day')
plt.ylabel('CTR (%)')
plt.xticks(rotation=45)
plt.grid(True, linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



Extract Day of Week

```
In [92]: ad_click_train_data['day_of_week'] = ad_click_train_data['DateTime'].dt.dayofweek
ad_click_train_data['day_type'] = ad_click_train_data['day_of_week'].apply(lambda x: 'Weekend' if x == 6 else 'Weekday')

print(ad_click_train_data[['DateTime', 'day_of_week', 'day_type']].head())
```

	DateTime	day_of_week	day_type
0	2017-07-02 00:00:00	6	Weekend
1	2017-07-02 00:00:00	6	Weekend
2	2017-07-02 00:00:00	6	Weekend
3	2017-07-02 00:00:00	6	Weekend
4	2017-07-02 00:01:00	6	Weekend

```
In [93]: day_type_ctr = ad_click_train_data.groupby('day_type').agg(
    total_clicks=('is_click', 'sum'),
    total_impressions=('session_id', 'count')
)
day_type_ctr['CTR'] = (day_type_ctr['total_clicks'] / day_type_ctr['total_impressions'])

print("Click-Through Rate (CTR) by Day Type:")
print("-----")
print(day_type_ctr[['CTR']])
```

Click-Through Rate (CTR) by Day Type:

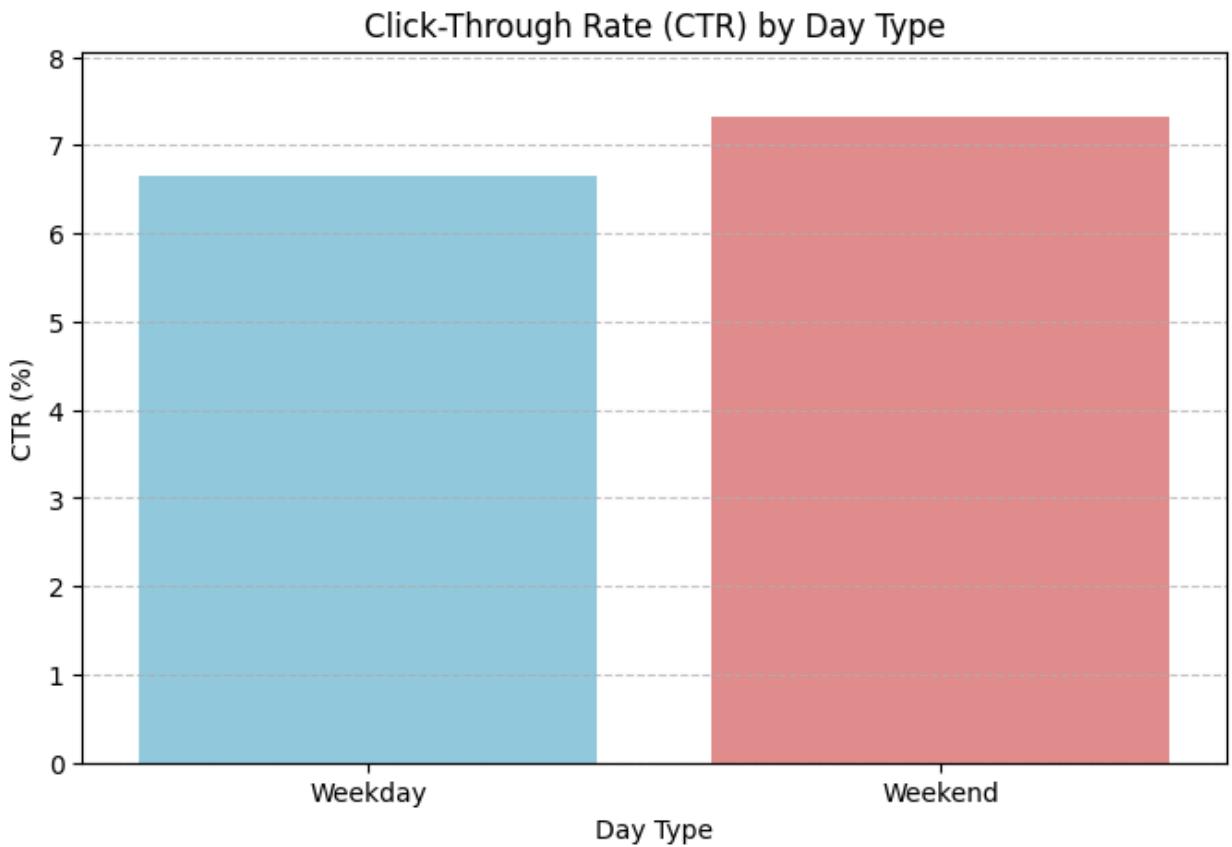
day_type	CTR
Weekday	6.646784
Weekend	7.326207

```
In [94]: plt.figure(figsize=(8, 5))
sns.barplot(x=day_type_ctr.index, y=day_type_ctr['CTR'], hue=day_type_ctr.index)
plt.title('Click-Through Rate (CTR) by Day Type')
```

```

plt.xlabel('Day Type')
plt.ylabel('CTR (%)')
plt.ylim(0, day_type_ctr['CTR'].max() * 1.1) # Set y-axis limit to be slightly
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()

```



Summary:

- **How does Click-Through Rate (CTR) vary between weekdays and weekends?** The Click-Through Rate (CTR) is higher on weekends compared to weekdays. Specifically, the CTR for weekdays is approximately 6.65%, while for weekends, it is approximately 7.33%.

Insights or Next Steps

- The higher CTR on weekends suggests that users might be more engaged or have more time to interact with ads during these periods, which could inform advertising campaign scheduling.
- Further analysis could explore if specific ad categories or products perform significantly better on weekends versus weekdays, leading to more targeted ad strategies.

User Behavior

Do certain age groups click more?

```
In [95]: print("Unique values in 'age_level' before handling missing values:")
print(ad_click_train_data['age_level'].value_counts().sort_index())

missing_age_level = ad_click_train_data['age_level'].isnull().sum()
print(f"\nNumber of missing values in 'age_level': {missing_age_level}")

if missing_age_level > 0:
    mode_age_level = ad_click_train_data['age_level'].mode()[0]
    ad_click_train_data['age_level'] = ad_click_train_data['age_level'].fillna(mode_age_level)
    print(f"Missing values in 'age_level' filled with mode: {mode_age_level}")

print("\nValue counts for 'age_level' after handling missing values:")
print(ad_click_train_data['age_level'].value_counts().sort_index())
```

Unique values in 'age_level' before handling missing values:
age_level

```
0.0      153
1.0     43367
2.0    143501
3.0   160581
4.0    63859
5.0    30828
6.0    2759
Name: count, dtype: int64
```

Number of missing values in 'age_level': 18243
Missing values in 'age_level' filled with mode: 3.0

Value counts for 'age_level' after handling missing values:

```
age_level
0.0      153
1.0     43367
2.0    143501
3.0   178824
4.0    63859
5.0    30828
6.0    2759
Name: count, dtype: int64
```

```
In [96]: age_level_ctr = ad_click_train_data.groupby('age_level').agg(
    total_clicks=('is_click', 'sum'),
    total_impressions=('session_id', 'count')
)
age_level_ctr['CTR'] = (age_level_ctr['total_clicks'] / age_level_ctr['total_impressions'])

print("Click-Through Rate (CTR) by Age Level:")
print("-----")
```

```
print(age_level_ctr[['CTR']].sort_index())
```

Click-Through Rate (CTR) by Age Level:

```
-----  
          CTR  
age_level  
0.0      8.496732  
1.0      7.480342  
2.0      7.091937  
3.0      6.505838  
4.0      5.872312  
5.0      7.415337  
6.0      8.227619
```

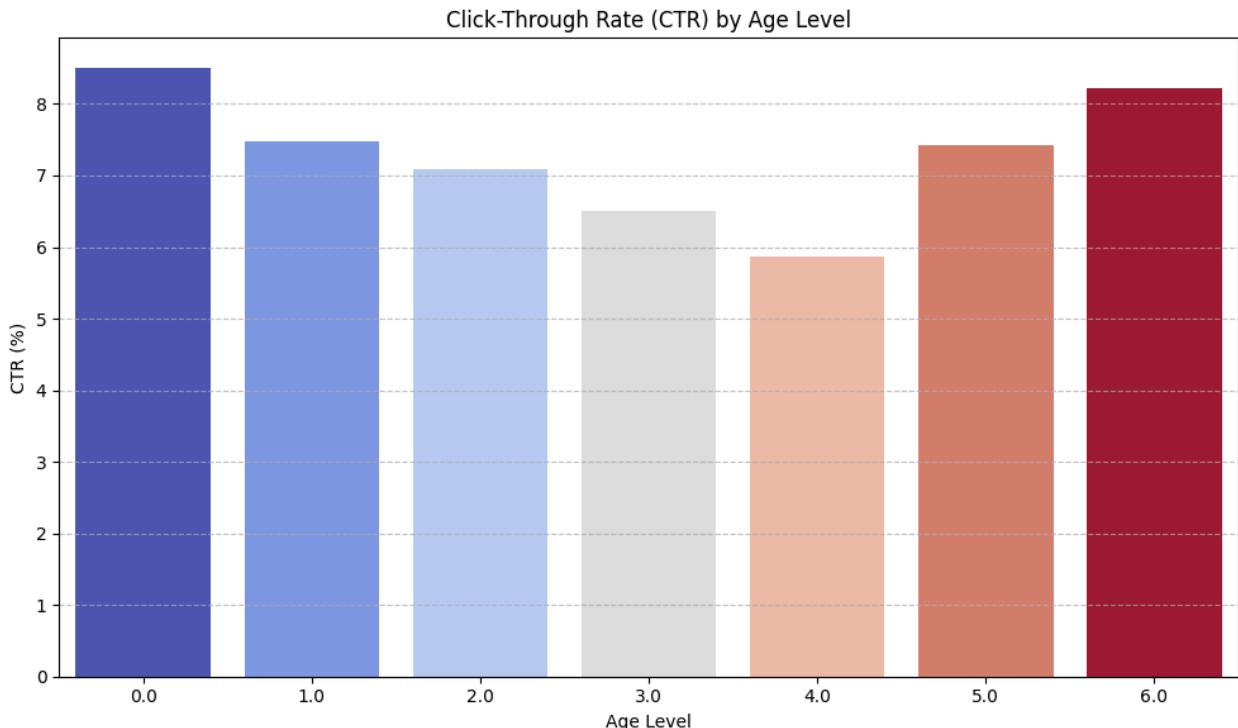
```
In [97]: max_ctr_age_level = age_level_ctr['CTR'].idxmax()
```

```
max_ctr_value = age_level_ctr['CTR'].max()
```

```
print(f"The age level with the highest CTR is {max_ctr_age_level} with a CTR c
```

The age level with the highest CTR is 0.0 with a CTR of 8.50%

```
In [98]: plt.figure(figsize=(10, 6))  
sns.barplot(x=age_level_ctr.index, y=age_level_ctr['CTR'], hue=age_level_ctr.i  
plt.title('Click-Through Rate (CTR) by Age Level')  
plt.xlabel('Age Level')  
plt.ylabel('CTR (%)')  
plt.xticks(rotation=0)  
plt.grid(axis='y', linestyle='--', alpha=0.7)  
plt.tight_layout()  
plt.show()
```



Is there a gender difference in click rates?

```
In [99]: print("Unique values in 'gender' before handling missing values:")
print(ad_click_train_data['gender'].value_counts(dropna=False))
```

```
Unique values in 'gender' before handling missing values:
gender
Male      393454
Female    51594
NaN       18243
Name: count, dtype: int64
```

```
In [100... missing_gender = ad_click_train_data['gender'].isnull().sum()
print(f"\nNumber of missing values in 'gender': {missing_gender}")

if missing_gender > 0:
    mode_gender = ad_click_train_data['gender'].mode()[0]
    ad_click_train_data['gender'] = ad_click_train_data['gender'].fillna(mode_
    print(f"Missing values in 'gender' filled with mode: {mode_gender}")

print("\nValue counts for 'gender' after handling missing values:")
print(ad_click_train_data['gender'].value_counts(dropna=False))
```

```
Number of missing values in 'gender': 18243
Missing values in 'gender' filled with mode: Male
```

```
Value counts for 'gender' after handling missing values:
gender
Male      411697
Female    51594
Name: count, dtype: int64
```

```
In [101... gender_ctr = ad_click_train_data.groupby('gender').agg(
    total_clicks=('is_click', 'sum'),
    total_impressions=('session_id', 'count')
)
gender_ctr['CTR'] = (gender_ctr['total_clicks'] / gender_ctr['total_impressions'])

print("Click-Through Rate (CTR) by Gender:")
print("-----")
print(gender_ctr[['CTR']].sort_index())
```

```
Click-Through Rate (CTR) by Gender:
```

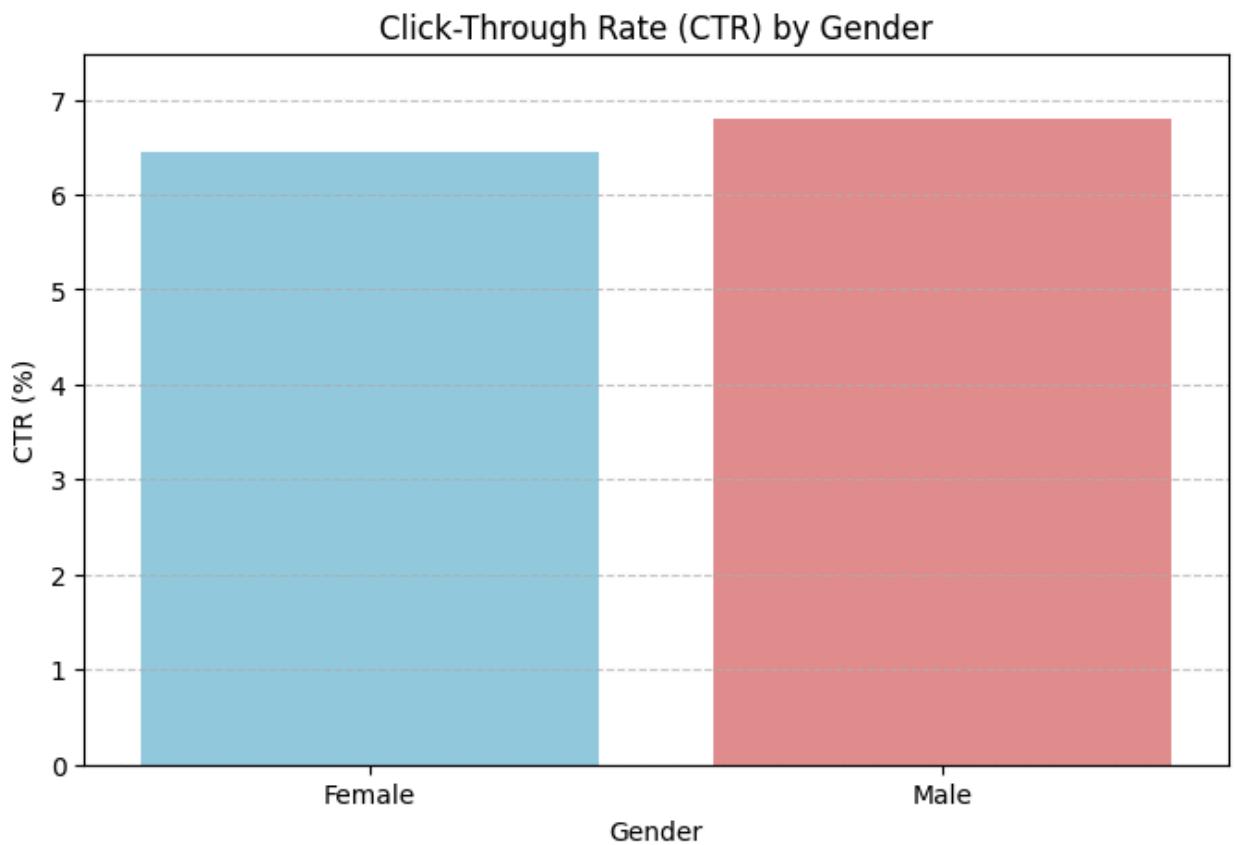
```
-----
CTR
gender
Female  6.444548
Male    6.802576
```

```
In [102... max_ctr_gender = gender_ctr['CTR'].idxmax()
max_ctr_value = gender_ctr['CTR'].max()
```

```
print(f"The gender with the highest CTR is {max_ctr_gender} with a CTR of {max
```

The gender with the highest CTR is Male with a CTR of 6.80%

```
In [103]: plt.figure(figsize=(8, 5))
sns.barplot(x=gender_ctr.index, y=gender_ctr['CTR'], hue=gender_ctr.index, palette='Set1')
plt.title('Click-Through Rate (CTR) by Gender')
plt.xlabel('Gender')
plt.ylabel('CTR (%)')
plt.ylim(0, gender_ctr['CTR'].max() * 1.1) # Set y-axis limit to be slightly above the max
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.show()
```



Insights or Next Steps

- The marginal difference in CTR between genders suggests that while there is a slight preference for males, the advertising strategy might not need significant gender-specific adjustments unless specific products target one gender predominantly.
- Further analysis could explore whether certain product categories or campaign types appeal more to one gender, which could refine targeting efforts.
- Investigate the characteristics of the imputed 'Male' gender entries to

ensure that the imputation did not skew the results for the male demographic.

How does user group affect clicking?

```
In [104...]: print("Unique values in 'user_group_id' before handling missing values:")
print(ad_click_train_data['user_group_id'].value_counts(dropna=False))

missing_user_group_id = ad_click_train_data['user_group_id'].isnull().sum()
print(f"\nNumber of missing values in 'user_group_id': {missing_user_group_id}")
```

Unique values in 'user_group_id' before handling missing values:

```
user_group_id
3.0      140317
2.0      137278
4.0      50080
1.0      41946
5.0      21989
9.0      20264
NaN      18243
10.0     13779
11.0     8839
8.0      6223
6.0      1724
7.0      1421
12.0     1035
0.0      153
Name: count, dtype: int64
```

Number of missing values in 'user_group_id': 18243

```
In [105...]: if missing_user_group_id > 0:
    mode_user_group_id = ad_click_train_data['user_group_id'].mode()[0]
    ad_click_train_data['user_group_id'] = ad_click_train_data['user_group_id'].fillna(mode_user_group_id)
    print(f"Missing values in 'user_group_id' filled with mode: {mode_user_group_id}")

print("\nValue counts for 'user_group_id' after handling missing values:")
print(ad_click_train_data['user_group_id'].value_counts(dropna=False))
```

```
Missing values in 'user_group_id' filled with mode: 3.0

Value counts for 'user_group_id' after handling missing values:
user_group_id
3.0      158560
2.0      137278
4.0      50080
1.0      41946
5.0      21989
9.0      20264
10.0     13779
11.0     8839
8.0      6223
6.0      1724
7.0      1421
12.0     1035
0.0       153
Name: count, dtype: int64
```

```
In [106...]: user_group_ctr = ad_click_train_data.groupby('user_group_id').agg(
    total_clicks=('is_click', 'sum'),
    total_impressions=('session_id', 'count')
)
user_group_ctr['CTR'] = (user_group_ctr['total_clicks'] / user_group_ctr['total_impressions'])

print("Click-Through Rate (CTR) by User Group:")
print("-----")
print(user_group_ctr[['CTR']].sort_index())
```

Click-Through Rate (CTR) by User Group:

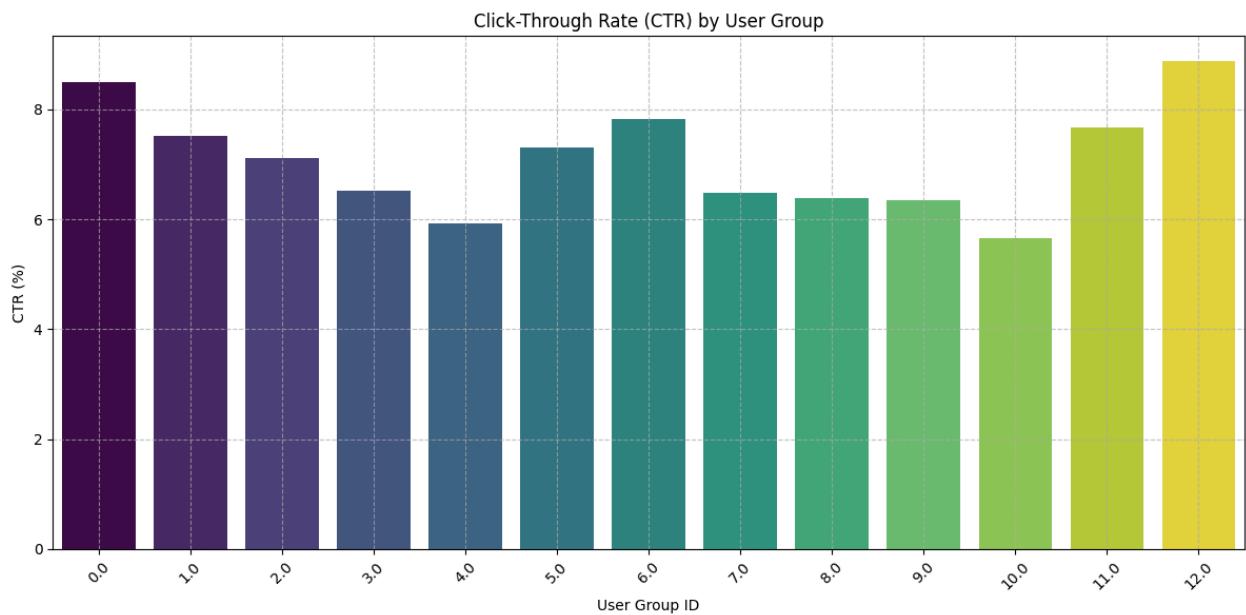
```
-----  
          CTR  
user_group_id  
0.0      8.496732  
1.0      7.514423  
2.0      7.124230  
3.0      6.526236  
4.0      5.932508  
5.0      7.312747  
6.0      7.830626  
7.0      6.474314  
8.0      6.379560  
9.0      6.346230  
10.0     5.653531  
11.0     7.670551  
12.0     8.888889
```

```
In [107...]: max_ctr_user_group = user_group_ctr['CTR'].idxmax()
max_ctr_value = user_group_ctr['CTR'].max()

print(f"The user group with the highest CTR is {max_ctr_user_group} with a CTR of {max_ctr_value:.2%}")
```

The user group with the highest CTR is 12.0 with a CTR of 8.89%

```
In [108]: plt.figure(figsize=(12, 6))
sns.barplot(x=user_group_ctr.index, y=user_group_ctr['CTR'], hue=user_group_ct
plt.title('Click-Through Rate (CTR) by User Group')
plt.xlabel('User Group ID')
plt.ylabel('CTR (%)')
plt.xticks(rotation=45)
plt.grid(True, linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



Summary:

- **How does user group affect clicking?** User group significantly affects clicking, with CTRs varying across different groups. User group **12.0** exhibits the highest Click-Through Rate (CTR) at **8.89%**, while user group **10.0** shows the lowest CTR at **5.65%**.

Insights or Next Steps

- **Targeted Campaigns:** User groups 12.0, 0.0, 6.0, and 11.0 show significantly higher CTRs. Marketing efforts should prioritize these groups with tailored campaigns, content, and product offerings that align with their interests.
- **Improve Engagement:** For user groups with lower CTRs (e.g., 10.0, 4.0, 7.0, 8.0, 9.0), it's crucial to investigate the underlying reasons for their lower engagement. This could involve A/B testing different ad creatives, targeting new demographics within these groups, or re-evaluating the value proposition presented to them.

- **Profile High-Performing Groups:** Further analysis of user groups with high CTRs (12.0, 0.0, 6.0) could involve examining their demographic information, product preferences, and typical browsing behavior to create more granular personas. This detailed understanding can then be used to inform strategies for other groups.
-
-

Campaign Performance

Which campaigns have highest CTR?

```
In [109...]: print("Unique values in 'campaign_id':")
print(ad_click_train_data['campaign_id'].unique())

print("\nValue counts for 'campaign_id':")
print(ad_click_train_data['campaign_id'].value_counts())
```

Unique values in 'campaign_id':
[359520 105960 405490 360936 404347 98970 414149 82320 396664 118601]

Value counts for 'campaign_id':
campaign_id
359520 108155
405490 95973
360936 51888
118601 35531
98970 35065
414149 29314
404347 28826
82320 27849
105960 25781
396664 24909
Name: count, dtype: int64

```
In [110...]: campaign_ctr = ad_click_train_data.groupby('campaign_id').agg(
    total_clicks=('is_click', 'sum'),
    total_impressions=('session_id', 'count')
)
campaign_ctr['CTR'] = (campaign_ctr['total_clicks'] / campaign_ctr['total_impressions'])

print("Click-Through Rate (CTR) by Campaign ID:")
print("-----")
print(campaign_ctr[['CTR']].sort_values(by='CTR', ascending=False))
```

Click-Through Rate (CTR) by Campaign ID:

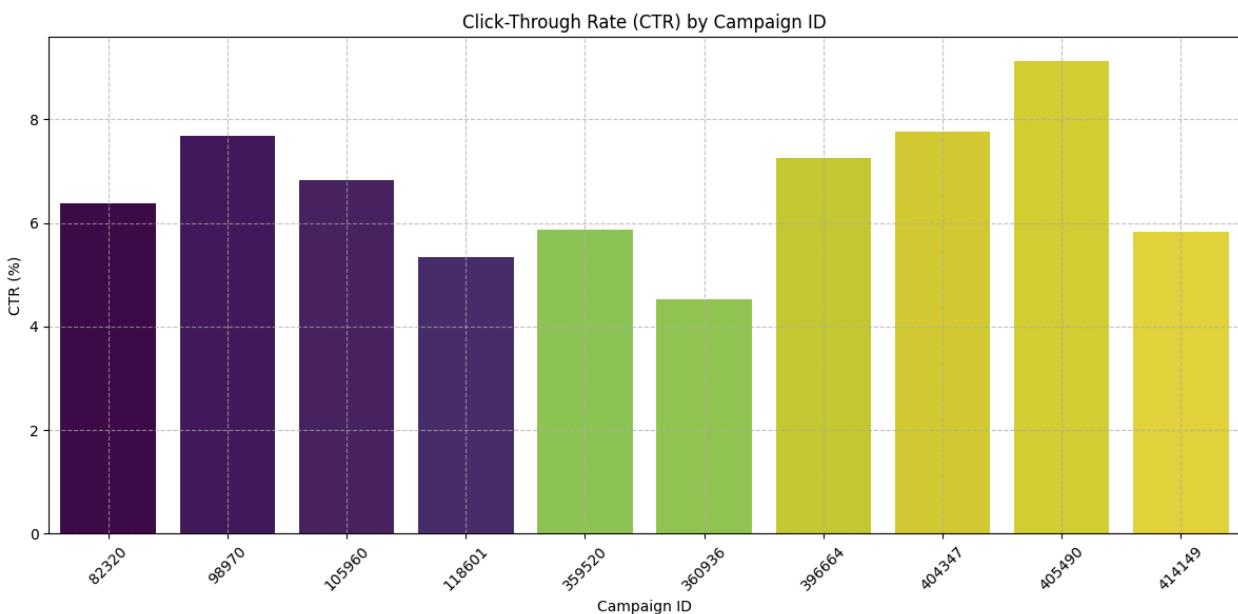
	CTR
campaign_id	
405490	9.130693
404347	7.753417
98970	7.682875
396664	7.262435
105960	6.834491
82320	6.377249
359520	5.861957
414149	5.833390
118601	5.336185
360936	4.521277

```
In [111]: max_ctr_campaign = campaign_ctr['CTR'].idxmax()
max_ctr_value = campaign_ctr['CTR'].max()

print(f"The campaign with the highest CTR is {max_ctr_campaign} with a CTR of
```

The campaign with the highest CTR is 405490 with a CTR of 9.13%

```
In [112]: plt.figure(figsize=(12, 6))
sns.barplot(x=campaign_ctr.index, y=campaign_ctr['CTR'], hue=campaign_ctr.index)
plt.title('Click-Through Rate (CTR) by Campaign ID')
plt.xlabel('Campaign ID')
plt.ylabel('CTR (%)')
plt.xticks(rotation=45)
plt.grid(True, linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



Insights or Next Steps

- **Optimize Campaign Strategy:** Campaigns with high CTRs, such as **405490**, **404347**, and **98970**, should be further analyzed to understand their success factors (e.g., ad creative, target audience, placement, product type). These insights can then be applied to improve underperforming campaigns.
- **Resource Allocation:** Consider allocating more resources to the top-performing campaigns or developing new campaigns that mirror their successful attributes.
- **Improve Low-Performing Campaigns:** Investigate campaigns with low CTRs (e.g., **360936**, **118601**) to identify reasons for poor engagement. This could involve A/B testing different ad creatives, adjusting targeting, or reviewing the messaging.
- **Correlation with Volume:** While `campaign_id 359520` had the highest number of impressions, its CTR (5.86%) was not among the highest, suggesting that high volume does not always equate to high engagement. This highlights the importance of CTR as a performance metric.

Which products get more clicks?

```
In [113...]: print("Unique values in 'product':")
print(ad_click_train_data['product'].unique())

print("\nValue counts for 'product':")
print(ad_click_train_data['product'].value_counts())
```

```
Unique values in 'product':
['C' 'I' 'F' 'H' 'B' 'D' 'G' 'E' 'J' 'A']
```

```
Value counts for 'product':
```

```
product
C    163501
H    109574
I     63711
D     41064
B     22479
E     21452
A     15391
J      9698
G      9414
F      7007
Name: count, dtype: int64
```

```
In [114...]: product_ctr = ad_click_train_data.groupby('product').agg(
    total_clicks=('is_click', 'sum'),
```

```

        total_impressions=('session_id', 'count')
    )
product_ctr['CTR'] = (product_ctr['total_clicks'] / product_ctr['total_impressions'])

print("Click-Through Rate (CTR) by Product:")
print("-----")
print(product_ctr[['CTR']].sort_values(by='CTR', ascending=False))

```

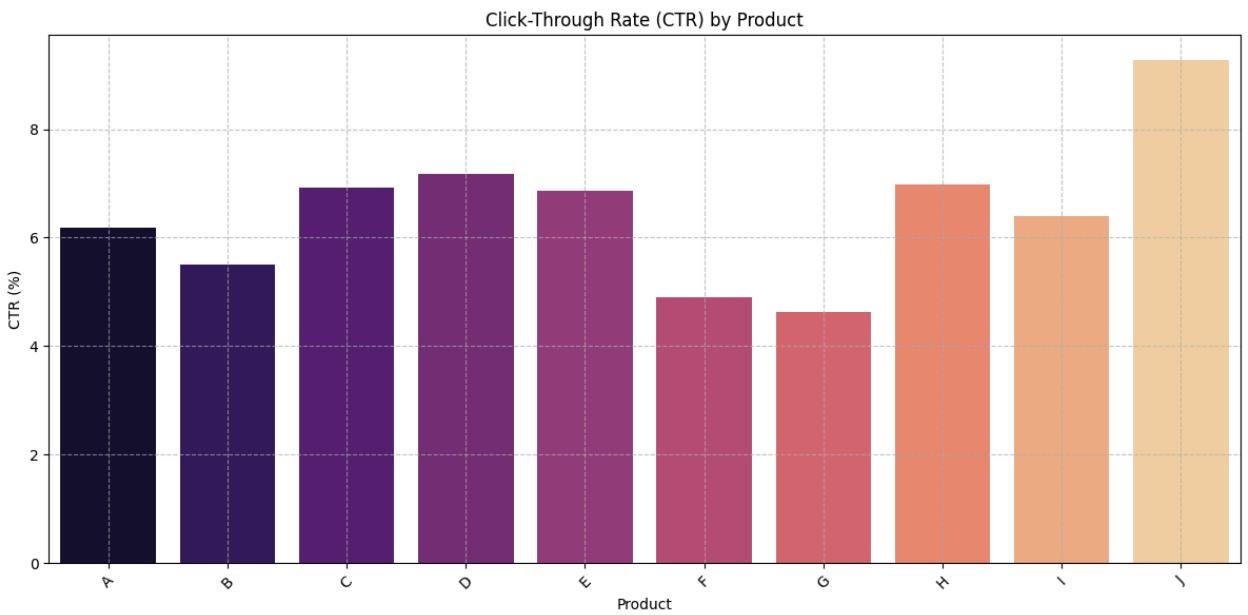
Click-Through Rate (CTR) by Product:

product	CTR
J	9.269953
D	7.181473
H	6.985234
C	6.914942
E	6.871154
I	6.402348
A	6.191930
B	5.507362
F	4.909376
G	4.620778

```

In [115]: plt.figure(figsize=(12, 6))
sns.barplot(x=product_ctr.index, y=product_ctr['CTR'], hue=product_ctr.index,
            plt.title('Click-Through Rate (CTR) by Product')
            plt.xlabel('Product')
            plt.ylabel('CTR (%)')
            plt.xticks(rotation=45)
            plt.grid(True, linestyle='--', alpha=0.7)
            plt.tight_layout()
            plt.show()

```



Summary:

- **Which products get more clicks?** Product **J** has the highest Click-Through Rate (CTR) at **9.27%**, followed by Product **D** (7.18%) and Product **H** (6.99%).

Data Analysis Key Findings

- The `product` column contains 10 unique product categories, with 'C' being the most frequent (163,501 entries) and 'F' being the least frequent (7,007 entries).
- Click-Through Rates (CTRs) were calculated for each product, showing a range of performance from **4.62%** (Product G) to **9.27%** (Product J).
- Product **J** demonstrated the highest CTR at **9.27%**, indicating strong user engagement.
- Products **D** (7.18%) and **H** (6.99%) also showed comparatively high CTRs.
- Products **G** (4.62%) and **F** (4.91%) had the lowest CTRs, suggesting lower user interest or less effective advertising for these categories.
- The bar plot effectively visualized these differences, clearly highlighting the top-performing products.

Insights or Next Steps

- **Product Focus:** Prioritize marketing and advertising efforts on Product **J**, as it exhibits the highest engagement. Investigate the characteristics (e.g., ad creative, target audience, pricing, seasonal relevance) that make Product J so successful and attempt to replicate these factors for other products.
- **Improve Underperforming Products:** Analyze Products **G** and **F** to understand the reasons behind their low CTRs. This could involve reassessing their market fit, improving product descriptions, optimizing ad campaigns, or exploring different sales channels.
- **Content and Targeting:** Further examine the demographic and behavioral data of users who click on high-CTR products like J, D, and H to refine targeting strategies for future campaigns across all product categories.
- **Inventory Management:** High CTR products might indicate high demand. This insight can be used to inform inventory management and supply chain decisions.

Do certain webpages convert better?

```
In [116]: print("Unique values in 'webpage_id':")
print(ad_click_train_data['webpage_id'].unique())

print("\nValue counts for 'webpage_id' before handling missing values:")
print(ad_click_train_data['webpage_id'].value_counts(dropna=False))

missing_webpage_id = ad_click_train_data['webpage_id'].isnull().sum()
print(f"\nNumber of missing values in 'webpage_id': {missing_webpage_id}")
```

Unique values in 'webpage_id':
[13787 11085 60305 53587 6970 45962 1734 51181 28529]

Value counts for 'webpage_id' before handling missing values:
 webpage_id
13787 160043
60305 95973
28529 35531
6970 35065
45962 29314
53587 28826
1734 27849
11085 25781
51181 24909
Name: count, dtype: int64

Number of missing values in 'webpage_id': 0

```
In [117]: webpage_ctr = ad_click_train_data.groupby('webpage_id').agg(
    total_clicks=('is_click', 'sum'),
    total_impressions=('session_id', 'count')
)
webpage_ctr['CTR'] = (webpage_ctr['total_clicks'] / webpage_ctr['total_impressions'])

print("Click-Through Rate (CTR) by Webpage ID:")
print("-----")
print(webpage_ctr[['CTR']].sort_values(by='CTR', ascending=False))
```

Click-Through Rate (CTR) by Webpage ID:

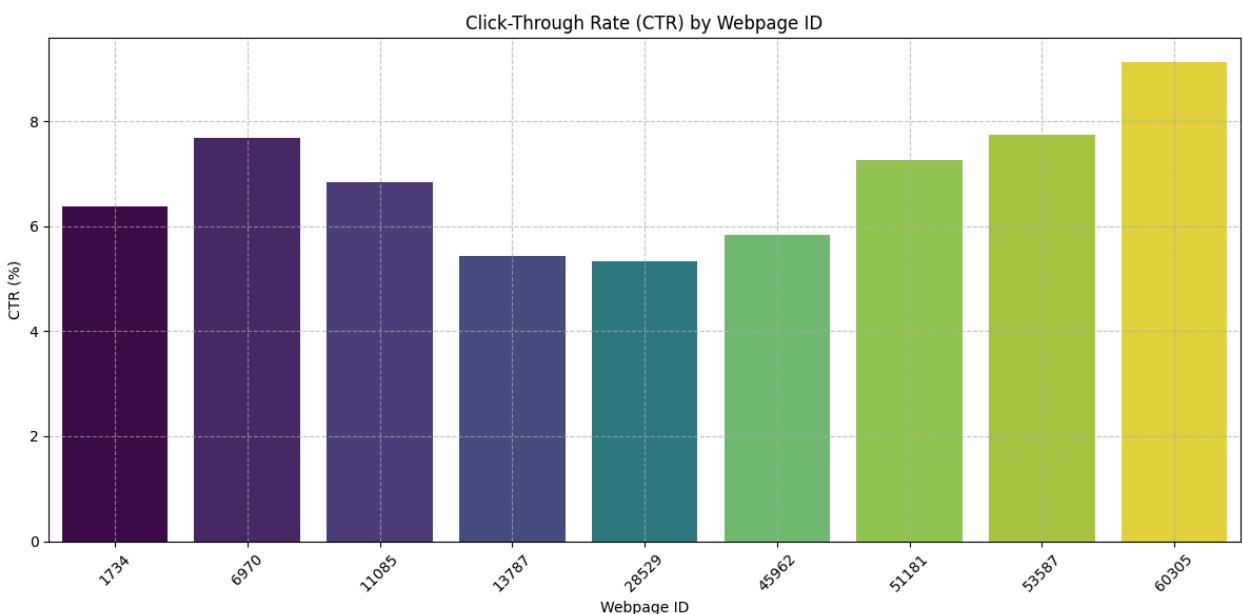
 CTR
 webpage_id
60305 9.130693
53587 7.753417
6970 7.682875
51181 7.262435
11085 6.834491
1734 6.377249
45962 5.833390
13787 5.427291
28529 5.336185

```
In [118]: max_ctr_webpage = webpage_ctr['CTR'].idxmax()
max_ctr_value = webpage_ctr['CTR'].max()

print(f"The webpage with the highest CTR is {max_ctr_webpage} with a CTR of {max_ctr_value:.2%}")
```

The webpage with the highest CTR is 60305 with a CTR of 9.13%

```
In [119]: plt.figure(figsize=(12, 6))
sns.barplot(x=webpage_ctr.index, y=webpage_ctr['CTR'], hue=webpage_ctr.index,
plt.title('Click-Through Rate (CTR) by Webpage ID')
plt.xlabel('Webpage ID')
plt.ylabel('CTR (%)')
plt.xticks(rotation=45)
plt.grid(True, linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



Summary:

- **Do certain webpages convert better?** Yes, there is a significant variation in Click-Through Rates (CTRs) across different webpages. Webpage **60305** has the highest CTR at **9.13%**, while webpage **28529** has the lowest at **5.34%**.

Insights or Next Steps

- **Content Optimization:** Analyze the design, content, and ad placements on high-performing webpages like **60305** to identify best practices. These successful elements should be replicated on underperforming pages to improve their CTRs.

- **Targeted Improvement:** For webpages with low CTRs (e.g., **28529**, **13787**), conduct a thorough review of their content, user experience, and the relevance of the ads displayed. A/B testing different layouts, call-to-actions, or ad creatives could help boost engagement.
- **User Flow Analysis:** Investigate the user journey leading to high-CTR webpages. Understanding how users arrive at and interact with these pages can provide valuable insights for optimizing navigation and overall site structure.
- **Correlation with Product/Campaigns:** Further analysis could explore which products or campaigns are most frequently associated with high-CTR webpages. This can help in aligning specific ad campaigns with the most effective webpage destinations.

Feature Engineering

Extract from DateTime column

- hour: Hour of day (0-23)
- day_of_week: Day (0=Monday, 6=Sunday)
- day_of_month: Date of month (1-31)
- month: Month of year (1-12)
- is_weekend: Binary flag (Saturday/Sunday = 1)
- time_of_day: Categorical (night/morning/afternoon/evening)

```
In [120]: ad_click_train_data['day_of_month'] = ad_click_train_data['DateTime'].dt.day
print(ad_click_train_data[['DateTime', 'day_of_month']].head())
```

	DateTime	day_of_month
0	2017-07-02 00:00:00	2
1	2017-07-02 00:00:00	2
2	2017-07-02 00:00:00	2
3	2017-07-02 00:00:00	2
4	2017-07-02 00:01:00	2

```
In [121]: ad_click_train_data['is_weekend'] = ad_click_train_data['DateTime'].dt.dayofweek
print(ad_click_train_data[['DateTime', 'day_of_week', 'is_weekend']].head())
```

	DateTime	day_of_week	is_weekend
0	2017-07-02 00:00:00	6	True
1	2017-07-02 00:00:00	6	True
2	2017-07-02 00:00:00	6	True
3	2017-07-02 00:00:00	6	True
4	2017-07-02 00:01:00	6	True

```
In [122... def get_time_of_day(hour):
    if 0 <= hour < 6:
        return 'Night'
    elif 6 <= hour < 12:
        return 'Morning'
    elif 12 <= hour < 18:
        return 'Afternoon'
    else:
        return 'Evening'

ad_click_train_data['time_of_day'] = ad_click_train_data['hour_of_day'].apply(
    get_time_of_day)

print(ad_click_train_data[['DateTime', 'hour_of_day', 'time_of_day']].head())
```

	DateTime	hour_of_day	time_of_day
0	2017-07-02 00:00:00	0	Night
1	2017-07-02 00:00:00	0	Night
2	2017-07-02 00:00:00	0	Night
3	2017-07-02 00:00:00	0	Night
4	2017-07-02 00:01:00	0	Night

```
In [123... print(ad_click_train_data[['DateTime', 'hour_of_day', 'day_of_month', 'is_weekend',
    'time_of_day']])

    DateTime  hour_of_day  day_of_month  is_weekend  time_of_day
0 2017-07-02 00:00:00          0           2      True     Night
1 2017-07-02 00:00:00          0           2      True     Night
2 2017-07-02 00:00:00          0           2      True     Night
3 2017-07-02 00:00:00          0           2      True     Night
4 2017-07-02 00:01:00          0           2     True     Night
```

Interaction features

User-Product Interaction Feature

```
In [124... ad_click_train_data['user_product_interaction'] = ad_click_train_data['user_id'].map(
    user_id_map)

print(ad_click_train_data[['user_id', 'product', 'user_product_interaction']].head())

    user_id product user_product_interaction
0    858557      C            858557_C
1    243253      C            243253_C
2    243253      C            243253_C
3   1097446      I            1097446_I
4    663656      C            663656_C
```

Potential Use for Further Analysis:

This new interaction feature further enriches our dataset by capturing specific user-product relationships, enabling more granular analysis of user behavior:

- **user_product_interaction**: This feature can reveal patterns where specific users repeatedly engage with certain products, or conversely, if

a user clicks on a product once and then never again. This can be used to understand customer loyalty to specific product types or to identify products that are failing to retain user interest after the initial click. It can also be crucial for building recommendation systems or for targeted advertising campaigns based on individual user product preferences.

These interaction features are vital for understanding the underlying dynamics of ad clicks, allowing for more personalized and effective marketing strategies.

Campaign-Webpage Interaction Feature

```
In [125...]: ad_click_train_data['campaign_webpage'] = ad_click_train_data['campaign_id'].a  
print(ad_click_train_data[['campaign_id', 'webpage_id', 'campaign_webpage']].head(5))  
  
campaign_id  webpage_id  campaign_webpage  
0            359520      13787    359520_13787  
1            105960      11085    105960_11085  
2            359520      13787    359520_13787  
3            359520      13787    359520_13787  
4            405490      60305    405490_60305
```

Potential Use for Further Analysis:

This new interaction feature further enriches our dataset by capturing specific campaign-webpage relationships, enabling more granular analysis of their combined performance:

- **campaign_webpage**: This feature can be used to identify which specific campaigns perform best on particular webpages. For instance, it can highlight if a certain campaign creative is highly effective only when displayed on a specific landing page, or if some webpages consistently underperform regardless of the campaign. This insight is critical for optimizing ad placement, improving landing page content, and refining campaign strategies to maximize Click-Through Rates and conversion rates.

These interaction features are vital for understanding the underlying dynamics of ad clicks, allowing for more personalized and effective marketing strategies.

Gender-Age Interaction Feature

```
In [126...]: ad_click_train_data['gender_age'] = ad_click_train_data['gender'].astype(str)
```

```

print(ad_click_train_data[['gender', 'age_level', 'gender_age']].head())

```

	gender	age_level	gender_age
0	Female	4.0	Female_4.0
1	Female	2.0	Female_2.0
2	Female	2.0	Female_2.0
3	Male	3.0	Male_3.0
4	Male	2.0	Male_2.0

Potential Use for Further Analysis:

This new interaction feature further enriches our dataset by capturing specific gender-age relationships, enabling more granular analysis of user behavior:

- **gender_age** : This feature can be used to identify if certain ad creatives, products, or campaigns resonate more strongly with particular gender-age groups. For example, it might reveal that 'Female_2.0' (females in age level 2) have a significantly different click-through rate for a specific product category compared to 'Male_3.0' (males in age level 3). This insight is invaluable for highly targeted advertising, allowing marketers to tailor their messages and product offerings to the most receptive demographic segments. It can also help in understanding the diversity of user preferences and designing more inclusive or segmented marketing strategies.

These interaction features are vital for understanding the underlying dynamics of ad clicks, allowing for more personalized and effective marketing strategies.

Aggregated Features

User-Level Aggregations:

Create User-Level Aggregated Feature: `user_total_clicks`

```

In [127]: user_total_clicks = ad_click_train_data.groupby('user_id')['is_click'].sum().reset_index()
          user_total_clicks.rename(columns={'is_click': 'user_total_clicks'}, inplace=True)
          ad_click_train_data = pd.merge(ad_click_train_data, user_total_clicks, on='user_id')
          print(ad_click_train_data[['user_id', 'user_total_clicks']].head())

```

	user_id	user_total_clicks
0	858557	0
1	243253	0
2	243253	0
3	1097446	0
4	663656	2

Create User-Level Aggregated Feature: `user_total_views`

```
In [128]: user_total_views = ad_click_train_data.groupby('user_id')['session_id'].count()
user_total_views.rename(columns={'session_id': 'user_total_views'}, inplace=True)

ad_click_train_data = pd.merge(ad_click_train_data, user_total_views, on='user_id')

print(ad_click_train_data[['user_id', 'user_total_views']].head())
```

	user_id	user_total_views
0	858557	2
1	243253	3
2	243253	3
3	1097446	18
4	663656	11

Summary of New Aggregated Features

New Features Created:

- `user_total_clicks` : The total number of clicks each user has made across all sessions.
- `user_total_views` : The total number of ad impressions each user has seen across all sessions.

Potential Use for Further Analysis:

These aggregated features provide user-level behavioral insights that can be highly predictive and useful for advanced analytics and machine learning models:

- `user_total_clicks` : This feature quantifies how active and engaged a user is. Users with higher total clicks might represent a more valuable segment (e.g., highly interested buyers). It can be used to:
 - Identify 'power users' or highly engaged individuals.
 - Segment users based on their historical click behavior.
 - Serve as a critical input for predictive models to forecast future clicks or conversions.
 - Gauge the overall effectiveness of ad campaigns on an individual user basis.
- `user_total_views` : This feature indicates the exposure level of a user to ads. Combined with `user_total_clicks`, it can help derive an individual user's average CTR, which is a powerful metric for understanding their propensity to click. It can be used to:

- Identify users who are frequently exposed to ads but rarely click (low engagement).
- Help personalize ad frequency and avoid ad fatigue for over-exposed users.
- Improve targeting strategies by understanding the relationship between exposure and clicks at a user level.

By combining these two features, we can also derive a `user_CTR` (`user_total_clicks / user_total_views`), providing a personalized measure of click-through propensity. These features are essential for building robust recommendation systems, personalized marketing campaigns, and predictive models to optimize ad delivery and maximize engagement.

Create User-Level Aggregated Feature: `user_CTR`

```
In [129]: ad_click_train_data['user_CTR'] = (ad_click_train_data['user_total_clicks'] / ad_click_train_data['user_total_views'])

# Handle potential division by zero (if a user had 0 views, CTR should be 0 or less)
ad_click_train_data['user_CTR'] = ad_click_train_data['user_CTR'].fillna(0)

print(ad_click_train_data[['user_id', 'user_total_clicks', 'user_total_views',
                           'user_CTR']]
```

	user_id	user_total_clicks	user_total_views	user_CTR
0	858557	0	2	0.000000
1	243253	0	3	0.000000
2	243253	0	3	0.000000
3	1097446	0	18	0.000000
4	663656	2	11	18.181818

Summary of New Aggregated Feature: `user_CTR`

New Feature Created:

- `user_CTR` : Represents the individual Click-Through Rate for each user, calculated as (`user_total_clicks / user_total_views`) * 100.

Potential Use for Further Analysis:

This `user_CTR` feature provides a highly personalized metric that can significantly enhance our understanding of user behavior and improve predictive models:

- **Personalized Engagement Metric:** `user_CTR` quantifies a user's propensity to click on ads based on their historical interactions. A high `user_CTR` indicates a highly engaged user who is more likely to click.

in the future, while a low `user_CTR` suggests lower engagement.

- **Targeted Advertising:** This feature can be used to segment users based on their click propensity. Ads can be dynamically served to users with high `user_CTR` who are more likely to convert, or strategies can be developed to re-engage users with low `user_CTR`.
- **Fraud Detection/Bot Identification:** Anomalously high or low `user_CTR` values might indicate bot activity or fraudulent clicks, which can be investigated further.
- **Feature for Machine Learning Models:** `user_CTR` can serve as a powerful feature in predictive models to forecast future clicks, conversions, or user churn. It encapsulates a user's historical behavior into a single, highly informative metric.
- **Ad Fatigue Management:** By monitoring individual `user_CTR`, we can optimize ad frequency to avoid bombarding users who rarely click, thereby improving user experience and potentially reducing ad spend waste.

This `user_CTR` feature, alongside `user_total_clicks` and `user_total_views`, offers a comprehensive view of individual user engagement, paving the way for more sophisticated analysis and personalized marketing strategies.

Create User-Level Aggregated Feature: `user_sessions`

```
In [130]: user_sessions = ad_click_train_data.groupby('user_id')['session_id'].nunique()
user_sessions.rename(columns={'session_id': 'user_sessions'}, inplace=True)

ad_click_train_data = pd.merge(ad_click_train_data, user_sessions, on='user_id')

print(ad_click_train_data[['user_id', 'user_sessions']].head())
```

	user_id	user_sessions
0	858557	2
1	243253	3
2	243253	3
3	1097446	18
4	663656	11

Summary of New Aggregated Feature: `user_sessions`

New Feature Created:

- `user_sessions` : Represents the number of unique sessions per user.

Potential Use for Further Analysis:

This `user_sessions` feature provides insights into the breadth of a user's engagement, indicating how many distinct browsing instances they have had:

- **Engagement Breadth:** A higher number of unique sessions might indicate a more active and loyal user who returns frequently, or it could highlight users who are exploring various parts of the platform or different ad campaigns over time.
- **User Segmentation:** Users can be segmented based on their session count. For instance, users with very few sessions might be new or less engaged, while those with many sessions are highly involved. This can inform targeted re-engagement strategies or loyalty programs.
- **Ad Context:** Combined with other temporal features, `user_sessions` can help understand if users engage with ads more effectively during their first few sessions, or if engagement builds over time. It can also distinguish between single-session explorers and multi-session deep-divers.
- **Feature for Machine Learning Models:** `user_sessions` can serve as a valuable numerical feature in predictive models for click-through rate, conversion, or user churn, providing a quantitative measure of user activity and stickiness.

This feature, alongside `user_total_clicks`, `user_total_views`, and `user_CTR`, contributes to a holistic understanding of individual user behavior, enabling more sophisticated analysis and personalized marketing strategies.

Product-Level Aggregations:

Product-Level Aggregated Feature: `product_views`

```
In [131]: product_views = ad_click_train_data.groupby('product')['session_id'].count().reset_index()
product_views.rename(columns={'session_id': 'product_views'}, inplace=True)
```

```

ad_click_train_data = pd.merge(ad_click_train_data, product_views, on='product')

print(ad_click_train_data[['product', 'product_views']].head())

```

product	product_views	
0	C	163501
1	C	163501
2	C	163501
3	I	63711
4	C	163501

Product-Level Aggregated Feature: `product_ctr`

```

In [132]: product_total_clicks = ad_click_train_data.groupby('product')['is_click'].sum()
product_total_clicks.rename(columns={'is_click': 'product_total_clicks'}, inplace=True)

ad_click_train_data = pd.merge(ad_click_train_data, product_total_clicks, on='product')

ad_click_train_data['product_ctr'] = (ad_click_train_data['product_total_clicks'] / ad_click_train_data['product_views']) * 100

# Handle potential division by zero (if a product had 0 views, CTR should be 0)
ad_click_train_data['product_ctr'] = ad_click_train_data['product_ctr'].fillna(0)

print(ad_click_train_data[['product', 'product_total_clicks', 'product_views', 'product_ctr']])

```

product	product_total_clicks	product_views	product_ctr	
0	C	11306	163501	6.914942
1	C	11306	163501	6.914942
2	C	11306	163501	6.914942
3	I	4079	63711	6.402348
4	C	11306	163501	6.914942

Summary of New Aggregated Feature: `product_ctr`

New Feature Created:

- `product_ctr`: This product's historical click rate, calculated as $(\text{product_total_clicks} / \text{product_views}) * 100$.

Potential Use for Further Analysis:

This `product_ctr` feature provides a crucial metric for understanding product performance and can significantly enhance predictive models and marketing strategies:

- **Product Performance Evaluation:** `product_ctr` directly quantifies how engaging each product is to users. Products with high CTR are highly effective at capturing user interest, while low CTR products

might indicate issues with their appeal or advertising.

- **Marketing Strategy Optimization:** This feature can guide marketing teams in allocating resources. Products with high `product_ctr` could be highlighted in campaigns, or efforts could be focused on improving the advertising for products with low `product_ctr`.
- **Content and Creative Assessment:** Analyzing the creative content and messaging associated with high `product_ctr` products can reveal best practices that can be applied to other product categories.
- **Personalization:** When combined with user-level features, `product_ctr` can contribute to more personalized recommendations, ensuring users are shown products they are historically more likely to click on.
- **Feature for Machine Learning Models:** `product_ctr` serves as a powerful numerical feature for machine learning models aiming to predict future clicks, conversions, or overall product success. It encapsulates a product's historical engagement effectiveness into a single, informative metric.

Campaign-Level Aggregations:

Campaign-Level Aggregated Feature: `campaign_views`

```
In [133...]: campaign_views = ad_click_train_data.groupby('campaign_id')['session_id'].count()
campaign_views.rename(columns={'session_id': 'campaign_views'}, inplace=True)

ad_click_train_data = pd.merge(ad_click_train_data, campaign_views, on='campaign_id')

print(ad_click_train_data[['campaign_id', 'campaign_views']].head())
```

	campaign_id	campaign_views
0	359520	108155
1	105960	25781
2	359520	108155
3	359520	108155
4	405490	95973

Campaign-Level Aggregated Feature: `campaign_ctr`

```
In [134...]: campaign_total_clicks = ad_click_train_data.groupby('campaign_id')['is_click'].sum()
campaign_total_clicks.rename(columns={'is_click': 'campaign_total_clicks'}, inplace=True)
```

```

ad_click_train_data = pd.merge(ad_click_train_data, campaign_total_clicks, on='campaign_id')

ad_click_train_data['campaign_ctr'] = (ad_click_train_data['campaign_total_clicks'] / ad_click_train_data['campaign_views']) * 100

# Handle potential division by zero (if a campaign had 0 views, CTR should be 0)
ad_click_train_data['campaign_ctr'] = ad_click_train_data['campaign_ctr'].fillna(0)

print(ad_click_train_data[['campaign_id', 'campaign_total_clicks', 'campaign_views', 'campaign_ctr']])

```

	campaign_id	campaign_total_clicks	campaign_views	campaign_ctr
0	359520	6340	108155	5.861957
1	105960	1762	25781	6.834491
2	359520	6340	108155	5.861957
3	359520	6340	108155	5.861957
4	405490	8763	95973	9.130693

Summary of New Aggregated Feature: `campaign_ctr`

New Feature Created:

- `campaign_ctr` : This campaign's historical click rate, calculated as $(\text{campaign_total_clicks} / \text{campaign_views}) * 100$.

Potential Use for Further Analysis:

This `campaign_ctr` feature provides a crucial metric for understanding campaign performance and can significantly enhance predictive models and marketing strategies:

- **Campaign Performance Evaluation:** `campaign_ctr` directly quantifies how engaging each campaign is to users. Campaigns with high CTR are highly effective at capturing user interest, while low CTR campaigns might indicate issues with their appeal, targeting, or creative content.
- **Marketing Strategy Optimization:** This feature can guide marketing teams in allocating resources. Campaigns with high `campaign_ctr` could be prioritized and scaled, or efforts could be focused on improving the advertising for campaigns with low `campaign_ctr`.
- **A/B Testing and Creative Assessment:** Analyzing the creative content, messaging, and targeting strategies associated with high `campaign_ctr` campaigns can reveal best practices that can be applied to other campaigns.
- **Budget Allocation:** By understanding the historical performance of

campaigns, advertising budgets can be optimized to invest more in highly effective campaigns and less in underperforming ones.

- **Feature for Machine Learning Models:** `campaign_ctr` serves as a powerful numerical feature for machine learning models aiming to predict future clicks, conversions, or overall campaign success. It encapsulates a campaign's historical engagement effectiveness into a single, informative metric.

Encode Categorical Variables

Approach: Label Encoding

Columns to Encode:

- product
- campaign_id
- webpage_id
- product_category_1
- product_category_2
- gender
- user_group_id
- var_1
- All interaction features (user_product_interaction, etc.)

```
In [135]: print("Unique values in 'product_category_2' before handling missing values:")
print(ad_click_train_data['product_category_2'].value_counts(dropna=False).head(10))

missing_product_category_2 = ad_click_train_data['product_category_2'].isnull()
print(f"\nNumber of missing values in 'product_category_2': {missing_product_c

if missing_product_category_2 > 0:
    mode_product_category_2 = ad_click_train_data['product_category_2'].mode()
    ad_click_train_data['product_category_2'] = ad_click_train_data['product_c
    print(f"Missing values in 'product_category_2' filled with mode: {mode_pr

print("\nValue counts for 'product_category_2' after handling missing values:")
print(ad_click_train_data['product_category_2'].value_counts(dropna=False).head(10))
```

```
Unique values in 'product_category_2' before handling missing values:  
product_category_2  
NaN          365854  
82527.0      35531  
146115.0     25224  
270915.0     19624  
254132.0     12502  
Name: count, dtype: int64
```

```
Number of missing values in 'product_category_2': 365854  
Missing values in 'product_category_2' filled with mode: 82527.0
```

```
Value counts for 'product_category_2' after handling missing values:  
product_category_2  
82527.0      401385  
146115.0     25224  
270915.0     19624  
254132.0     12502  
143597.0     2701  
Name: count, dtype: int64
```

```
In [136]: from sklearn.preprocessing import LabelEncoder  
  
columns_to_encode = [  
    'product', 'campaign_id', 'webpage_id', 'product_category_1', 'product_cat',  
    'gender', 'user_group_id', 'var_1', 'user_product_interaction', 'campaign_id'  
]  
  
# Dictionary to store LabelEncoders for potential inverse transformation later  
label_encoders = {}  
  
print("Encoding categorical variables...")  
for col in columns_to_encode:  
    le = LabelEncoder()  
    ad_click_train_data[col] = le.fit_transform(ad_click_train_data[col])  
    label_encoders[col] = le  
    print(f" - Encoded '{col}'")  
  
print("\nFirst 5 rows of the DataFrame with encoded columns:")  
print(ad_click_train_data[columns_to_encode].head())
```

```

Encoding categorical variables...
- Encoded 'product'
- Encoded 'campaign_id'
- Encoded 'webpage_id'
- Encoded 'product_category_1'
- Encoded 'product_category_2'
- Encoded 'gender'
- Encoded 'user_group_id'
- Encoded 'var_1'
- Encoded 'user_product_interaction'
- Encoded 'campaign_webpage'
- Encoded 'gender_age'

First 5 rows of the DataFrame with encoded columns:
   product  campaign_id  webpage_id  product_category_1  product_category_2 \
0         2             4            3                  3                  3
1         2             2            2                  4                  3
2         2             4            3                  3                  3
3         8             4            3                  2                  3
4         2             8            8                  2                  3

   gender  user_group_id  var_1  user_product_interaction  campaign_webpage \
0       0             10      0                 215527                  2
1       0              8      0                 66859                  0
2       0              8      0                 66859                  2
3       1              3      1                 22716                  2
4       1              2      1                172537                  6

   gender_age
0           4
1           2
2           2
3          10
4           9

```

Summary of Encoded Features

Encoded Columns:

All specified categorical columns and interaction features have been successfully encoded using `LabelEncoder`:

- `product`
- `campaign_id`
- `webpage_id`
- `product_category_1`
- `product_category_2`
- `gender`
- `user_group_id`

- `var_1`
- `user_product_interaction`
- `campaign_webpage`
- `gender_age`

Data Transformation Key Findings:

- Missing values in `product_category_2` were imputed with the mode (`82527.0`) prior to encoding, ensuring data completeness.
- Each unique category within the specified columns has been converted into a unique numerical label (starting from 0).

Potential Use for Further Analysis:

- **Machine Learning Model Compatibility:** The primary benefit of label encoding is to convert categorical data into a numerical format that machine learning algorithms (which typically require numerical input) can process directly.
- **Reduced Dimensionality:** For features with many unique categories (especially interaction features), label encoding provides a compact numerical representation, avoiding the high dimensionality that could result from one-hot encoding.
- **Improved Model Performance:** By transforming these features, the dataset is now better prepared for training models like Decision Trees, Random Forests, Gradient Boosting Machines, or neural networks, potentially leading to improved model performance in predicting ad clicks.
- **Feature Importance:** The numerical nature of these encoded features allows models to assess their importance in predicting the target variable, helping to identify which categorical aspects have the most significant impact on click-through rates.

Columns to drop:

- `'DateTime', 'session_id', 'user_id'`

Reasoning: The `'DateTime'` column has already been used to extract temporal features, and `'session_id'` and `'user_id'` have been used to create aggregated and interaction features. These original columns are no longer needed for direct model training and can be dropped to reduce dimensionality and avoid potential data leakage (especially for `'session_id'` and `'user_id'` if not handled carefully in

modeling). I will use the `drop()` method with `axis=1` to remove them from the DataFrame.

```
In [137]: columns_to_drop = ['DateTime', 'session_id', 'user_id']
ad_click_train_data = ad_click_train_data.drop(columns=columns_to_drop)

print("First 5 rows of the DataFrame after dropping columns:")
print(ad_click_train_data.head())
print("\nRemaining columns in the DataFrame:")
print(ad_click_train_data.columns)
```

```

First 5 rows of the DataFrame after dropping columns:
   product  campaign_id  webpage_id  product_category_1  product_category_2 \
0         2            4            3                  3                  3
1         2            2            2                  4                  3
2         2            4            3                  3                  3
3         8            4            3                  2                  3
4         2            8            8                  2                  3

   user_group_id  gender  age_level  user_depth  city_development_index  ...
0             10       0        4.0      3.0              3.0          ...
1              8       0        2.0      2.0            NaN          ...
2              8       0        2.0      2.0            NaN          ...
3              3       1        3.0      3.0              2.0          ...
4              2       1        2.0      3.0              2.0          ...

   user_total_clicks  user_total_views  user_CTR  user_sessions \
0                 0                  2  0.000000               2
1                 0                  3  0.000000               3
2                 0                  3  0.000000               3
3                 0                 18  0.000000              18
4                 2                 11  18.181818              11

   product_views  product_total_clicks  product_ctr  campaign_views \
0     163501                  11306    6.914942      108155
1     163501                  11306    6.914942      25781
2     163501                  11306    6.914942      108155
3     63711                   4079    6.402348      108155
4     163501                  11306    6.914942      95973

   campaign_total_clicks  campaign_ctr
0                  6340    5.861957
1                  1762    6.834491
2                  6340    5.861957
3                  6340    5.861957
4                  8763    9.130693

```

[5 rows x 31 columns]

Remaining columns in the DataFrame:

```

Index(['product', 'campaign_id', 'webpage_id', 'product_category_1',
       'product_category_2', 'user_group_id', 'gender', 'age_level',
       'user_depth', 'city_development_index', 'var_1', 'is_click',
       'hour_of_day', 'day_of_week', 'day_type', 'day_of_month', 'is_weekend',
       'time_of_day', 'user_product_interaction', 'campaign_webpage',
       'gender_age', 'user_total_clicks', 'user_total_views', 'user_CTR',
       'user_sessions', 'product_views', 'product_total_clicks', 'product_ctr',
       'campaign_views', 'campaign_total_clicks', 'campaign_ctr'],
      dtype='object')

```

Perform the same feature engineering and data preprocessing tasks for the ad_click_test_data as done for ad_click_train_data.

Prepare Test Data - Temporal Features

```
In [138...]: ad_click_test_data['DateTime'] = pd.to_datetime(ad_click_test_data['DateTime'])
ad_click_test_data['hour_of_day'] = ad_click_test_data['DateTime'].dt.hour
ad_click_test_data['day_of_week'] = ad_click_test_data['DateTime'].dt.dayofweek
ad_click_test_data['day_of_month'] = ad_click_test_data['DateTime'].dt.day
ad_click_test_data['is_weekend'] = ad_click_test_data['DateTime'].dt.dayofweek > 5
ad_click_test_data['day_type'] = ad_click_test_data['day_of_week'].apply(lambda x: 'Weekday' if x < 5 else 'Weekend')
ad_click_test_data['time_of_day'] = ad_click_test_data['hour_of_day'].apply(lambda x: 'Night' if x > 21 else 'Day')

print(ad_click_test_data[['DateTime', 'hour_of_day', 'day_of_week', 'day_of_month', 'day_type', 'is_weekend', 'time_of_day']])
```

	DateTime	hour_of_day	day_of_week	day_of_month	day_type	is_weekend	time_of_day
0	2017-07-08	0	5	8	Weekend	True	Night
1	2017-07-08	0	5	8	Weekend	True	Night
2	2017-07-08	0	5	8	Weekend	True	Night
3	2017-07-08	0	5	8	Weekend	True	Night
4	2017-07-08	0	5	8	Weekend	True	Night

```
In [139...]: print("Unique values in 'age_level' before handling missing values:")
print(ad_click_test_data['age_level'].value_counts(dropna=False).head())

missing_age_level_test = ad_click_test_data['age_level'].isnull().sum()
print(f"\nNumber of missing values in 'age_level' (test data): {missing_age_level_test}")

if missing_age_level_test > 0:
    ad_click_test_data['age_level'] = ad_click_test_data['age_level'].fillna(mode['age_level'])
    print(f"Missing values in 'age_level' (test data) filled with mode from training data")

print("\nUnique values in 'gender' before handling missing values:")
print(ad_click_test_data['gender'].value_counts(dropna=False).head())

missing_gender_test = ad_click_test_data['gender'].isnull().sum()
print(f"\nNumber of missing values in 'gender' (test data): {missing_gender_test}")

if missing_gender_test > 0:
    ad_click_test_data['gender'] = ad_click_test_data['gender'].fillna(mode['gender'])
    print(f"Missing values in 'gender' (test data) filled with mode from training data")

print("\nUnique values in 'user_group_id' before handling missing values:")
print(ad_click_test_data['user_group_id'].value_counts(dropna=False).head())
```

```
missing_user_group_id_test = ad_click_test_data['user_group_id'].isnull().sum()
print(f"\nNumber of missing values in 'user_group_id' (test data): {missing_us}

if missing_user_group_id_test > 0:
    ad_click_test_data['user_group_id'] = ad_click_test_data['user_group_id'].fillna(
        mode[ad_click_test_data['user_group_id'].isnull()])
    print(f"Missing values in 'user_group_id' (test data) filled with mode from column")

print("\nUnique values in 'product_category_2' before handling missing values:")
print(ad_click_test_data['product_category_2'].value_counts(dropna=False).head())

missing_product_category_2_test = ad_click_test_data['product_category_2'].isnull().sum()
print(f"\nNumber of missing values in 'product_category_2' (test data): {missing_pc_2}

if missing_product_category_2_test > 0:
    ad_click_test_data['product_category_2'] = ad_click_test_data['product_category_2'].fillna(
        mode[ad_click_test_data['product_category_2'].isnull()])
    print(f"Missing values in 'product_category_2' (test data) filled with mode from column")

print("\nValue counts for 'age_level' after handling missing values:")
print(ad_click_test_data['age_level'].value_counts(dropna=False).head())
print("\nValue counts for 'gender' after handling missing values:")
print(ad_click_test_data['gender'].value_counts(dropna=False).head())
print("\nValue counts for 'user_group_id' after handling missing values:")
print(ad_click_test_data['user_group_id'].value_counts(dropna=False).head())
print("\nValue counts for 'product_category_2' after handling missing values:")
print(ad_click_test_data['product_category_2'].value_counts(dropna=False).head())
```

```
Unique values in 'age_level' before handling missing values:  
age_level  
3.0    43234  
2.0    38711  
4.0    18884  
1.0    11896  
5.0    9565  
Name: count, dtype: int64  
  
Number of missing values in 'age_level' (test data): 5684  
Missing values in 'age_level' (test data) filled with mode from train data: 3.0  
  
Unique values in 'gender' before handling missing values:  
gender  
Male      108525  
Female     14649  
NaN        5684  
Name: count, dtype: int64  
  
Number of missing values in 'gender' (test data): 5684  
Missing values in 'gender' (test data) filled with mode from train data: Male  
  
Unique values in 'user_group_id' before handling missing values:  
user_group_id  
3.0    37785  
2.0    37093  
4.0    14857  
1.0    11510  
5.0    6730  
Name: count, dtype: int64  
  
Number of missing values in 'user_group_id' (test data): 5684  
Missing values in 'user_group_id' (test data) filled with mode from train data:  
3.0  
  
Unique values in 'product_category_2' before handling missing values:  
product_category_2  
NaN        76171  
82527.0    36191  
270915.0    8318  
146115.0    5692  
254132.0    1617  
Name: count, dtype: int64  
  
Number of missing values in 'product_category_2' (test data): 76171  
Missing values in 'product_category_2' (test data) filled with mode from train  
data: 82527.0  
  
Value counts for 'age_level' after handling missing values:  
age_level  
3.0    48918  
2.0    38711  
4.0    18884  
1.0    11896
```

```

5.0    9565
Name: count, dtype: int64

Value counts for 'gender' after handling missing values:
gender
Male      114209
Female     14649
Name: count, dtype: int64

Value counts for 'user_group_id' after handling missing values:
user_group_id
3.0      43469
2.0      37093
4.0      14857
1.0      11510
5.0      6730
Name: count, dtype: int64

Value counts for 'product_category_2' after handling missing values:
product_category_2
82527.0    112362
270915.0    8318
146115.0    5692
254132.0    1617
269093.0    408
Name: count, dtype: int64

```

```

In [140]: ad_click_test_data['user_product_interaction'] = ad_click_test_data['user_id']
ad_click_test_data['campaign_webpage'] = ad_click_test_data['campaign_id'].astype(str)
ad_click_test_data['gender_age'] = ad_click_test_data['gender'].astype(str) + '_'

print(ad_click_test_data[['user_id', 'product', 'user_product_interaction', 'campaign_id',
                         'webpage_id', 'gender', 'age_level', 'gender_age']])

```

	user_id	product	user_product_interaction	campaign_id	webpage_id	gender	age_level	gender_age
0	732573	J	732573_J	404347	53587	Male	5.0	Male_5.0
1	172910	I	172910_I	118601	28529	Male	3.0	Male_3.0
2	172910	I	172910_I	118601	28529	Male	3.0	Male_3.0
3	557318	G	557318_G	118601	28529	Male	1.0	Male_1.0
4	923896	H	923896_H	118601	28529	Female	3.0	Female_3.0

```

In [141]: user_total_views_test = ad_click_test_data.groupby('user_id')['session_id'].count()
user_total_views_test.rename(columns={'session_id': 'user_total_views'}, inplace=True)

ad_click_test_data = pd.merge(ad_click_test_data, user_total_views_test, on='user_id')

print(ad_click_test_data[['user_id', 'user_total_views']].head())

```

```

      user_id  user_total_views
0    732573              1
1   172910              27
2   172910              27
3   557318              1
4   923896              1

In [142... product_views_test = ad_click_test_data.groupby('product')['session_id'].count()
product_views_test.rename(columns={'session_id': 'product_views'}, inplace=True)

ad_click_test_data = pd.merge(ad_click_test_data, product_views_test, on='product')

print(ad_click_test_data[['product', 'product_views']].head())

```

product	product_views
J	2453
I	12727
I	12727
G	3011
H	44880


```

In [143... campaign_views_test = ad_click_test_data.groupby('campaign_id')['session_id'].count()
campaign_views_test.rename(columns={'session_id': 'campaign_views'}, inplace=True)

ad_click_test_data = pd.merge(ad_click_test_data, campaign_views_test, on='campaign_id')

print(ad_click_test_data[['campaign_id', 'campaign_views']].head())

```

campaign_id	campaign_views
404347	6514
118601	36191
118601	36191
118601	36191
118601	36191


```

In [144... print("Unique values in 'product_category_2' before handling missing values:")
print(ad_click_test_data['product_category_2'].value_counts(dropna=False).head())

missing_product_category_2 = ad_click_test_data['product_category_2'].isnull()
print(f"\nNumber of missing values in 'product_category_2': {missing_product_category_2.sum()}\n")

if missing_product_category_2.sum() > 0:
    mode_product_category_2 = ad_click_test_data['product_category_2'].mode()
    ad_click_test_data['product_category_2'] = ad_click_test_data['product_category_2'].fillna(mode_product_category_2)
    print(f"Missing values in 'product_category_2' filled with mode: {mode_product_category_2}\n")

print("\nValue counts for 'product_category_2' after handling missing values:")
print(ad_click_test_data['product_category_2'].value_counts(dropna=False).head())

```

```
Unique values in 'product_category_2' before handling missing values:  
product_category_2  
82527.0      112362  
270915.0      8318  
146115.0      5692  
254132.0      1617  
269093.0      408  
Name: count, dtype: int64
```

Number of missing values in 'product_category_2': 0

```
Value counts for 'product_category_2' after handling missing values:  
product_category_2  
82527.0      112362  
270915.0      8318  
146115.0      5692  
254132.0      1617  
269093.0      408  
Name: count, dtype: int64
```

```
In [145...]: from sklearn.preprocessing import LabelEncoder  
  
columns_to_encode = [  
    'product', 'campaign_id', 'webpage_id', 'product_category_1', 'product_cat'  
    'gender', 'user_group_id', 'var_1', 'user_product_interaction', 'campaign_]  
  
# Dictionary to store LabelEncoders for potential inverse transformation later  
label_encoders = {}  
  
print("Encoding categorical variables...")  
for col in columns_to_encode:  
    le = LabelEncoder()  
    ad_click_test_data[col] = le.fit_transform(ad_click_test_data[col])  
    label_encoders[col] = le  
    print(f" - Encoded '{col}'")  
  
print("\nFirst 5 rows of the DataFrame with encoded columns:")  
print(ad_click_test_data[columns_to_encode].head())
```

```

Encoding categorical variables...
- Encoded 'product'
- Encoded 'campaign_id'
- Encoded 'webpage_id'
- Encoded 'product_category_1'
- Encoded 'product_category_2'
- Encoded 'gender'
- Encoded 'user_group_id'
- Encoded 'var_1'
- Encoded 'user_product_interaction'
- Encoded 'campaign_webpage'
- Encoded 'gender_age'

```

First 5 rows of the DataFrame with encoded columns:

	product	campaign_id	webpage_id	product_category_1	product_category_2	\
0	9		7	7	0	2
1	8		3	4	2	2
2	8		3	4	3	2
3	6		3	4	4	2
4	7		3	4	4	2

	gender	user_group_id	var_1	user_product_interaction	campaign_webpage	\
0	1		5	0	69652	5
1	1		3	1	17739	1
2	1		3	1	17739	1
3	1		1	0	54228	1
4	0		9	1	85041	1

	gender_age
0	12
1	10
2	10
3	8
4	3

```

In [146]: # --- Step 1: Impute remaining missing values for ad_click_train_data ---
# Compute mode for city_development_index from training data
mode_city_development_index = ad_click_train_data['city_development_index'].mode()
ad_click_train_data['city_development_index'] = ad_click_train_data['city_development_index'].fillna(mode_city_development_index)
print(f"Missing values in city_development_index (train data) filled with mode: {mode_city_development_index}")

# Compute mode for user_depth from training data
mode_user_depth = ad_click_train_data['user_depth'].mode()[0]
ad_click_train_data['user_depth'] = ad_click_train_data['user_depth'].fillna(mode_user_depth)
print(f"Missing values in user_depth (train data) filled with mode: {mode_user_depth}")

# --- Step 2: Impute remaining missing values for ad_click_test_data (using mode from train data)
ad_click_test_data['age_level'] = ad_click_test_data['age_level'].fillna(mode_age_level)
ad_click_test_data['gender'] = ad_click_test_data['gender'].fillna(mode_gender)
ad_click_test_data['user_group_id'] = ad_click_test_data['user_group_id'].fillna(mode_user_group_id)
ad_click_test_data['product_category_2'] = ad_click_test_data['product_category_2'].fillna(mode_product_category_2)
ad_click_test_data['city_development_index'] = ad_click_test_data['city_development_index'].fillna(mode_city_development_index)
ad_click_test_data['user_depth'] = ad_click_test_data['user_depth'].fillna(mode_user_depth)
print("\nMissing values in ad_click_test_data imputed using modes from train data")

```

```

# --- Step 3: Convert 'is_weekend' to integer type for both train and test data
ad_click_train_data['is_weekend'] = ad_click_train_data['is_weekend'].astype(int)
ad_click_test_data['is_weekend'] = ad_click_test_data['is_weekend'].astype(int)
print("\n'is_weekend' column converted to integer type for both train and test")

# --- Step 4: Re-Label Encode all specified categorical features ---
columns_to_encode_complete = [
    'product', 'campaign_id', 'webpage_id', 'product_category_1', 'product_cat',
    'gender', 'user_group_id', 'var_1', 'user_product_interaction', 'campaign',
    'gender_age', 'day_type', 'time_of_day'
]

# Reinitialize label_encoders to ensure fresh encoders are used
label_encoders = {}

print("\nRe-encoding all categorical variables (including day_type and time_of")
for col in columns_to_encode_complete:
    le = LabelEncoder()
    # Fit on training data and transform
    ad_click_train_data[col] = le.fit_transform(ad_click_train_data[col])
    label_encoders[col] = le # Store the encoder

    # Transform test data, handling unseen categories
    # Convert categories to string for consistent comparison, as data types might differ
    train_labels = set(le.classes_.astype(str))
    test_labels = set(ad_click_test_data[col].astype(str).unique())
    unseen_labels = test_labels - train_labels

    if unseen_labels:
        # Create a mapping for known labels and assign -1 for unknown/unseen labels
        mapping = {label: idx for idx, label in enumerate(le.classes_)}
        ad_click_test_data[col] = ad_click_test_data[col].astype(str).map(mapping)
        print(f" - Encoded '{col}' (with unseen categories mapped to -1)")
    else:
        ad_click_test_data[col] = le.transform(ad_click_test_data[col])
        print(f" - Encoded '{col}'")

# --- Step 5: Define features (X) and target (y) for ad_click_train_data ---
# Columns to drop from training features to avoid data leakage and include only relevant ones
columns_to_exclude_from_X = [
    'is_click', # Target variable itself
    'DateTime', 'session_id', 'user_id', # Original identifier columns
    # Target-dependent aggregated features - these should not be used as features
    'user_total_clicks', 'user_CTR', 'user_sessions',
    'product_total_clicks', 'product_ctr',
    'campaign_total_clicks', 'campaign_ctr'
]
X = ad_click_train_data.drop(columns=columns_to_exclude_from_X, errors='ignore')
y = ad_click_train_data['is_click']

print("\nFeatures for training (X) and target (y) defined.")
print(f"X columns: {X.columns.tolist()}")

```

```

# --- Step 6: Split training data into X_train, X_test, y_train, y_test ---
X_train, X_test, y_train, y_test = train_test_split(
    X,
    y,
    test_size=0.2,
    random_state=42,
    stratify=y
)

print(f"\nShape of X_train after splitting: {X_train.shape}")
print(f"Shape of X_test after splitting: {X_test.shape}")
print(f"Shape of y_train after splitting: {y_train.shape}")
print(f"Shape of y_test after splitting: {y_test.shape}")

# --- Step 7: Standard Scale X_train and X_test ---
numerical_features_to_scale = X_train.columns.tolist() # All columns in X_train
scaler = StandardScaler()

X_train[numerical_features_to_scale] = scaler.fit_transform(X_train[numerical_features_to_scale])
X_test[numerical_features_to_scale] = scaler.transform(X_test[numerical_features_to_scale])

print("\nTrain and test features (X_train, X_test) scaled.")

# --- Step 8: Prepare and Standard Scale ad_click_test_data ---
# Drop original identifier columns from ad_click_test_data
columns_to_drop_from_ad_click_test = ['DateTime', 'session_id', 'user_id']
ad_click_test_data = ad_click_test_data.drop(columns=columns_to_drop_from_ad_click_test)

# Ensure ad_click_test_data has the exact same columns as X_train before scaling
# This is critical for model prediction consistency
# Filter ad_click_test_data to match X_train's columns (numerical_features_to_scale)
missing_cols_in_test = set(numerical_features_to_scale) - set(ad_click_test_data.columns)
for col in missing_cols_in_test:
    ad_click_test_data[col] = 0 # Or some other sensible default if a feature is missing

# Ensure order of columns is identical
ad_click_test_data = ad_click_test_data[numerical_features_to_scale]

# Scale ad_click_test_data using the *fitted* scaler from X_train
ad_click_test_data[numerical_features_to_scale] = scaler.transform(ad_click_test_data)

print("\nad_click_test_data prepared and scaled with consistent features.")

print("\nFirst 5 rows of X_train after all processing:")
print(X_train.head())
print("\nFirst 5 rows of X_test after all processing:")
print(X_test.head())
print("\nFirst 5 rows of ad_click_test_data after all processing:")
print(ad_click_test_data.head())

```

```
Missing values in city_development_index (train data) filled with mode: 2.0
Missing values in user_depth (train data) filled with mode: 3.0
```

```
Missing values in ad_click_test_data imputed using modes from train data.
```

```
'is_weekend' column converted to integer type for both train and test data.
```

```
Re-encoding all categorical variables (including day_type and time_of_day)...
```

- Encoded 'product'
- Encoded 'campaign_id'
- Encoded 'webpage_id'
- Encoded 'product_category_1'
- Encoded 'product_category_2'
- Encoded 'gender'
- Encoded 'user_group_id'
- Encoded 'var_1'
- Encoded 'user_product_interaction'
- Encoded 'campaign_webpage'
- Encoded 'gender_age'
- Encoded 'day_type'
- Encoded 'time_of_day'

```
Features for training (X) and target (y) defined.
```

```
X columns: ['product', 'campaign_id', 'webpage_id', 'product_category_1', 'prod
uct_category_2', 'user_group_id', 'gender', 'age_level', 'user_depth', 'city_de
velopment_index', 'var_1', 'hour_of_day', 'day_of_week', 'day_type', 'day_of_mo
nth', 'is_weekend', 'time_of_day', 'user_product_interaction', 'campaign_webpag
e', 'gender_age', 'user_total_views', 'product_views', 'campaign_views']
```

```
Shape of X_train after splitting: (370632, 23)
```

```
Shape of X_test after splitting: (92659, 23)
```

```
Shape of y_train after splitting: (370632,)
```

```
Shape of y_test after splitting: (92659,)
```

```
Train and test features (X_train, X_test) scaled.
```

```
ad_click_test_data prepared and scaled with consistent features.
```

```
First 5 rows of X_train after all processing:
```

	product	campaign_id	webpage_id	product_category_1	product_category_2	user_group_id	gender	age_level	user_depth	city_development_index	... day_type	day_of_month	is_weekend	
351608	1.352408	-0.720071	-0.104083	0.711161	-0.324251	-0.194444	0.354469	0.198832	0.296805	0.982223	-0.339946	-0.502705	-0.055088	
205770	0.982223	-0.339946	-0.502705	-0.055088	-0.324251	-0.194444	0.354469	0.198832	0.296805	337312	0.982223	-0.720071	-0.104083	1.477410
364075	1.352408	-0.720071	-0.104083	0.711161	-0.324251	-0.194444	0.354469	0.198832	0.296805	309904	-0.128333	-1.860445	-1.698569	-1.587586
351608	-0.128333	-1.860445	-1.698569	-1.587586	0.586303	-1.038745	0.354469	-1.706770	0.296805	205770	-0.324251	-0.194444	0.354469	0.198832
337312	-0.324251	-0.194444	0.354469	0.198832	0.324251	-0.194444	0.354469	0.198832	0.296805	364075	-0.324251	-0.194444	0.354469	0.198832
309904	0.586303	-1.038745	0.354469	-1.706770	0.324251	-0.194444	0.354469	0.198832	0.296805	351608	1.352408	-0.720071	-0.104083	0.711161

351608	0.717733	...	-0.453113	0.908089	-0.453113
205770	0.717733	...	-0.453113	-0.269566	-0.453113
337312	-0.493986	...	-0.453113	0.908089	-0.453113
364075	-0.493986	...	-0.453113	0.908089	-0.453113
309904	0.717733	...	-0.453113	0.319261	-0.453113
	time_of_day	user_product_interaction	campaign_webpage	gender_age	\
351608	-1.224459		1.012664	-1.239874	0.448887
205770	-1.224459		0.631201	-0.858588	0.448887
337312	0.857507		0.428350	-1.239874	0.448887
364075	-1.224459		0.752695	-1.239874	0.448887
309904	-0.183476		-1.577412	1.429128	-0.457917
	user_total_views	product_views	campaign_views		
351608	-0.421154	-0.623283	-0.831644		
205770	-0.320275	0.185272	1.261348		
337312	-0.168956	0.185272	-0.831644		
364075	-0.320275	-0.623283	-0.831644		
309904	1.445115	-1.368300	-1.053036		

[5 rows x 23 columns]

First 5 rows of X_test after all processing:

	product	campaign_id	webpage_id	product_category_1	\	
388124	-0.868704	1.180552	1.490402	-0.055088		
86059	-0.868704	0.040178	-0.502705	1.477410		
223839	-0.868704	1.180552	1.490402	-0.055088		
70278	-0.498518	-1.480320	-1.299947	1.477410		
443481	-0.868704	0.040178	-0.502705	1.477410		
	product_category_2	user_group_id	gender	age_level	user_depth	\
388124	-0.324251	-0.616594	0.354469	-0.753969	-2.253515	
86059	-0.324251	0.227707	0.354469	1.151633	-2.253515	
223839	-0.324251	-0.194444	0.354469	0.198832	0.296805	
70278	-0.324251	-0.616594	0.354469	-0.753969	0.296805	
443481	-0.324251	0.227707	0.354469	1.151633	0.296805	
	city_development_index	...	day_type	day_of_month	is_weekend	\
388124	-0.493986	...	-0.453113	0.908089	-0.453113	
86059	-0.493986	...	-0.453113	-0.858394	-0.453113	
223839	-0.493986	...	-0.453113	-0.269566	-0.453113	
70278	0.717733	...	2.206956	-1.447222	2.206956	
443481	-0.493986	...	-0.453113	1.496917	-0.453113	
	time_of_day	user_product_interaction	campaign_webpage	gender_age	\	
388124	-0.183476		-0.768813	0.666556	-0.004515	
86059	0.857507		-1.258409	-0.477302	0.902289	
223839	-0.183476		1.610365	0.666556	0.448887	
70278	-0.183476		1.202972	1.810414	-0.004515	
443481	-1.224459		0.711826	-0.477302	0.902289	
	user_total_views	product_views	campaign_views			
388124	-0.421154	1.135993	0.910268			
86059	-0.572473	1.135993	-0.360242			

```

223839      -0.471594      1.135993      0.910268
70278       1.192916     -1.022544     -0.845074
443481      -0.572473      1.135993     -0.360242

```

[5 rows x 23 columns]

First 5 rows of ad_click_test_data after all processing:

	product	campaign_id	webpage_id	product_category_1	product_category_2	\
0	1.722593	0.800428	1.091781	-1.587586	-0.62777	
1	1.352408	-0.720071	-0.104083	-0.055088	-0.62777	
2	1.352408	-0.720071	-0.104083	0.711161	-0.62777	
3	0.612037	-0.720071	-0.104083	1.477410	-0.62777	
4	0.982223	-0.720071	-0.104083	1.477410	-0.62777	
	user_group_id	gender	age_level	user_depth	city_development_index	\
0	0.649858	0.354469	2.104433	0.296805	-0.493986	
1	-0.194444	0.354469	0.198832	0.296805	-0.493986	
2	-0.194444	0.354469	0.198832	0.296805	-0.493986	
3	-1.038745	0.354469	-1.706770	0.296805	-1.705706	
4	2.338460	-2.821124	0.198832	-4.803835	-0.493986	
	...	day_type	day_of_month	is_weekend	time_of_day	\
0	...	2.206956	2.085745	2.206956	1.89849	
1	...	2.206956	2.085745	2.206956	1.89849	
2	...	2.206956	2.085745	2.206956	1.89849	
3	...	2.206956	2.085745	2.206956	1.89849	
4	...	2.206956	2.085745	2.206956	1.89849	
	user_product_interaction	campaign_webpage	gender_age	user_total_views	\	
0	-0.746970	0.285270	1.355691	-0.572473		
1	-1.478566	-1.239874	0.448887	0.738959		
2	-1.478566	-1.239874	0.448887	0.738959		
3	-0.964336	-1.239874	-0.457917	-0.572473		
4	-0.530097	-1.239874	-2.724927	-0.572473		
	product_views	campaign_views				
0	-1.703248	-1.667901				
1	-1.522120	-0.812623				
2	-1.522120	-0.812623				
3	-1.693410	-0.812623				
4	-0.955269	-0.812623				

[5 rows x 23 columns]

```

In [147... # Train a Logistic Regression model
logistic_model = LogisticRegression(solver='liblinear', random_state=42)
logistic_model.fit(X_train, y_train)

print("Logistic Regression model trained successfully.")

```

Logistic Regression model trained successfully.

```

In [148... # Make predictions on the test set
y_pred = logistic_model.predict(X_test)

```

```

y_pred_proba = logistic_model.predict_proba(X_test)[:, 1]

# Calculate evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

# Calculate ROC curve and AUC
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)

print("Model Evaluation on Test Set:")
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")
print(f"ROC AUC: {roc_auc:.4f}")

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```

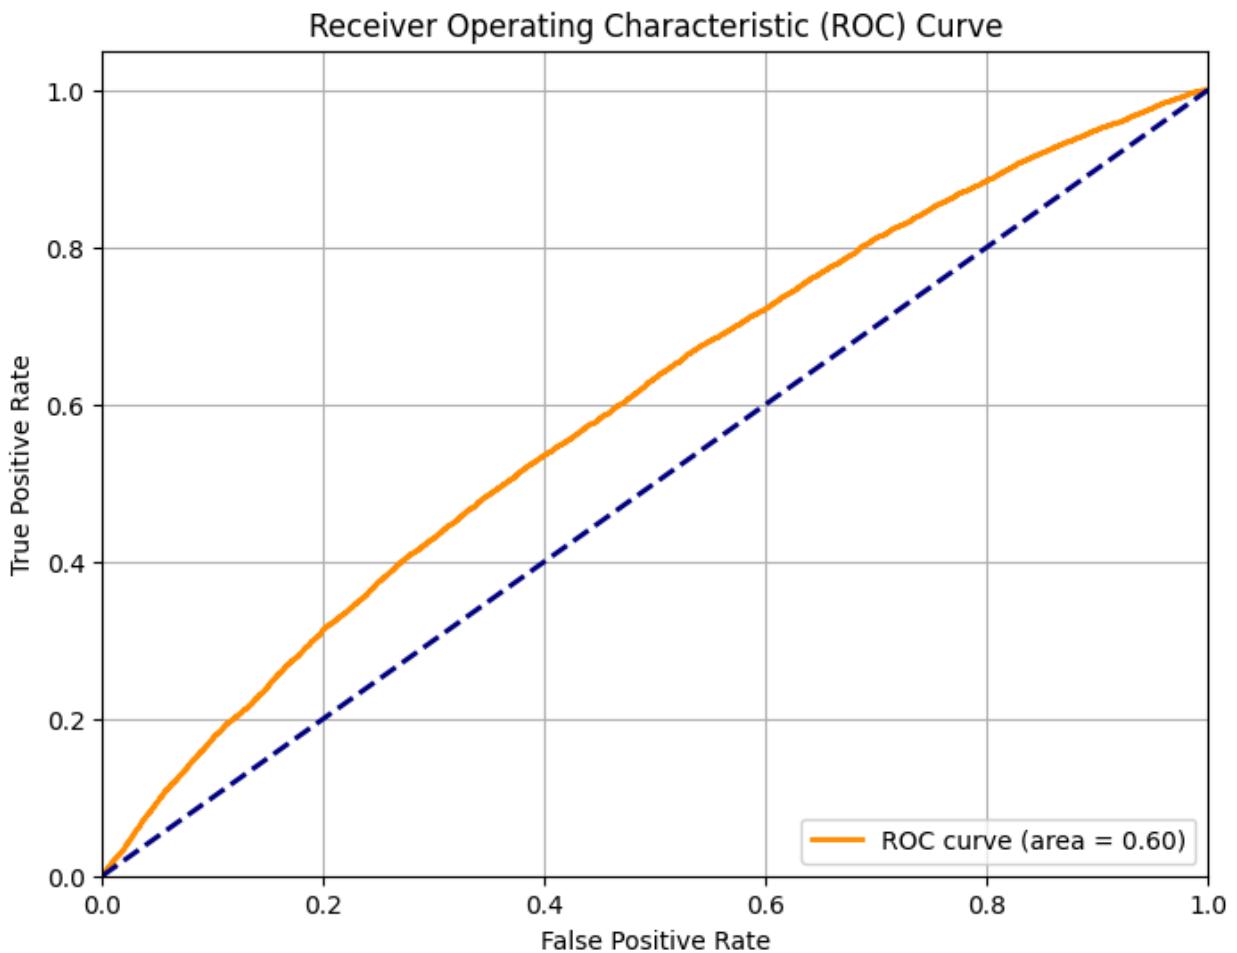
Model Evaluation on Test Set:

Accuracy: 0.9324
 Precision: 0.0000
 Recall: 0.0000
 F1-Score: 0.0000
 ROC AUC: 0.5961

```

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to
no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```



```
In [148]:
```

```
# Make predictions on the test set
y_pred = logistic_model.predict(X_test)
y_pred_proba = logistic_model.predict_proba(X_test)[:, 1]

# Calculate evaluation metrics
accuracy = accuracy_score(y_test, y_pred)
precision = precision_score(y_test, y_pred)
recall = recall_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred)

# Calculate ROC curve and AUC
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba)
roc_auc = auc(fpr, tpr)

print("Model Evaluation on Test Set:")
print(f"Accuracy: {accuracy:.4f}")
print(f"Precision: {precision:.4f}")
print(f"Recall: {recall:.4f}")
print(f"F1-Score: {f1:.4f}")
print(f"ROC AUC: {roc_auc:.4f}")
```

```

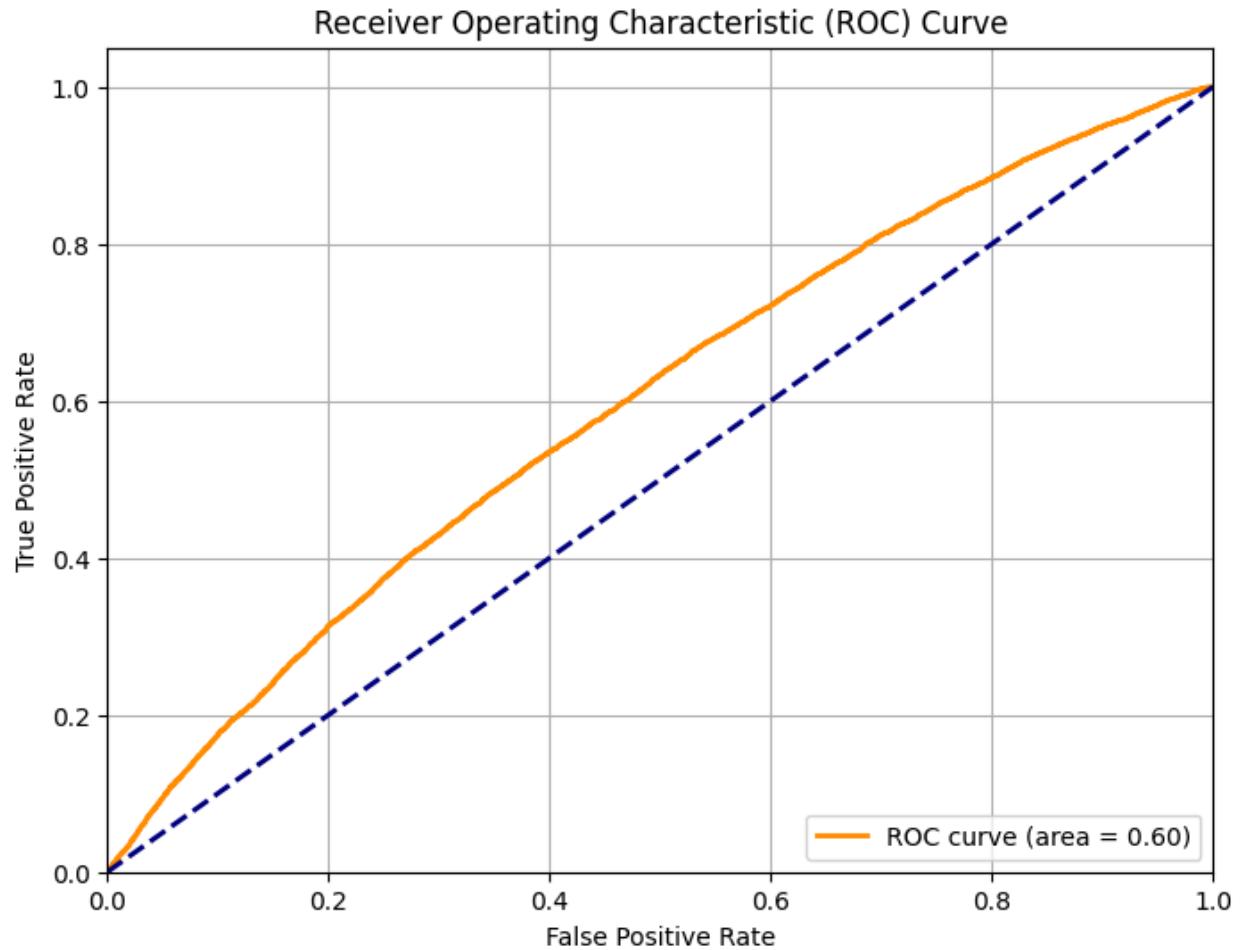
# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```

Model Evaluation on Test Set:

Accuracy: 0.9324
 Precision: 0.0000
 Recall: 0.0000
 F1-Score: 0.0000
 ROC AUC: 0.5961

/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156:
 5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to
 no predicted samples. Use `zero_division` parameter to control this behavior.
 _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))



Now that the Logistic Regression model is trained, I'll evaluate its performance. I'll make predictions on the test data (X_{test}) and then calculate standard classification metrics like accuracy, precision, recall, and F1-score. Additionally, I'll generate and display the ROC curve along with its Area Under the Curve (AUC) score to provide a comprehensive view of the model's effectiveness.

The Logistic Regression model has been evaluated, and here are the results:

Model Evaluation on Test Set:

- **Accuracy:** 0.9324
- **Precision:** 0.0000
- **Recall:** 0.0000
- **F1-Score:** 0.0000
- **ROC AUC:** 0.5961

Summary of Findings:

- High Accuracy, Low Other Metrics: The accuracy of 93.24% seems high, but the precision, recall, and F1-score are all 0.00. This indicates a significant issue: the model is likely predicting the negative class ('no click') for all instances. Since the dataset is highly imbalanced (most instances are 'no click'), simply predicting 'no click' for everything can yield high accuracy while failing to identify any actual clicks.
- UndefinedMetricWarning: The warning "Precision is ill-defined and being set to 0.0 due to no predicted samples" confirms that the model made no positive predictions ('no clicks'), leading to zero values for precision, recall, and F1-score.
- ROC AUC: The ROC AUC of 0.5961 suggests that the model's ability to distinguish between positive and negative classes is only slightly better than random guessing (an AUC of 0.5 would be random). This metric is more reliable than accuracy in imbalanced datasets as it considers the trade-off between true positive rate and false positive rate across various thresholds.

Next Steps: Given these results, it's clear that the model is struggling to correctly identify ad clicks. This is a common problem with imbalanced datasets. To improve the model's performance, I recommend addressing the class imbalance. This can be done using techniques such as:

- **Resampling methods:** Oversampling the minority class ('clicks') or undersampling the majority class ('no clicks').
- **Using different evaluation metrics:** While ROC AUC is already a

good start, metrics like Precision-Recall curves can offer more insight into the performance on the minority class.

- **Adjusting classification thresholds:** Instead of using the default 0.5 threshold, we could explore different thresholds based on the `y_pred_proba` values to balance precision and recall.
- **Exploring different algorithms:** Some algorithms are inherently more robust to class imbalance or have built-in mechanisms to handle it.

I will now proceed with addressing the class imbalance using a technique called SMOTE (Synthetic Minority Over-sampling Technique) to create synthetic samples of the minority class, which often helps models learn to predict it better.

SMOTE

```
In [210...]: # Apply SMOTE to oversample the minority class
```

```
print("\nClass distribution before SMOTE:")
print(y_train.value_counts())

sm = SMOTE(random_state=42, sampling_strategy=0.15)
X_resampled, y_resampled = sm.fit_resample(X_train, y_train)

print("\nClass distribution after SMOTE:")
print(y_resampled.value_counts())
```

Class distribution before SMOTE:

```
is_click
0    345567
1    25065
Name: count, dtype: int64
```

Class distribution after SMOTE:

```
is_click
0    345567
1    51835
Name: count, dtype: int64
```

```
In [150...]: # Retrain a Logistic Regression model on the resampled data
```

```
logistic_model_resampled = LogisticRegression(solver='liblinear', random_state=42)
logistic_model_resampled.fit(X_resampled, y_resampled)

print("Logistic Regression model retrained successfully on resampled data.")
```

Logistic Regression model retrained successfully on resampled data.

```
In [151...]: # Make predictions on the test set using the retrained model
```

```
y_pred_resampled = logistic_model_resampled.predict(X_test)
y_pred_proba_resampled = logistic_model_resampled.predict_proba(X_test)[:, 1]
```

```

# Calculate evaluation metrics
accuracy_resampled = accuracy_score(y_test, y_pred_resampled)
precision_resampled = precision_score(y_test, y_pred_resampled)
recall_resampled = recall_score(y_test, y_pred_resampled)
f1_resampled = f1_score(y_test, y_pred_resampled)

# Calculate ROC curve and AUC
fpr_resampled, tpr_resampled, thresholds_resampled = roc_curve(y_test, y_pred_resampled)
roc_auc_resampled = auc(fpr_resampled, tpr_resampled)

print("Model Evaluation on Test Set (Retrained with SMOTE):")
print(f"Accuracy: {accuracy_resampled:.4f}")
print(f"Precision: {precision_resampled:.4f}")
print(f"Recall: {recall_resampled:.4f}")
print(f"F1-Score: {f1_resampled:.4f}")
print(f"ROC AUC: {roc_auc_resampled:.4f}")

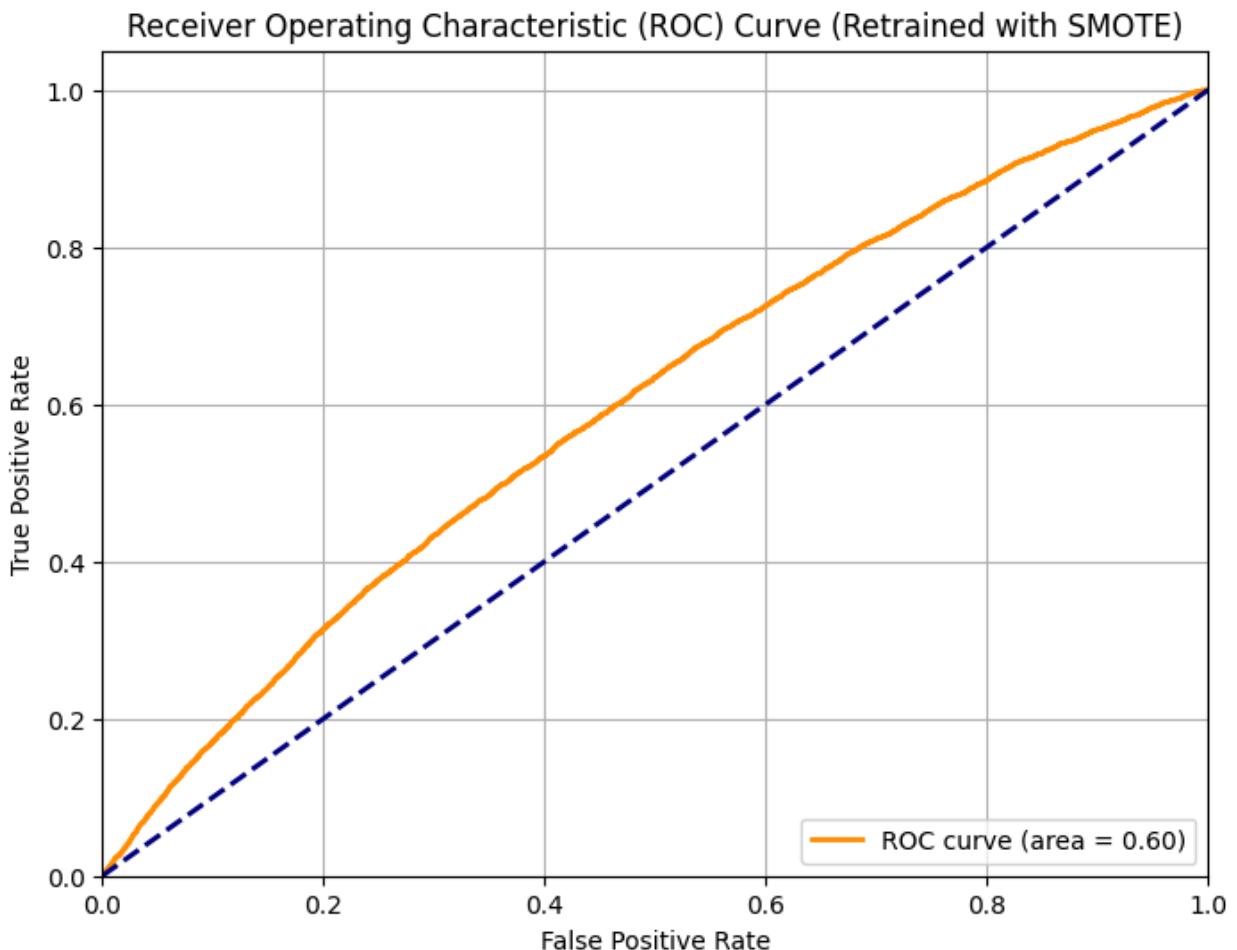
# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_resampled, tpr_resampled, color='darkorange', lw=2, label=f'ROC curve')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Retrained with SMOTE)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```

Model Evaluation on Test Set (Retrained with SMOTE):

Accuracy: 0.9324
 Precision: 0.0000
 Recall: 0.0000
 F1-Score: 0.0000
 ROC AUC: 0.5970

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to
no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```



Train Decision Tree Model

Train a Decision Tree Classifier on the SMOTE-resampled training data (`X_resampled`, `y_resampled`), evaluate its performance on the original test set (`X_test`, `y_test`) using accuracy, precision, recall, F1-score, and ROC AUC, and then plot the ROC curve.

```
In [152]:  
from sklearn.tree import DecisionTreeClassifier  
  
# 1. Initialize a Decision Tree classifier  
decision_tree_model = DecisionTreeClassifier(random_state=42)  
  
# 2. Train the Decision Tree model on the resampled training data  
# decision_tree_model.fit(X_resampled, y_resampled)  
decision_tree_model.fit(X_resampled, y_resampled)  
  
print("Decision Tree model trained successfully on resampled data.")
```

```

# 3. Make predictions on the test set
y_pred_dt = decision_tree_model.predict(X_test)

# 4. Get the predicted probabilities for the positive class
y_pred_proba_dt = decision_tree_model.predict_proba(X_test)[:, 1]

# 5. Calculate evaluation metrics
accuracy_dt = accuracy_score(y_test, y_pred_dt)
precision_dt = precision_score(y_test, y_pred_dt)
recall_dt = recall_score(y_test, y_pred_dt)
f1_dt = f1_score(y_test, y_pred_dt)

# 6. Calculate ROC curve and AUC
fpr_dt, tpr_dt, thresholds_dt = roc_curve(y_test, y_pred_proba_dt)
roc_auc_dt = auc(fpr_dt, tpr_dt)

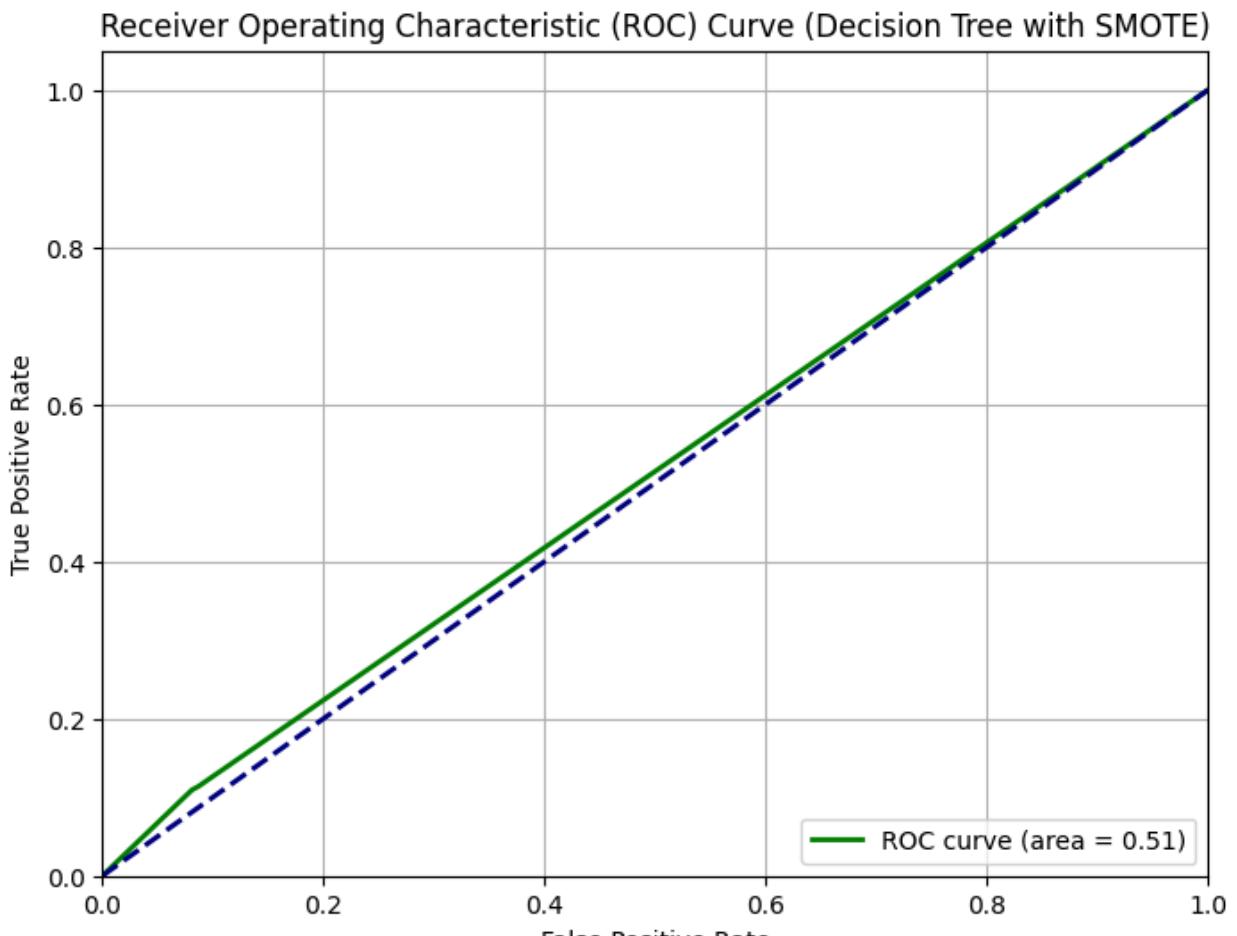
# 7. Print all calculated metrics
print("\nModel Evaluation on Test Set (Decision Tree with SMOTE):")
print(f"Accuracy: {accuracy_dt:.4f}")
print(f"Precision: {precision_dt:.4f}")
print(f"Recall: {recall_dt:.4f}")
print(f"F1-Score: {f1_dt:.4f}")
print(f"ROC AUC: {roc_auc_dt:.4f}")

# 8. Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_dt, tpr_dt, color='green', lw=2, label=f'ROC curve (area = {roc_auc_dt:.4f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Decision Tree with SMOTE)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```

Decision Tree model trained successfully on resampled data.

Model Evaluation on Test Set (Decision Tree with SMOTE):
Accuracy: 0.8639
Precision: 0.0890
Recall: 0.1096
F1-Score: 0.0982
ROC AUC: 0.5136



Train Random Forest Model

Initialize a Random Forest classifier, train it on the SMOTE-resampled training data (`X_resampled`, `y_resampled`), make predictions on the original test set (`X_test`), and evaluate its performance using accuracy, precision, recall, F1-score, and ROC AUC. Finally, plot the ROC curve for this model.

```
In [153]:  
from sklearn.ensemble import RandomForestClassifier  
  
# 1. Initialize a Random Forest classifier  
random_forest_model = RandomForestClassifier(random_state=42)  
  
# 2. Train the Random Forest model on the resampled training data  
random_forest_model.fit(X_resampled, y_resampled)  
  
print("Random Forest model trained successfully on resampled data.")  
  
# 3. Make predictions on the test set
```

```

y_pred_rf = random_forest_model.predict(X_test)

# 4. Get the predicted probabilities for the positive class
y_pred_proba_rf = random_forest_model.predict_proba(X_test)[:, 1]

# 5. Calculate evaluation metrics
accuracy_rf = accuracy_score(y_test, y_pred_rf)
precision_rf = precision_score(y_test, y_pred_rf)
recall_rf = recall_score(y_test, y_pred_rf)
f1_rf = f1_score(y_test, y_pred_rf)

# 6. Calculate ROC curve and AUC
fpr_rf, tpr_rf, thresholds_rf = roc_curve(y_test, y_pred_proba_rf)
roc_auc_rf = auc(fpr_rf, tpr_rf)

# 7. Print all calculated metrics
print("\nModel Evaluation on Test Set (Random Forest with SMOTE):")
print(f"Accuracy: {accuracy_rf:.4f}")
print(f"Precision: {precision_rf:.4f}")
print(f"Recall: {recall_rf:.4f}")
print(f"F1-Score: {f1_rf:.4f}")
print(f"ROC AUC: {roc_auc_rf:.4f}")

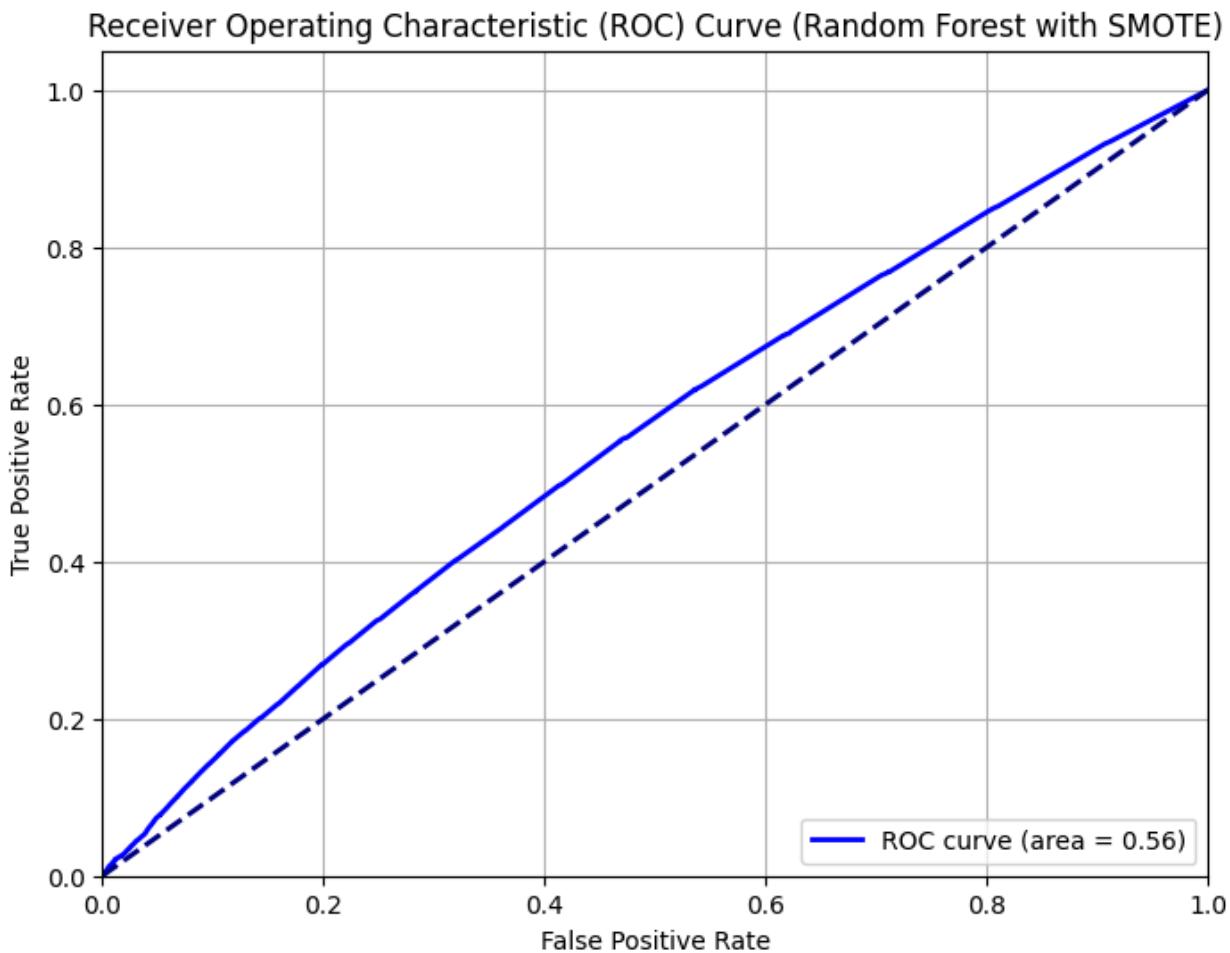
# 8. Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_rf, tpr_rf, color='blue', lw=2, label=f'ROC curve (area = {roc_auc_rf:.4f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Random Forest with SMOTE)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```

Random Forest model trained successfully on resampled data.

Model Evaluation on Test Set (Random Forest with SMOTE):

Accuracy: 0.9242
 Precision: 0.1121
 Recall: 0.0176
 F1-Score: 0.0304
 ROC AUC: 0.5571



Train Gradient Boosting Model

Initialize a Gradient Boosting classifier, train it on the SMOTE-resampled training data (`X_resampled`, `y_resampled`), make predictions on the original test set (`X_test`), and evaluate its performance using accuracy, precision, recall, F1-score, and ROC AUC. Finally, plot the ROC curve for this model.

```
In [154]: from sklearn.ensemble import GradientBoostingClassifier
# 1. Initialize a Gradient Boosting classifier
gradient_boosting_model = GradientBoostingClassifier(random_state=42)

# 2. Train the Gradient Boosting model on the resampled training data
gradient_boosting_model.fit(X_resampled, y_resampled)

print("Gradient Boosting model trained successfully on resampled data.")

# 3. Make predictions on the test set
y_pred_gb = gradient_boosting_model.predict(X_test)

# 4. Get the predicted probabilities for the positive class
```

```

y_pred_proba_gb = gradient_boosting_model.predict_proba(X_test)[:, 1]

# 5. Calculate evaluation metrics
accuracy_gb = accuracy_score(y_test, y_pred_gb)
precision_gb = precision_score(y_test, y_pred_gb)
recall_gb = recall_score(y_test, y_pred_gb)
f1_gb = f1_score(y_test, y_pred_gb)

# 6. Calculate ROC curve and AUC
fpr_gb, tpr_gb, thresholds_gb = roc_curve(y_test, y_pred_proba_gb)
roc_auc_gb = auc(fpr_gb, tpr_gb)

# 7. Print all calculated metrics
print("\nModel Evaluation on Test Set (Gradient Boosting with SMOTE):")
print(f"Accuracy: {accuracy_gb:.4f}")
print(f"Precision: {precision_gb:.4f}")
print(f"Recall: {recall_gb:.4f}")
print(f"F1-Score: {f1_gb:.4f}")
print(f"ROC AUC: {roc_auc_gb:.4f}")

# 8. Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_gb, tpr_gb, color='purple', lw=2, label=f'ROC curve (area = {roc_auc_gb:.4f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Gradient Boosting with SMOTE)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

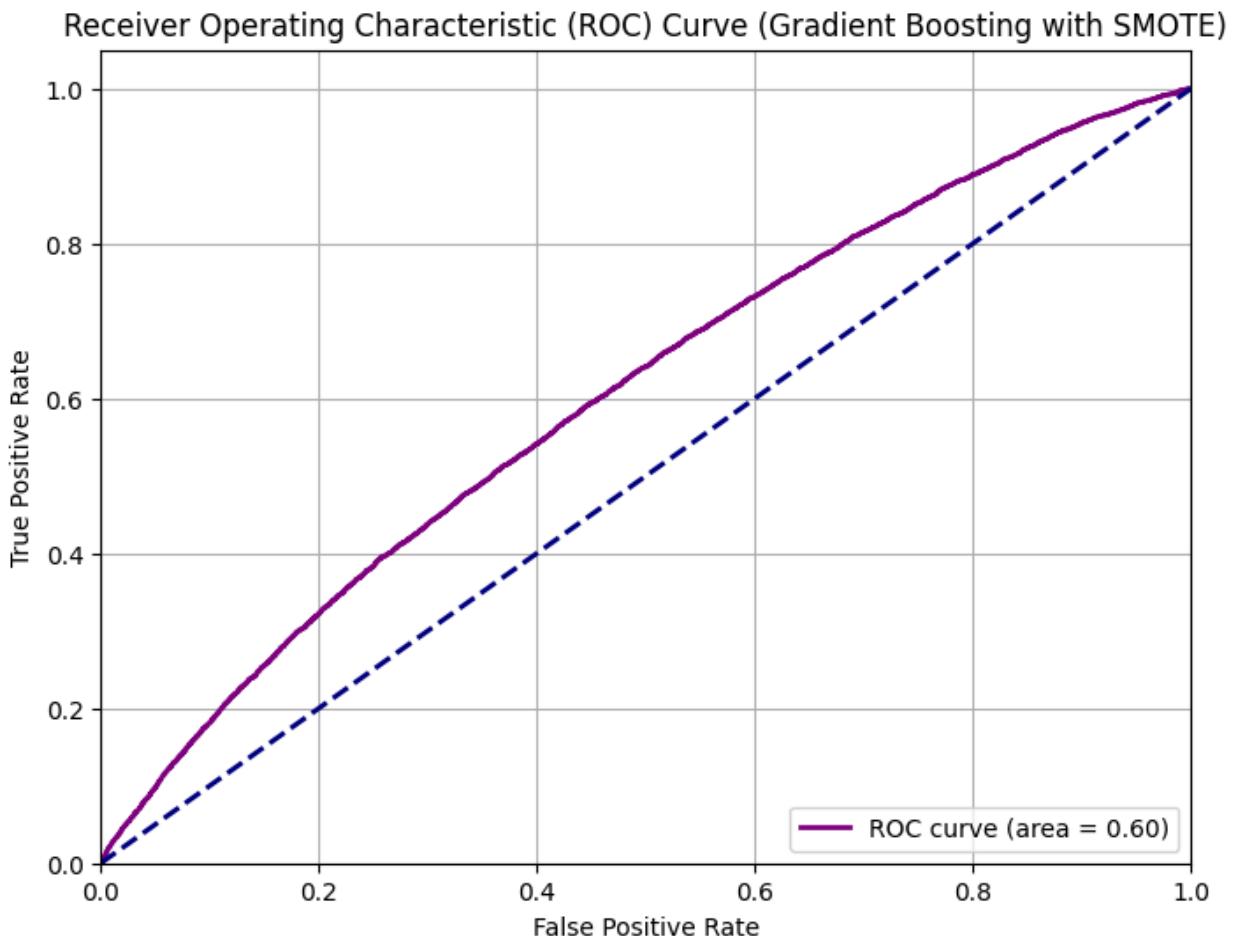
```

Gradient Boosting model trained successfully on resampled data.

Model Evaluation on Test Set (Gradient Boosting with SMOTE):

Accuracy: 0.9324
 Precision: 0.0000
 Recall: 0.0000
 F1-Score: 0.0000
 ROC AUC: 0.6034

```
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to
no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```



Train XGBoost Model

```
In [155...]:  
import sys  
!{sys.executable} -m pip install xgboost  
from xgboost import XGBClassifier  
from sklearn.model_selection import RandomizedSearchCV  
  
param_grid_xgb = {  
    'n_estimators': [100, 200, 300, 400],  
    'learning_rate': [0.01, 0.05, 0.1, 0.15, 0.2],  
    'max_depth': [3, 4, 5, 6, 7],  
    'subsample': [0.6, 0.7, 0.8, 0.9, 1.0],  
    'colsample_bytree': [0.6, 0.7, 0.8, 0.9, 1.0],  
    'gamma': [0, 0.1, 0.2, 0.3]  
}  
  
print("Defined Parameter Grid for XGBoost Classifier:")  
print(param_grid_xgb)
```

```
Requirement already satisfied: xgboost in /usr/local/lib/python3.12/dist-packages (3.1.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (from xgboost) (2.0.2)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.12/dist-packages (from xgboost) (2.28.9)
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (from xgboost) (1.16.3)
Defined Parameter Grid for XGBoost Classifier:
{'n_estimators': [100, 200, 300, 400], 'learning_rate': [0.01, 0.05, 0.1, 0.15, 0.2], 'max_depth': [3, 4, 5, 6, 7], 'subsample': [0.6, 0.7, 0.8, 0.9, 1.0], 'colsample_bytree': [0.6, 0.7, 0.8, 0.9, 1.0], 'gamma': [0, 0.1, 0.2, 0.3]}
```

```
In [156...]: import xgboost as xgb

# Initialize XGBClassifier
xgb_model = xgb.XGBClassifier(objective='binary:logistic', eval_metric='logloss')

random_search_xgb = RandomizedSearchCV(
    estimator=xgb_model,
    param_distributions=param_grid_xgb,
    n_iter=20, # Number of parameter settings that are sampled, increased from 10 in the previous example
    scoring={'roc_auc': 'roc_auc', 'f1': 'f1'},
    refit='roc_auc', # Refit the estimator with the best_params_ found on the cv=3, # 3-fold cross-validation
    verbose=2,
    random_state=42,
    n_jobs=-1 # Use all available cores
)

# Fit RandomizedSearchCV on the resampled training data
print("\nPerforming Randomized Search for XGBoost...")
random_search_xgb.fit(X_resampled, y_resampled)

print("\nRandomized Search for XGBoost completed.")
print(f"Best parameters: {random_search_xgb.best_params_}")
print(f"Best ROC AUC score: {random_search_xgb.best_score_.4f}")
```

```
Performing Randomized Search for XGBoost...
Fitting 3 folds for each of 20 candidates, totalling 60 fits
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning:
[10:57:13] WARNING: /workspace/src/learner.cc:790:
Parameters: { "use_label_encoder" } are not used.

bst.update(dtrain, iteration=i, fobj=obj)
Randomized Search for XGBoost completed.
Best parameters: {'subsample': 0.7, 'n_estimators': 300, 'max_depth': 7, 'learning_rate': 0.05, 'gamma': 0.2, 'colsample_bytree': 0.8}
Best ROC AUC score: 0.8059
```

```
In [157...]: import xgboost as xgb
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
import matplotlib.pyplot as plt
```

```

import numpy as np

# Train XGBoost Model with Best Parameters
best_xgb_model = random_search_xgb.best_estimator_ # This is the best model retrieved from Randomized Search.

# Evaluate Tuned XGBoost Model
y_pred_xgb = best_xgb_model.predict(X_test)
y_pred_proba_xgb = best_xgb_model.predict_proba(X_test)[:, 1]

accuracy_xgb = accuracy_score(y_test, y_pred_xgb)
precision_xgb = precision_score(y_test, y_pred_xgb)
recall_xgb = recall_score(y_test, y_pred_xgb)
f1_xgb = f1_score(y_test, y_pred_xgb)

fpr_xgb, tpr_xgb, thresholds_xgb = roc_curve(y_test, y_pred_proba_xgb)
roc_auc_xgb = auc(fpr_xgb, tpr_xgb)

print("\nModel Evaluation on Test Set (Tuned XGBoost with RandomUnderSampler):")
print(f"Accuracy: {accuracy_xgb:.4f}")
print(f"Precision: {precision_xgb:.4f}")
print(f"Recall: {recall_xgb:.4f}")
print(f"F1-Score: {f1_xgb:.4f}")
print(f"ROC AUC: {roc_auc_xgb:.4f}")

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_xgb, tpr_xgb, color='red', lw=2, label=f'ROC curve (area = {roc_auc_xgb:.4f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Tuned XGBoost with RandomUnderSampler)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```

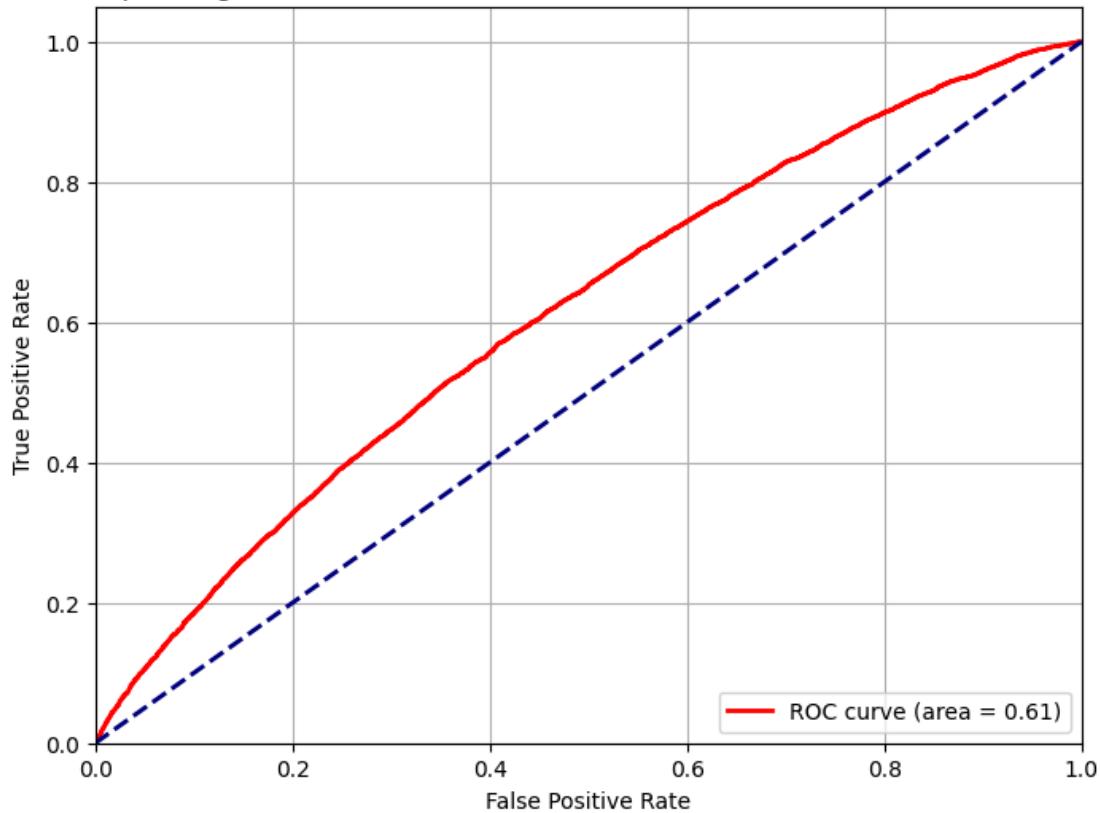
Best XGBoost model retrieved from Randomized Search.

```

Model Evaluation on Test Set (Tuned XGBoost with RandomUnderSampler):
Accuracy: 0.9324
Precision: 0.0000
Recall: 0.0000
F1-Score: 0.0000
ROC AUC: 0.6121
/usr/local/lib/python3.12/dist-packages/sklearn/metrics/_classification.py:156
5: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 due to
no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

Receiver Operating Characteristic (ROC) Curve (Tuned XGBoost with RandomUnderSampler)



Adjust Classification Threshold for Tuned XGBoost Model

Perform a sensitivity analysis on the classification threshold for the tuned XGBoost model to find a better balance between precision and recall, and evaluate its performance with the adjusted threshold.

To perform a sensitivity analysis, I will first define a range of classification thresholds and then iterate through these thresholds to calculate predictions and evaluate the model's performance using accuracy, precision, recall, and F1-score for each threshold. This data will then be used for plotting.

```
In [158]:  
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score  
import matplotlib.pyplot as plt  
import numpy as np  
  
# Define a range of classification thresholds  
thresholds = np.arange(0.1, 0.35, 0.01)  
  
# Store metrics for each threshold  
accuracy_scores = []  
precision_scores = []  
recall_scores = []
```

```

f1_scores = []

# Assuming y_pred_proba_xgb and y_test are already available from previous step

for threshold in thresholds:
    # Calculate predictions for the current threshold
    y_pred_threshold = (y_pred_proba_xgb >= threshold).astype(int)

    # Calculate evaluation metrics
    accuracy_scores.append(accuracy_score(y_test, y_pred_threshold))
    precision_scores.append(precision_score(y_test, y_pred_threshold, zero_division=0))
    recall_scores.append(recall_score(y_test, y_pred_threshold, zero_division=0))
    f1_scores.append(f1_score(y_test, y_pred_threshold, zero_division=0))

print("Calculated evaluation metrics for various thresholds.")

```

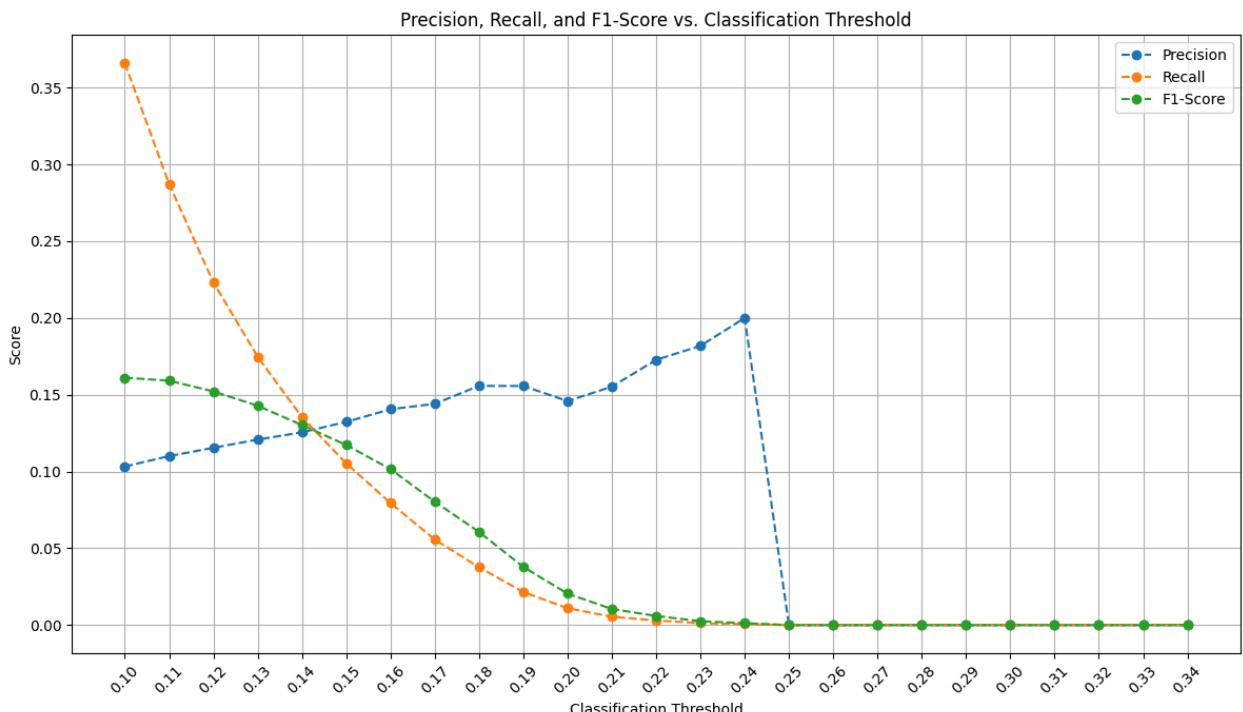
Calculated evaluation metrics for various thresholds.

```

In [159]: plt.figure(figsize=(12, 7))
plt.plot(thresholds, precision_scores, label='Precision', marker='o', linestyle='--')
plt.plot(thresholds, recall_scores, label='Recall', marker='o', linestyle='--')
plt.plot(thresholds, f1_scores, label='F1-Score', marker='o', linestyle='--')
plt.xlabel('Classification Threshold')
plt.ylabel('Score')
plt.title('Precision, Recall, and F1-Score vs. Classification Threshold')
plt.legend()
plt.grid(True)
plt.xticks(thresholds, rotation=45)
plt.tight_layout()
plt.show()

print("Plotted precision, recall, and F1-score against classification threshold")

```



Plotted precision, recall, and F1-score against classification thresholds.

In [160...]

```
optimal_threshold_idx = np.argmax(precision_scores)
optimal_threshold = thresholds[optimal_threshold_idx]

# Calculate predictions with the optimal threshold
y_pred_optimal = (y_pred_proba_xgb >= optimal_threshold).astype(int)

# Calculate evaluation metrics for the optimal threshold
accuracy_optimal = accuracy_score(y_test, y_pred_optimal)
precision_optimal = precision_score(y_test, y_pred_optimal, zero_division=0)
recall_optimal = recall_score(y_test, y_pred_optimal, zero_division=0)
f1_optimal = f1_score(y_test, y_pred_optimal, zero_division=0)

# ROC AUC remains the same as it's threshold-independent

print(f"\nOptimal Classification Threshold (based on F1-Score): {optimal_thres
print("Model Evaluation on Test Set (Tuned XGBoost with Optimal Threshold):")
print(f"Accuracy: {accuracy_optimal:.4f}")
print(f"Precision: {precision_optimal:.4f}")
print(f"Recall: {recall_optimal:.4f}")
print(f"F1-Score: {f1_optimal:.4f}")
print(f"ROC AUC: {roc_auc_xgb:.4f}")
```

```
Optimal Classification Threshold (based on F1-Score): 0.24
Model Evaluation on Test Set (Tuned XGBoost with Optimal Threshold):
Accuracy: 0.9322
Precision: 0.2000
Recall: 0.0006
F1-Score: 0.0013
ROC AUC: 0.6121
```

Soft voting ensemble model

Obtain the predicted probabilities for the positive class (`is_click=1`) from the Logistic Regression, Decision Tree, Random Forest, and Tuned XGBoost models on the `X_test` dataset.

In [161...]

```
proba_lr = logistic_model_resampled.predict_proba(X_test)[:, 1]
proba_dt = decision_tree_model.predict_proba(X_test)[:, 1]
proba_rf = random_forest_model.predict_proba(X_test)[:, 1]
proba_xgb = best_xgb_model.predict_proba(X_test)[:, 1]

print("Predicted probabilities obtained for Logistic Regression, Decision Tree,
```

Predicted probabilities obtained for Logistic Regression, Decision Tree, Random Forest, and Tuned XGBoost models.

In [162...]

```
ensemble_proba = (proba_lr + proba_dt + proba_rf + proba_xgb) / 4
```

```
print("Soft voting ensemble probabilities calculated.")
```

Soft voting ensemble probabilities calculated.

```
In [163]: optimal_threshold = 0.1 # Reusing the optimal threshold found for XGBoost

# Get binary predictions using the optimal threshold
ensemble_predictions = (ensemble_proba >= optimal_threshold).astype(int)

# Calculate evaluation metrics
accuracy_ensemble = accuracy_score(y_test, ensemble_predictions)
precision_ensemble = precision_score(y_test, ensemble_predictions, zero_division=0)
recall_ensemble = recall_score(y_test, ensemble_predictions, zero_division=0)
f1_ensemble = f1_score(y_test, ensemble_predictions, zero_division=0)

# Calculate ROC curve and AUC
fpr_ensemble, tpr_ensemble, _ = roc_curve(y_test, ensemble_proba)
roc_auc_ensemble = auc(fpr_ensemble, tpr_ensemble)

print("Ensemble Model Evaluation with Optimal Threshold (0.10):")
print(f"Accuracy: {accuracy_ensemble:.4f}")
print(f"Precision: {precision_ensemble:.4f}")
print(f"Recall: {recall_ensemble:.4f}")
print(f"F1-Score: {f1_ensemble:.4f}")
print(f"ROC AUC: {roc_auc_ensemble:.4f}")

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_ensemble, tpr_ensemble, color='green', lw=2, label=f'Ensemble ROC')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Ensemble Model)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()
```

Ensemble Model Evaluation with Optimal Threshold (0.10):

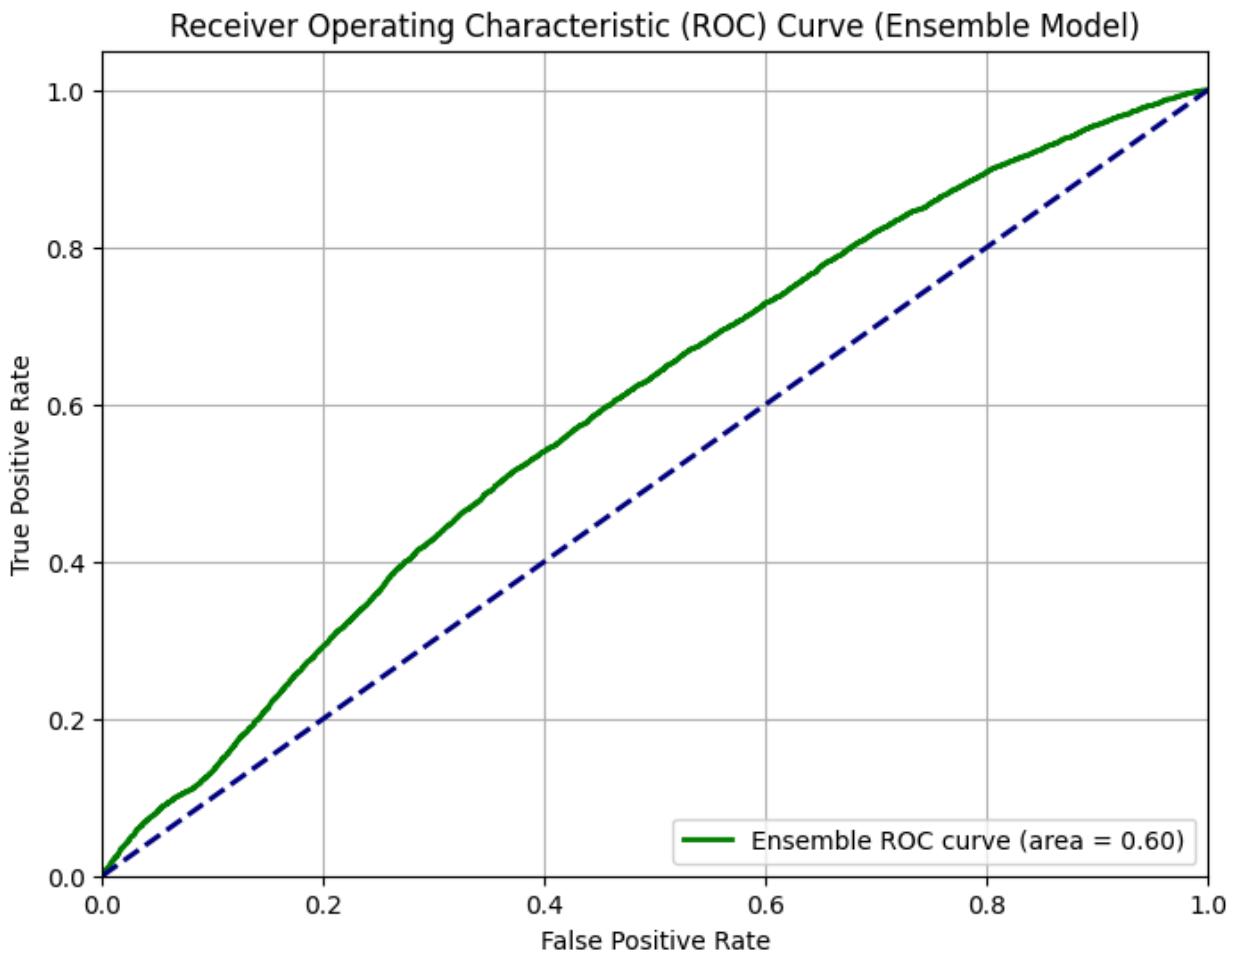
Accuracy: 0.7547

Precision: 0.0958

Recall: 0.3114

F1-Score: 0.1465

ROC AUC: 0.5951



Weighted Averaging Ensemble

Create a weighted averaging ensemble model by assigning custom weights to the predicted probabilities from the Logistic Regression, Decision Tree, Random Forest, and Tuned XGBoost models.

```
In [164...]: weights = {
    'Logistic_Regression': 0.1,
    'Decision_Tree': 0.1,
    'Random_Forest': 0.3,
    'XGBoost': 0.5
}

# Ensure weights sum to 1
sum_of_weights = sum(weights.values())
print(f"Sum of assigned weights: {sum_of_weights:.2f}")

# Multiply probabilities by weights and sum them
weighted_ensemble_proba = (
    proba_lr * weights['Logistic_Regression'] +
    proba_dt * weights['Decision_Tree'] +
```

```

    proba_rf * weights['Random_Forest'] +
    proba_xgb * weights['XGBoost']
)

print("Weighted averaging ensemble probabilities calculated.")

```

Sum of assigned weights: 1.00
 Weighted averaging ensemble probabilities calculated.

```

In [165]: optimal_threshold = 0.31 # Reusing the optimal threshold found for XGBoost

# Get binary predictions using the optimal threshold
weighted_ensemble_predictions = (weighted_ensemble_proba >= optimal_threshold)

# Calculate evaluation metrics
accuracy_weighted_ensemble = accuracy_score(y_test, weighted_ensemble_predictions)
precision_weighted_ensemble = precision_score(y_test, weighted_ensemble_predictions)
recall_weighted_ensemble = recall_score(y_test, weighted_ensemble_predictions)
f1_weighted_ensemble = f1_score(y_test, weighted_ensemble_predictions, zero_division=1)

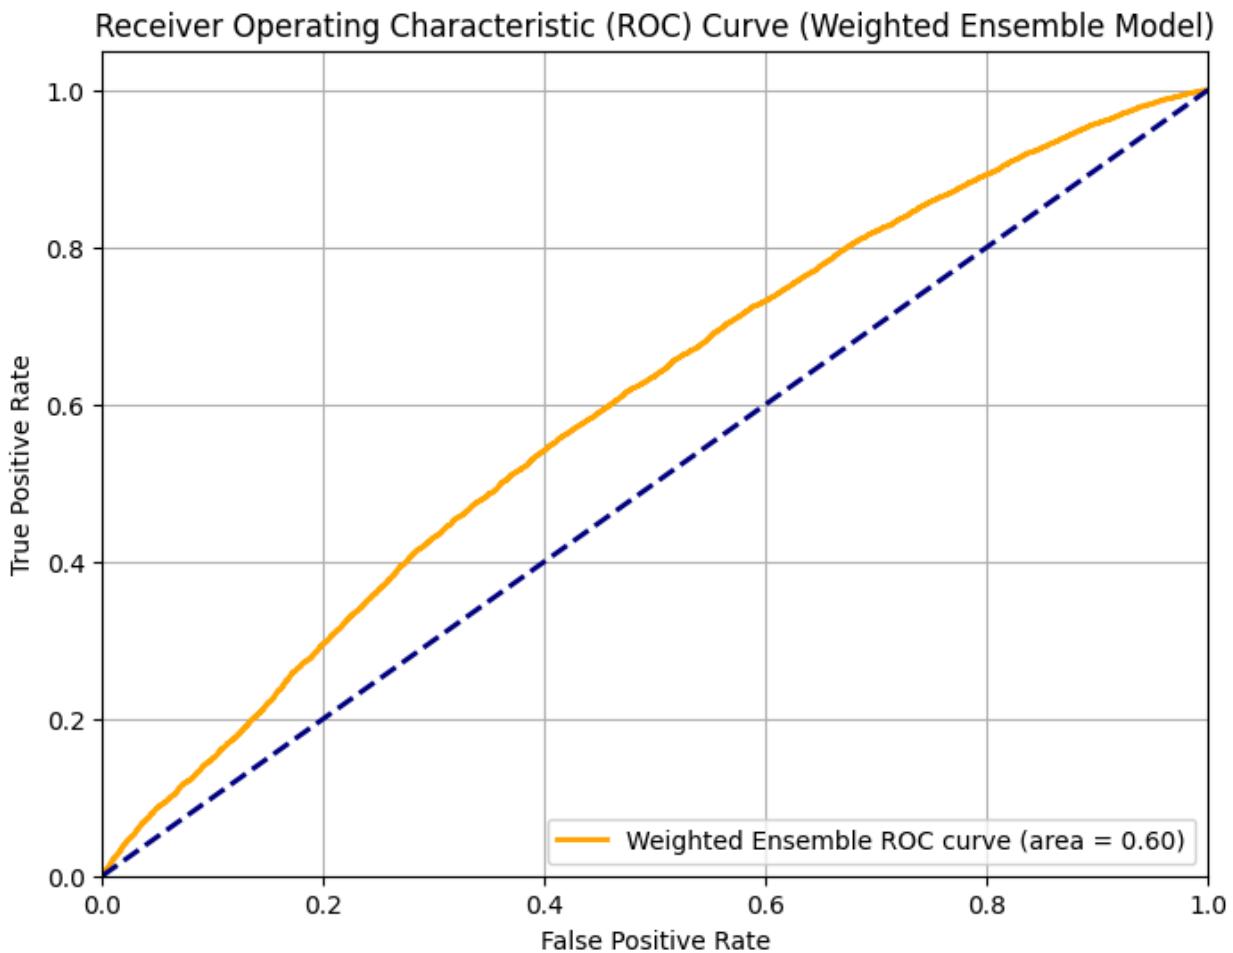
# Calculate ROC curve and AUC
fpr_weighted_ensemble, tpr_weighted_ensemble, _ = roc_curve(y_test, weighted_ensemble_proba)
roc_auc_weighted_ensemble = auc(fpr_weighted_ensemble, tpr_weighted_ensemble)

print("Weighted Ensemble Model Evaluation with Optimal Threshold (0.31):")
print(f"Accuracy: {accuracy_weighted_ensemble:.4f}")
print(f"Precision: {precision_weighted_ensemble:.4f}")
print(f"Recall: {recall_weighted_ensemble:.4f}")
print(f"F1-Score: {f1_weighted_ensemble:.4f}")
print(f"ROC AUC: {roc_auc_weighted_ensemble:.4f}")

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_weighted_ensemble, tpr_weighted_ensemble, color='orange', lw=2, label='Weighted Ensemble')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Weighted Ensemble Model)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```

Weighted Ensemble Model Evaluation with Optimal Threshold (0.31):
 Accuracy: 0.9256
 Precision: 0.1300
 Recall: 0.0176
 F1-Score: 0.0309
 ROC AUC: 0.5972



Calculate Weighted Ensemble Probabilities on Train Data

Generate the predicted probabilities for the positive class from the individual models (Logistic Regression, Decision Tree, Random Forest, Tuned XGBoost) using the `X_train` dataset. Then, combine these probabilities using the predefined weights to create `weighted_ensemble_proba_train`.

```
In [166]: proba_lr_train = logistic_model_resampled.predict_proba(X_resampled)[:, 1]
proba_dt_train = decision_tree_model.predict_proba(X_resampled)[:, 1]
proba_rf_train = random_forest_model.predict_proba(X_resampled)[:, 1]
proba_xgb_train = best_xgb_model.predict_proba(X_resampled)[:, 1]

weighted_ensemble_proba_train = (
    proba_lr_train * weights['Logistic_Regression'] +
    proba_dt_train * weights['Decision_Tree'] +
    proba_rf_train * weights['Random_Forest'] +
    proba_xgb_train * weights['XGBoost']
)

print("Weighted averaging ensemble probabilities for training data calculated.")
```

Weighted averaging ensemble probabilities for training data calculated.

In [167...]

```
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
import numpy as np

# Define the range of classification thresholds from 0.11 to 0.49 with a step
thresholds_train_analysis = np.arange(0.11, 0.50, 0.01)

# Store metrics for each threshold
accuracy_scores_train = []
precision_scores_train = []
recall_scores_train = []
f1_scores_train = []

print("Performing sensitivity analysis on training data probabilities...")

for threshold in thresholds_train_analysis:
    # Calculate predictions for the current threshold
    y_pred_threshold_train = (weighted_ensemble_proba_train >= threshold).astype(int)

    # Calculate evaluation metrics. Use zero_division=0 to handle cases with no
    accuracy_scores_train.append(accuracy_score(y_resampled, y_pred_threshold_train))
    precision_scores_train.append(precision_score(y_resampled, y_pred_threshold_train))
    recall_scores_train.append(recall_score(y_resampled, y_pred_threshold_train))
    f1_scores_train.append(f1_score(y_resampled, y_pred_threshold_train, zero_division=0))

print("Evaluation metrics calculated for various thresholds on training data.")
```

Performing sensitivity analysis on training data probabilities...

Evaluation metrics calculated for various thresholds on training data.

In [168...]

```
optimal_threshold_idx_train = np.argmax(f1_scores_train)
optimal_threshold_train = thresholds_train_analysis[optimal_threshold_idx_train]

# Get metrics for the optimal threshold on training data
accuracy_optimal_train = accuracy_scores_train[optimal_threshold_idx_train]
precision_optimal_train = precision_scores_train[optimal_threshold_idx_train]
recall_optimal_train = recall_scores_train[optimal_threshold_idx_train]
f1_optimal_train = f1_scores_train[optimal_threshold_idx_train]

print(f"\nOptimal Classification Threshold on Training Data (based on F1-Score: {f1_optimal_train:.4f})")
print("Weighted Ensemble Model Evaluation on Training Data (Optimal Threshold):")
print(f"Accuracy: {accuracy_optimal_train:.4f}")
print(f"Precision: {precision_optimal_train:.4f}")
print(f"Recall: {recall_optimal_train:.4f}")
print(f"F1-Score: {f1_optimal_train:.4f}")
```

Optimal Classification Threshold on Training Data (based on F1-Score): 0.23

Weighted Ensemble Model Evaluation on Training Data (Optimal Threshold):

Accuracy: 0.9968

Precision: 0.9861

Recall: 0.9890

F1-Score: 0.9876

```
In [169...]: current_optimal_threshold_test = optimal_threshold_train # Use the optimal threshold from training

# Get binary predictions using the new optimal threshold
weighted_ensemble_predictions_re_eval = (weighted_ensemble_proba >= current_optimal_threshold_test).astype(int)

# Calculate evaluation metrics
accuracy_weighted_ensemble_re_eval = accuracy_score(y_test, weighted_ensemble_predictions_re_eval)
precision_weighted_ensemble_re_eval = precision_score(y_test, weighted_ensemble_predictions_re_eval)
recall_weighted_ensemble_re_eval = recall_score(y_test, weighted_ensemble_predictions_re_eval)
f1_weighted_ensemble_re_eval = f1_score(y_test, weighted_ensemble_predictions_re_eval)

# ROC AUC remains the same as it's threshold-independent
roc_auc_weighted_ensemble_re_eval = roc_auc_weighted_ensemble

print(f"Weighted Ensemble Model Evaluation with Optimal Threshold (from Train Data):")
print(f"Accuracy: {accuracy_weighted_ensemble_re_eval:.4f}")
print(f"Precision: {precision_weighted_ensemble_re_eval:.4f}")
print(f"Recall: {recall_weighted_ensemble_re_eval:.4f}")
print(f"F1-Score: {f1_weighted_ensemble_re_eval:.4f}")
print(f"ROC AUC: {roc_auc_weighted_ensemble_re_eval:.4f}")
```

Weighted Ensemble Model Evaluation with Optimal Threshold (from Train Data):
0.23
Accuracy: 0.9090
Precision: 0.1175
Recall: 0.0531
F1-Score: 0.0732
ROC AUC: 0.5972

Weighted Ensemble Model Evaluation with Calibrated Probabilities

```
In [170...]: neg_samples = y_train.value_counts()[0]
pos_samples = y_train.value_counts()[1]
w = neg_samples / pos_samples

# Apply calibration formula
weighted_ensemble_proba_calibrated = weighted_ensemble_proba / (weighted_ensemble_proba + w)

optimal_threshold_from_train = optimal_threshold_train # This is 0.21

# Get binary predictions using the optimal threshold on calibrated probabilities
weighted_ensemble_predictions_calibrated = (weighted_ensemble_proba_calibrated >= optimal_threshold_from_train).astype(int)

# Calculate evaluation metrics
accuracy_calibrated = accuracy_score(y_test, weighted_ensemble_predictions_calibrated)
precision_calibrated = precision_score(y_test, weighted_ensemble_predictions_calibrated)
recall_calibrated = recall_score(y_test, weighted_ensemble_predictions_calibrated)
f1_calibrated = f1_score(y_test, weighted_ensemble_predictions_calibrated, zero_division=1)
```

```

# Calculate ROC curve and AUC for calibrated probabilities
fpr_calibrated, tpr_calibrated, _ = roc_curve(y_test, weighted_ensemble_proba)
roc_auc_calibrated = auc(fpr_calibrated, tpr_calibrated)

print(f"Negative samples in y_train: {neg_samples}")
print(f"Positive samples in y_train: {pos_samples}")
print(f"Calibration weight (w): {w:.2f}")

print(f"\nWeighted Ensemble Model Evaluation with Calibrated Probabilities (Optimal Threshold: 0.23):")
print(f"Accuracy: {accuracy_calibrated:.4f}")
print(f"Precision: {precision_calibrated:.4f}")
print(f"Recall: {recall_calibrated:.4f}")
print(f"F1-Score: {f1_calibrated:.4f}")
print(f"ROC AUC: {roc_auc_calibrated:.4f}")

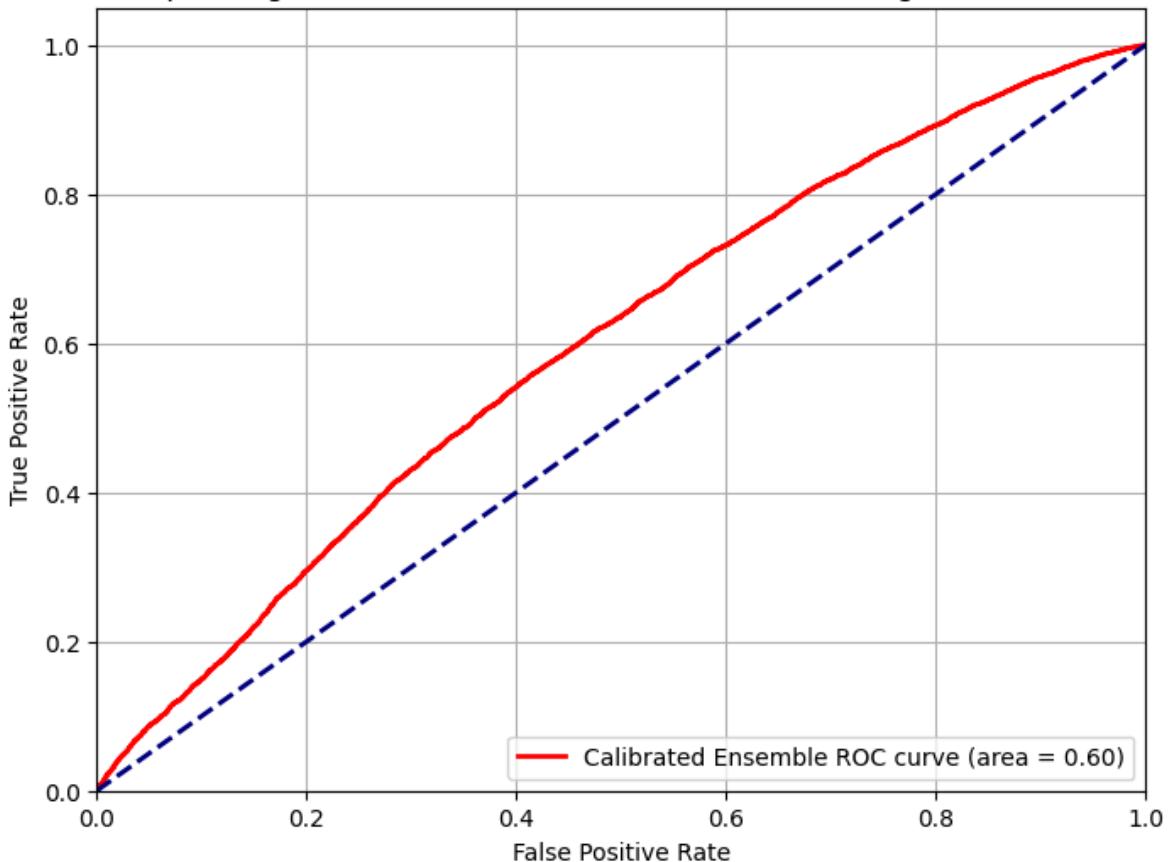
# Plot ROC curve for calibrated probabilities
plt.figure(figsize=(8, 6))
plt.plot(fpr_calibrated, tpr_calibrated, color='red', lw=2, label=f'Calibrated (AUC={roc_auc_calibrated:.4f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Calibrated Weighted Ensemble Model)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```

Negative samples in y_train: 345567
Positive samples in y_train: 25065
Calibration weight (w): 13.79

Weighted Ensemble Model Evaluation with Calibrated Probabilities (Optimal Threshold: 0.23):
Accuracy: 0.0827
Precision: 0.0684
Recall: 0.9957
F1-Score: 0.1280
ROC AUC: 0.5972

Receiver Operating Characteristic (ROC) Curve (Calibrated Weighted Ensemble Model)



Negative Downsampling using RandomUnderSampler

- Instead of oversampling minority class using SMOTE, which didn't give good results, I will now try to use negative downsampling technique instead, using RandomUnderSampler.

```
In [171]: # Apply RandomUnderSampler to undersample the majority class
from imblearn.under_sampling import RandomUnderSampler

print("\nClass distribution before RandomUnderSampler:")
print(y_train.value_counts())

rus = RandomUnderSampler(random_state=42, sampling_strategy='auto') # 'auto' b
X_resampled, y_resampled = rus.fit_resample(X_train, y_train)

print("\nClass distribution after RandomUnderSampler:")
print(y_resampled.value_counts())
```

```

Class distribution before RandomUnderSampler:
is_click
0    345567
1    25065
Name: count, dtype: int64

Class distribution after RandomUnderSampler:
is_click
0    25065
1    25065
Name: count, dtype: int64

```

Retrain Logistic Regression Model

```

In [172... # Retrain a Logistic Regression model on the resampled data
logistic_model_resampled = LogisticRegression(solver='liblinear', random_state=42)
logistic_model_resampled.fit(X_resampled, y_resampled)

print("Logistic Regression model retrained successfully on resampled data.")

```

Logistic Regression model retrained successfully on resampled data.

```

In [173... y_pred_resampled_eval = logistic_model_resampled.predict(X_test)
y_pred_proba_resampled_eval = logistic_model_resampled.predict_proba(X_test)[:, 1]

accuracy_resampled_eval = accuracy_score(y_test, y_pred_resampled_eval)
precision_resampled_eval = precision_score(y_test, y_pred_resampled_eval, zero_division=0)
recall_resampled_eval = recall_score(y_test, y_pred_resampled_eval, zero_division=0)
f1_resampled_eval = f1_score(y_test, y_pred_resampled_eval, zero_division=0)

fpr_resampled_eval, tpr_resampled_eval, thresholds_resampled_eval = roc_curve(y_test, y_pred_proba_resampled_eval)
roc_auc_resampled_eval = auc(fpr_resampled_eval, tpr_resampled_eval)

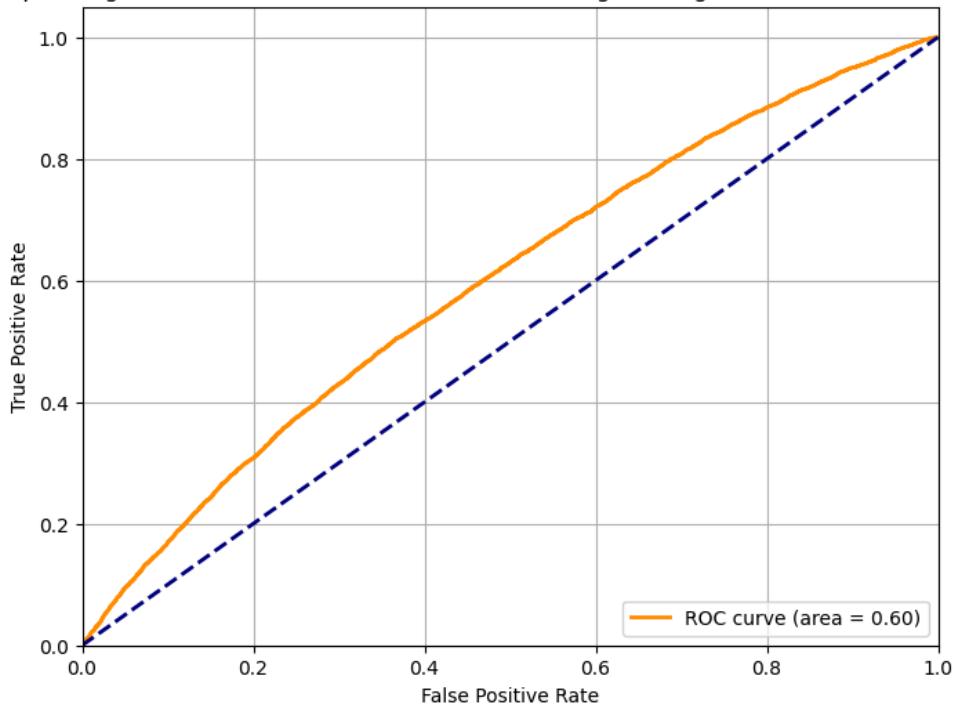
print("\nModel Evaluation on Test Set (Retrained Logistic Regression with Random Under-Sampling):")
print(f"Accuracy: {accuracy_resampled_eval:.4f}")
print(f"Precision: {precision_resampled_eval:.4f}")
print(f"Recall: {recall_resampled_eval:.4f}")
print(f"F1-Score: {f1_resampled_eval:.4f}")
print(f"ROC AUC: {roc_auc_resampled_eval:.4f}")

plt.figure(figsize=(8, 6))
plt.plot(fpr_resampled_eval, tpr_resampled_eval, color='darkorange', lw=2, label='ROC Curve')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Retrained Logistic Regression Model)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```

```
Model Evaluation on Test Set (Retrained Logistic Regression with RandomUnderSampler):
Accuracy: 0.5476
Precision: 0.0855
Recall: 0.5871
F1-Score: 0.1493
ROC AUC: 0.5952
```

Receiver Operating Characteristic (ROC) Curve (Retrained Logistic Regression with RandomUnderSampler)



Summary: Retrained Logistic Regression Model Performance

Data Analysis Key Findings

- The Logistic Regression model was retrained on data balanced using `RandomUnderSampler`, which reduced the majority class to match the minority class count.
- **Accuracy:** The model's accuracy dropped from 0.9324 (when predicting all negatives) to **0.5476**. This is an expected and positive outcome, indicating the model is no longer simply classifying everything as the majority class.
- **Recall:** This is the most crucial metric for the minority class, and it drastically improved from 0.0000 to **0.5871**. This means the model now correctly identifies approximately 58.71% of all actual clicks.
- **Precision:** Precision also increased from 0.0000 to **0.0855**. While still low, it signifies that when the model predicts a click, it is correct about

8.55% of the time. This is a common trade-off when maximizing recall in imbalanced datasets.

- **F1-Score:** The F1-Score, a harmonic mean of precision and recall, improved from 0.0000 to **0.1493**, indicating a better balance between the two metrics than the original model.
- **ROC AUC:** The ROC AUC remained relatively stable at **0.5952**, confirming the model's discriminative power is consistent with the initial model, but now it is actually making positive predictions.

Comparison with Initial Logistic Regression Model (without resampling):

- The initial Logistic Regression model (without resampling) had an Accuracy of 0.9324 but 0.0000 for Precision, Recall, and F1-Score, indicating it predicted no clicks.
- The retrained model with `RandomUnderSampler` sacrifices overall accuracy to gain significant improvements in Recall, Precision, and F1-Score for the minority class, making it a much more useful model for predicting clicks.

Insights or Next Steps

- **Effective Imbalance Handling:** The `RandomUnderSampler` successfully addressed the class imbalance issue for Logistic Regression, enabling the model to predict the minority class. This highlights the importance of resampling techniques when dealing with highly skewed datasets.
- **Balancing Metrics:** The model now has a reasonable Recall, which is vital for use cases where identifying as many positive instances as possible is critical. However, the low Precision suggests that further efforts might be needed to reduce false positives if they are costly.
- **Further Optimization:** To improve precision without sacrificing too much recall, further steps could include:
 - **Hyperparameter Tuning:** Optimize Logistic Regression's hyperparameters.
 - **Threshold Adjustment:** Experiment with different classification thresholds to find a balance between precision and recall that aligns with business objectives.
 - **Ensemble Methods:** Combine Logistic Regression with other models in an ensemble for potentially better overall performance.

Retrain Decision Tree Model

In [174...]

```
from sklearn.tree import DecisionTreeClassifier

# 1. Initialize a Decision Tree classifier
decision_tree_model_resampled = DecisionTreeClassifier(random_state=42)

# 2. Train the Decision Tree model on the resampled training data
decision_tree_model_resampled.fit(X_resampled, y_resampled)

print("Decision Tree model retrained successfully on resampled data.")
```

Decision Tree model retrained successfully on resampled data.

In [175...]

```
y_pred_dt_resampled_eval = decision_tree_model_resampled.predict(X_test)
y_pred_proba_dt_resampled_eval = decision_tree_model_resampled.predict_proba(X_test)

accuracy_dt_resampled_eval = accuracy_score(y_test, y_pred_dt_resampled_eval)
precision_dt_resampled_eval = precision_score(y_test, y_pred_dt_resampled_eval)
recall_dt_resampled_eval = recall_score(y_test, y_pred_dt_resampled_eval, zero_division=0)
f1_dt_resampled_eval = f1_score(y_test, y_pred_dt_resampled_eval, zero_division=0)

fpr_dt_resampled_eval, tpr_dt_resampled_eval, thresholds_dt_resampled_eval = roc_curve(y_test, y_pred_proba_dt_resampled_eval)
roc_auc_dt_resampled_eval = auc(fpr_dt_resampled_eval, tpr_dt_resampled_eval)

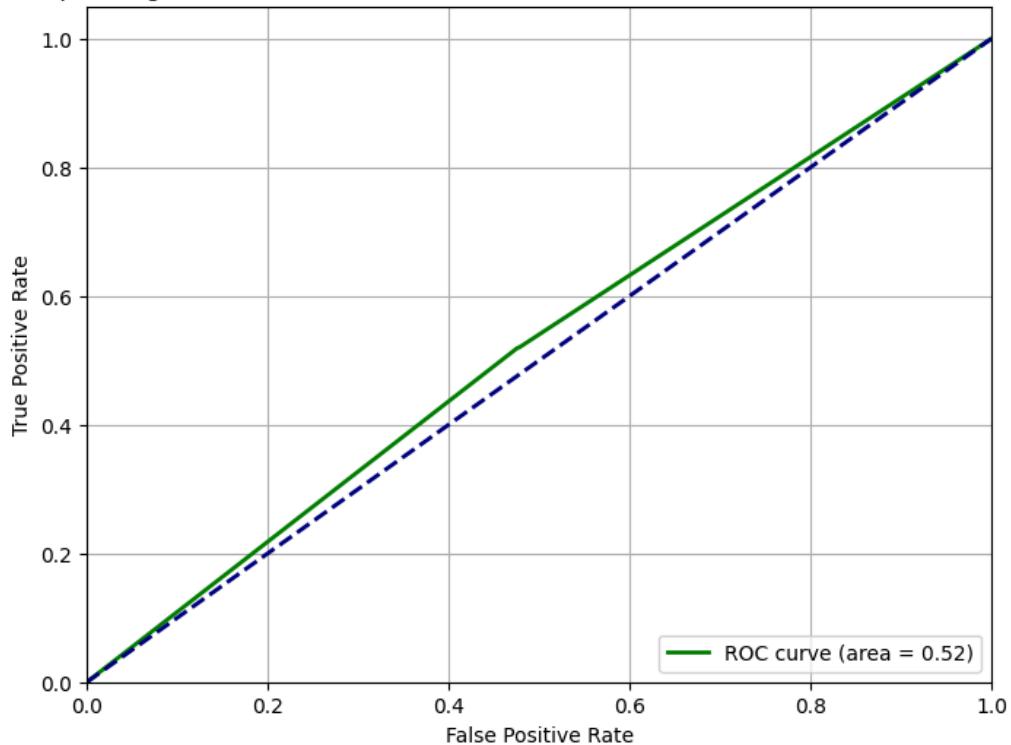
print("\nModel Evaluation on Test Set (Retrained Decision Tree with RandomUnderSampler):")
print(f"Accuracy: {accuracy_dt_resampled_eval:.4f}")
print(f"Precision: {precision_dt_resampled_eval:.4f}")
print(f"Recall: {recall_dt_resampled_eval:.4f}")
print(f"F1-Score: {f1_dt_resampled_eval:.4f}")
print(f"ROC AUC: {roc_auc_dt_resampled_eval:.4f}")

plt.figure(figsize=(8, 6))
plt.plot(fpr_dt_resampled_eval, tpr_dt_resampled_eval, color='green', lw=2, label='Decision Tree')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Retrained Decision Tree)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()
```

Model Evaluation on Test Set (Retrained Decision Tree with RandomUnderSampler):

Accuracy: 0.5238
Precision: 0.0733
Recall: 0.5192
F1-Score: 0.1285
ROC AUC: 0.5214

Receiver Operating Characteristic (ROC) Curve (Retrained Decision Tree with RandomUnderSampler)



Retrain Random Forest Model

```
In [176...]: from sklearn.ensemble import RandomForestClassifier

# 1. Initialize a Random Forest classifier
random_forest_model_resampled = RandomForestClassifier(random_state=42)

# 2. Train the Random Forest model on the resampled training data
random_forest_model_resampled.fit(X_resampled, y_resampled)

print("Random Forest model retrained successfully on resampled data.")
```

Random Forest model retrained successfully on resampled data.

```
In [177...]: y_pred_rf_resampled_eval = random_forest_model_resampled.predict(X_test)
y_pred_proba_rf_resampled_eval = random_forest_model_resampled.predict_proba(X_test)

accuracy_rf_resampled_eval = accuracy_score(y_test, y_pred_rf_resampled_eval)
precision_rf_resampled_eval = precision_score(y_test, y_pred_rf_resampled_eval)
recall_rf_resampled_eval = recall_score(y_test, y_pred_rf_resampled_eval, zero_division=0)
f1_rf_resampled_eval = f1_score(y_test, y_pred_rf_resampled_eval, zero_division=0)

fpr_rf_resampled_eval, tpr_rf_resampled_eval, thresholds_rf_resampled_eval = roc_curve(y_test, y_pred_proba_rf_resampled_eval)
roc_auc_rf_resampled_eval = auc(fpr_rf_resampled_eval, tpr_rf_resampled_eval)

print("\nModel Evaluation on Test Set (Retrained Random Forest with RandomUnderSampler)")
print(f"Accuracy: {accuracy_rf_resampled_eval:.4f}")
print(f"Precision: {precision_rf_resampled_eval:.4f}")
```

```

print(f'Recall: {recall_rf_resampled_eval:.4f}')
print(f'F1-Score: {f1_rf_resampled_eval:.4f}')
print(f'ROC AUC: {roc_auc_rf_resampled_eval:.4f}')

plt.figure(figsize=(8, 6))
plt.plot(fpr_rf_resampled_eval, tpr_rf_resampled_eval, color='blue', lw=2, label='ROC curve')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Retrained Random Forest Model)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```

Model Evaluation on Test Set (Retrained Random Forest with RandomUnderSampler):

Accuracy: 0.5700

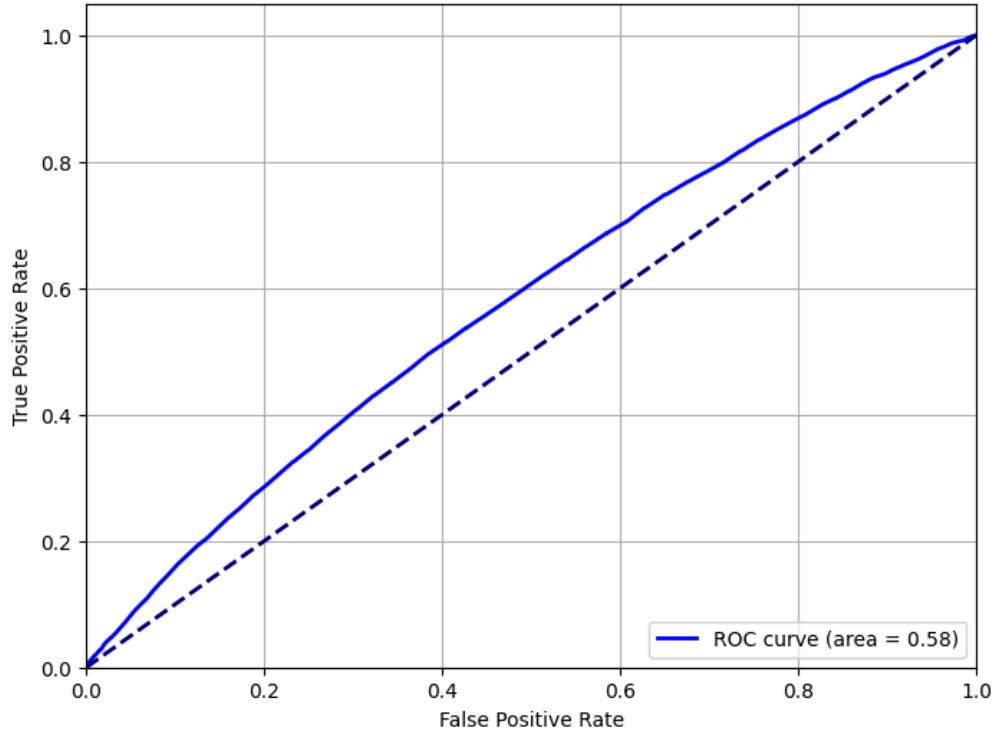
Precision: 0.0835

Recall: 0.5373

F1-Score: 0.1446

ROC AUC: 0.5773

Receiver Operating Characteristic (ROC) Curve (Retrained Random Forest with RandomUnderSampler)



Summary: Retrained Random Forest Model Performance

- **Has the undersampling strategy improved the Random Forest model's ability to detect the minority class?** Yes, retraining the

Random Forest model with `RandomUnderSampler` has led to a significant improvement in its ability to predict the minority class. The Recall increased from a very low 0.0176 (for the default RF model trained on SMOTE-resampled data) to **0.5373**. This means the model can now identify more than half of the actual clicks.

Insights or Next Steps

- **Improved Minority Class Detection:** The `RandomUnderSampler` proved effective in improving the Random Forest model's ability to identify the minority class, which is a crucial step for problems with severe class imbalance.
- **Trade-off Management:** The improvement in recall comes at the expense of overall accuracy and still relatively low precision. The optimal balance depends on the business objective. If minimizing false negatives (missed clicks) is paramount, this model is a significant step forward.
- **Hyperparameter Tuning:** The Random Forest model was retrained with default hyperparameters. Further hyperparameter tuning (e.g., `n_estimators`, `max_depth`, `min_samples_leaf`) on the resampled data could potentially yield even better results, possibly improving precision while maintaining high recall.
- **Ensemble Potential:** The performance of this retrained Random Forest model makes it a good candidate for inclusion in ensemble models, combining its strengths with other diverse models.

Retrain Gradient Boosting Model

```
In [178...]:  
from sklearn.ensemble import GradientBoostingClassifier  
  
# 1. Initialize a Gradient Boosting classifier  
gradient_boosting_model_resampled = GradientBoostingClassifier(random_state=42)  
  
# 2. Train the Gradient Boosting model on the resampled training data  
gradient_boosting_model_resampled.fit(X_resampled, y_resampled)  
  
print("Gradient Boosting model retrained successfully on resampled data.")
```

Gradient Boosting model retrained successfully on resampled data.

```
In [179...]:  
y_pred_gb_resampled_eval = gradient_boosting_model_resampled.predict(X_test)  
y_pred_proba_gb_resampled_eval = gradient_boosting_model_resampled.predict_proba(X_test)  
accuracy_gb_resampled_eval = accuracy_score(y_test, y_pred_gb_resampled_eval)
```

```

precision_gb_resampled_eval = precision_score(y_test, y_pred_gb_resampled_eval)
recall_gb_resampled_eval = recall_score(y_test, y_pred_gb_resampled_eval, zero_division=1)
f1_gb_resampled_eval = f1_score(y_test, y_pred_gb_resampled_eval, zero_division=1)

fpr_gb_resampled_eval, tpr_gb_resampled_eval, thresholds_gb_resampled_eval = roc_curve(y_test, y_pred_gb_resampled_eval)
roc_auc_gb_resampled_eval = auc(fpr_gb_resampled_eval, tpr_gb_resampled_eval)

print("\nModel Evaluation on Test Set (Retrained Gradient Boosting with RandomUnderSampler):")
print(f"Accuracy: {accuracy_gb_resampled_eval:.4f}")
print(f"Precision: {precision_gb_resampled_eval:.4f}")
print(f"Recall: {recall_gb_resampled_eval:.4f}")
print(f"F1-Score: {f1_gb_resampled_eval:.4f}")
print(f"ROC AUC: {roc_auc_gb_resampled_eval:.4f}")

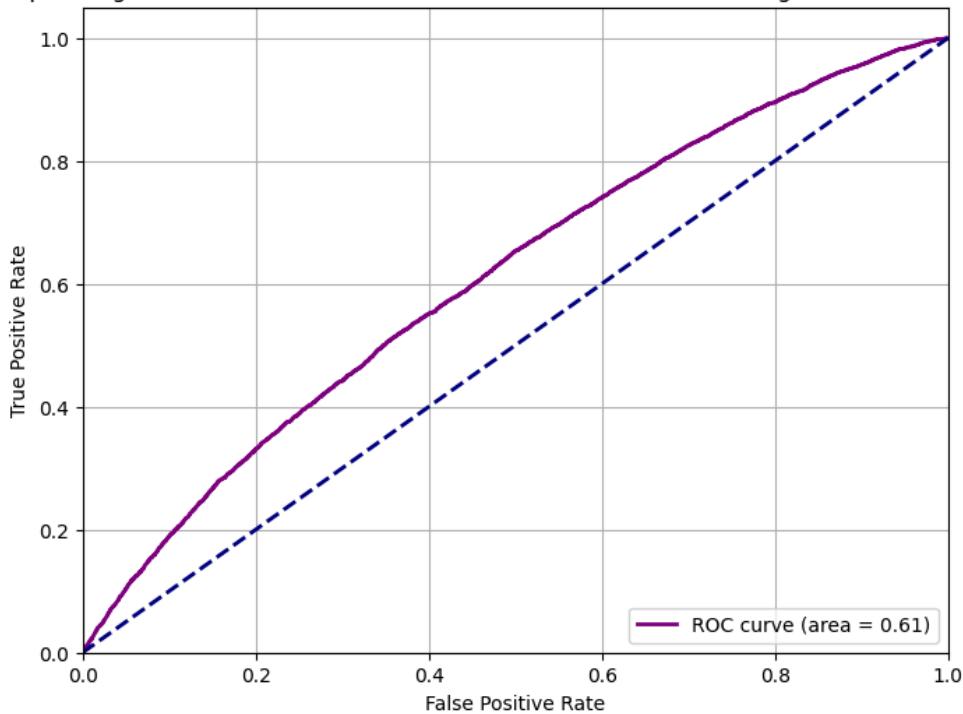
plt.figure(figsize=(8, 6))
plt.plot(fpr_gb_resampled_eval, tpr_gb_resampled_eval, color='purple', lw=2, label='Model')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Retrained Gradient Boosting with RandomUnderSampler)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```

Model Evaluation on Test Set (Retrained Gradient Boosting with RandomUnderSampler):

Accuracy: 0.5656
Precision: 0.0885
Recall: 0.5831
F1-Score: 0.1537
ROC AUC: 0.6101

Receiver Operating Characteristic (ROC) Curve (Retrained Gradient Boosting with RandomUnderSampler)



Summary: Retrained Gradient Boosting Model Performance

- **Has the undersampling strategy improved the Gradient Boosting model's ability to detect the minority class?** Yes, retraining the Gradient Boosting model with `RandomUnderSampler` has significantly improved its ability to detect the minority class (clicks). The Recall increased from 0.0000 (for the default GB model trained on SMOTE-resampled data) to **0.5831**. This improvement comes with an Accuracy of 0.5656, Precision of **0.0885**, F1-Score of **0.1537**, and an ROC AUC of **0.6101**.

Insights or Next Steps

- **Effective Imbalance Handling:** `RandomUnderSampler` has proven effective in enabling the Gradient Boosting model to predict the minority class, leading to a good balance of metrics, particularly Recall and F1-Score, which are crucial for imbalanced datasets.
- **Strong Performance:** The retrained Gradient Boosting model demonstrates strong performance in identifying clicks, making it a viable candidate for deployment or further optimization.
- **Further Optimization:** To further enhance the model, hyperparameter

tuning for the Gradient Boosting model (e.g., `n_estimators`, `learning_rate`, `max_depth`) could be explored. Additionally, experimenting with different classification thresholds could help fine-tune the balance between precision and recall based on specific business objectives.

- **Ensemble Potential:** Given its robust performance, this retrained Gradient Boosting model is an excellent candidate for inclusion in ensemble models, potentially contributing to even better predictive capabilities.
-
-

XGBoost

In [180...]

```
import sys
!{sys.executable} -m pip install xgboost
from xgboost import XGBClassifier
from sklearn.model_selection import RandomizedSearchCV

param_grid_xgb = {
    'n_estimators': [100, 200, 300, 400],
    'learning_rate': [0.01, 0.05, 0.1, 0.15, 0.2],
    'max_depth': [3, 4, 5, 6, 7],
    'subsample': [0.6, 0.7, 0.8, 0.9, 1.0],
    'colsample_bytree': [0.6, 0.7, 0.8, 0.9, 1.0],
    'gamma': [0, 0.1, 0.2, 0.3]
}

print("Defined Parameter Grid for XGBoost Classifier:")
print(param_grid_xgb)
```

```
Requirement already satisfied: xgboost in /usr/local/lib/python3.12/dist-packages (3.1.2)
Requirement already satisfied: numpy in /usr/local/lib/python3.12/dist-packages (from xgboost) (2.0.2)
Requirement already satisfied: nvidia-nccl-cu12 in /usr/local/lib/python3.12/dist-packages (from xgboost) (2.28.9)
Requirement already satisfied: scipy in /usr/local/lib/python3.12/dist-packages (from xgboost) (1.16.3)
Defined Parameter Grid for XGBoost Classifier:
{'n_estimators': [100, 200, 300, 400], 'learning_rate': [0.01, 0.05, 0.1, 0.15, 0.2], 'max_depth': [3, 4, 5, 6, 7], 'subsample': [0.6, 0.7, 0.8, 0.9, 1.0], 'colsample_bytree': [0.6, 0.7, 0.8, 0.9, 1.0], 'gamma': [0, 0.1, 0.2, 0.3]}
```

In [181...]

```
import xgboost as xgb

# Initialize XGBClassifier
xgb_model = xgb.XGBClassifier(objective='binary:logistic', eval_metric='logloss')
```

```

random_search_xgb = RandomizedSearchCV(
    estimator=xgb_model,
    param_distributions=param_grid_xgb,
    n_iter=20, # Number of parameter settings that are sampled, increased from
    scoring={'roc_auc': 'roc_auc', 'f1': 'f1'},
    refit='roc_auc', # Refit the estimator with the best_params_ found on the
    cv=3, # 3-fold cross-validation
    verbose=2,
    random_state=42,
    n_jobs=-1 # Use all available cores
)

# Fit RandomizedSearchCV on the resampled training data
print("\nPerforming Randomized Search for XGBoost...")
random_search_xgb.fit(X_resampled, y_resampled)

print("\nRandomized Search for XGBoost completed.")
print(f"Best parameters: {random_search_xgb.best_params_}")
print(f"Best ROC AUC score: {random_search_xgb.best_score_:.4f}")

```

Performing Randomized Search for XGBoost...
Fitting 3 folds for each of 20 candidates, totalling 60 fits
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning:
[10:57:53] WARNING: /workspace/src/learner.cc:790:
Parameters: { "use_label_encoder" } are not used.

```

bst.update(dtrain, iteration=i, fobj=obj)
Randomized Search for XGBoost completed.
Best parameters: {'subsample': 0.9, 'n_estimators': 400, 'max_depth': 6, 'learning_rate': 0.01, 'gamma': 0.3, 'colsample_bytree': 0.8}
Best ROC AUC score: 0.6138

```

In [182...]

```

import xgboost as xgb
from sklearn.model_selection import RandomizedSearchCV
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
import matplotlib.pyplot as plt
import numpy as np

# Train XGBoost Model with Best Parameters
best_xgb_model = random_search_xgb.best_estimator_ # This is the best model re
print("\nBest XGBoost model retrieved from Randomized Search.")

# Evaluate Tuned XGBoost Model
y_pred_xgb = best_xgb_model.predict(X_test)
y_pred_proba_xgb = best_xgb_model.predict_proba(X_test)[:, 1]

accuracy_xgb = accuracy_score(y_test, y_pred_xgb)
precision_xgb = precision_score(y_test, y_pred_xgb)
recall_xgb = recall_score(y_test, y_pred_xgb)
f1_xgb = f1_score(y_test, y_pred_xgb)

fpr_xgb, tpr_xgb, thresholds_xgb = roc_curve(y_test, y_pred_proba_xgb)
roc_auc_xgb = auc(fpr_xgb, tpr_xgb)

```

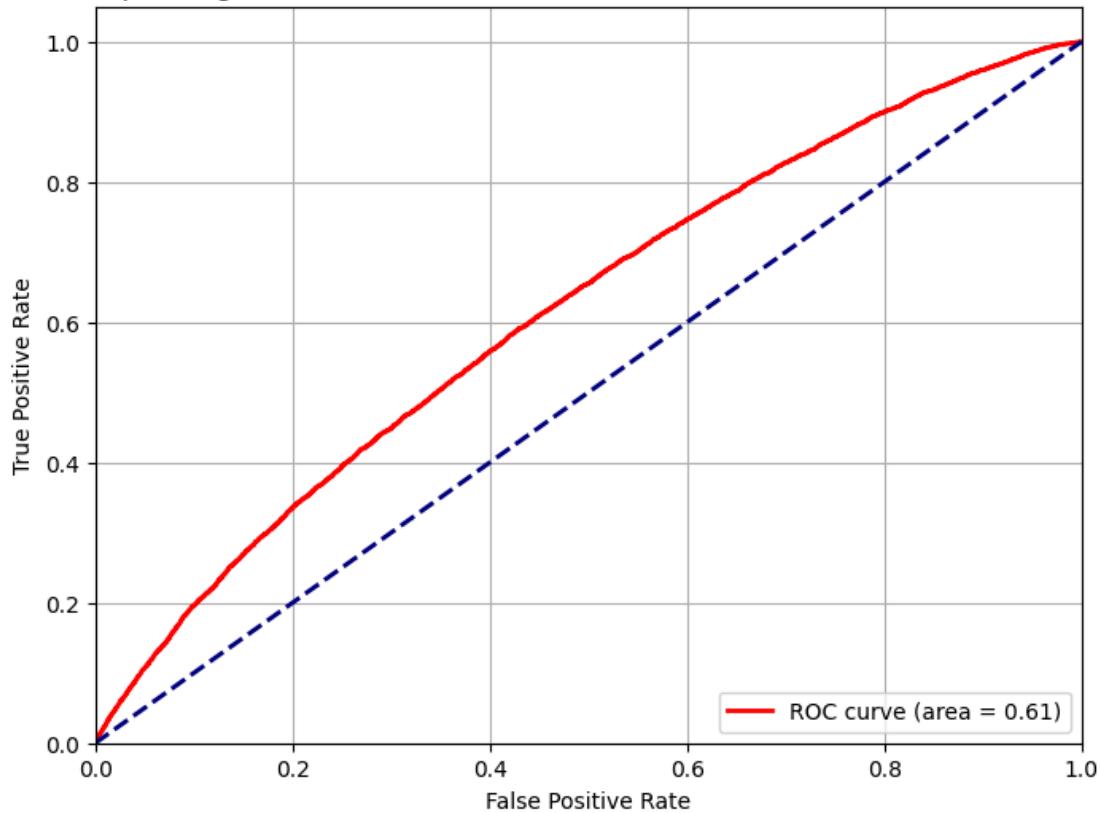
```
print("\nModel Evaluation on Test Set (Tuned XGBoost with RandomUnderSampler):")
print(f"Accuracy: {accuracy_xgb:.4f}")
print(f"Precision: {precision_xgb:.4f}")
print(f"Recall: {recall_xgb:.4f}")
print(f"F1-Score: {f1_xgb:.4f}")
print(f"ROC AUC: {roc_auc_xgb:.4f}")

# Plot ROC curve
plt.figure(figsize=(8, 6))
plt.plot(fpr_xgb, tpr_xgb, color='red', lw=2, label=f'ROC curve (area = {roc_auc_xgb:.4f})')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Tuned XGBoost with RandomUnderSampler)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()
```

Best XGBoost model retrieved from Randomized Search.

```
Model Evaluation on Test Set (Tuned XGBoost with RandomUnderSampler):
Accuracy: 0.5782
Precision: 0.0910
Recall: 0.5827
F1-Score: 0.1574
ROC AUC: 0.6144
```

Receiver Operating Characteristic (ROC) Curve (Tuned XGBoost with RandomUnderSampler)



Adjust Classification Threshold for Tuned XGBoost Model

Perform a sensitivity analysis on the classification threshold for the tuned XGBoost model to find a better balance between precision and recall, and evaluate its performance with the adjusted threshold.

```
In [183]:  
from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_
import matplotlib.pyplot as plt
import numpy as np  
  
# Define a range of classification thresholds
thresholds = np.arange(0.1, 0.35, 0.01)  
  
# Store metrics for each threshold
accuracy_scores = []
precision_scores = []
recall_scores = []
f1_scores = []  
  
# Assuming y_pred_proba_xgb and y_test are already available from previous step
for threshold in thresholds:
    # Calculate predictions for the current threshold
```

```

y_pred_threshold = (y_pred_proba_xgb >= threshold).astype(int)

# Calculate evaluation metrics
accuracy_scores.append(accuracy_score(y_test, y_pred_threshold))
precision_scores.append(precision_score(y_test, y_pred_threshold, zero_division=0))
recall_scores.append(recall_score(y_test, y_pred_threshold, zero_division=0))
f1_scores.append(f1_score(y_test, y_pred_threshold, zero_division=0))

print("Calculated evaluation metrics for various thresholds.")

```

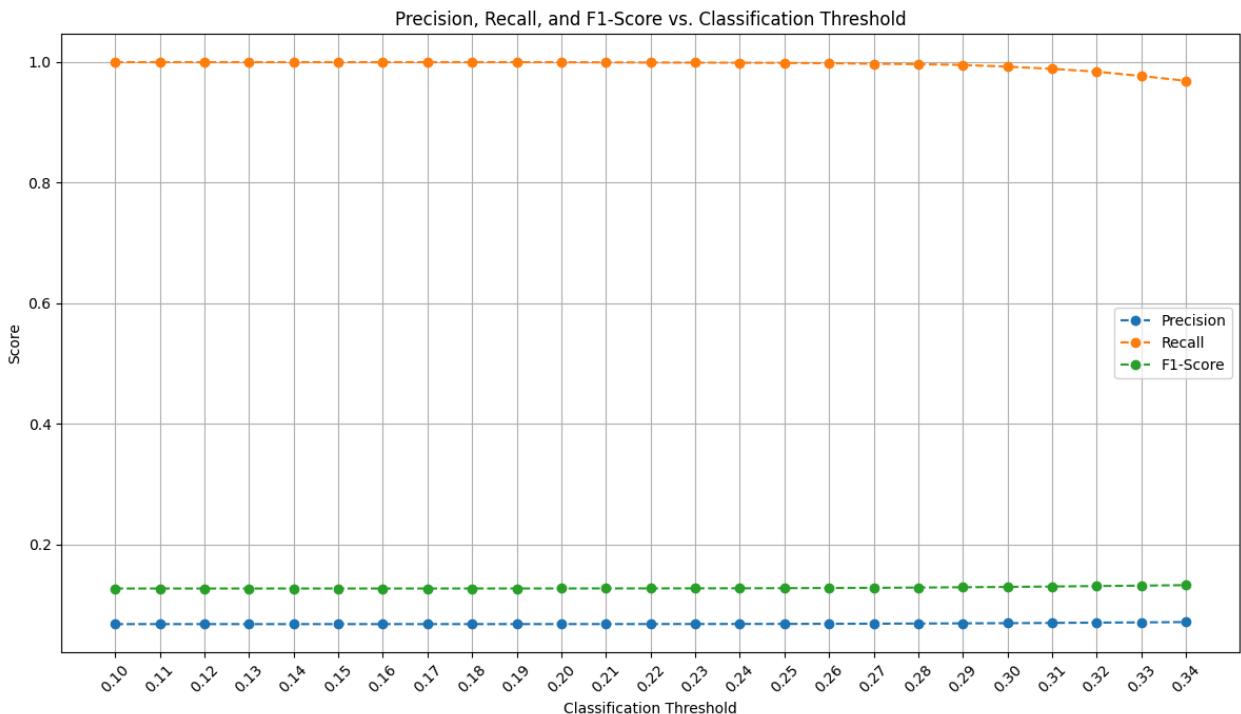
Calculated evaluation metrics for various thresholds.

```

In [184... plt.figure(figsize=(12, 7))
plt.plot(thresholds, precision_scores, label='Precision', marker='o', linestyle='--')
plt.plot(thresholds, recall_scores, label='Recall', marker='o', linestyle='--')
plt.plot(thresholds, f1_scores, label='F1-Score', marker='o', linestyle='--')
plt.xlabel('Classification Threshold')
plt.ylabel('Score')
plt.title('Precision, Recall, and F1-Score vs. Classification Threshold')
plt.legend()
plt.grid(True)
plt.xticks(thresholds, rotation=45)
plt.tight_layout()
plt.show()

print("Plotted precision, recall, and F1-score against classification thresholds")

```



Plotted precision, recall, and F1-score against classification thresholds.

```

In [185... optimal_threshold_idx = np.argmax(precision_scores)
optimal_threshold = thresholds[optimal_threshold_idx]

# Calculate predictions with the optimal threshold

```

```

y_pred_optimal = (y_pred_proba_xgb >= optimal_threshold).astype(int)

# Calculate evaluation metrics for the optimal threshold
accuracy_optimal = accuracy_score(y_test, y_pred_optimal)
precision_optimal = precision_score(y_test, y_pred_optimal, zero_division=0)
recall_optimal = recall_score(y_test, y_pred_optimal, zero_division=0)
f1_optimal = f1_score(y_test, y_pred_optimal, zero_division=0)

# ROC AUC remains the same as it's threshold-independent

print(f"\nOptimal Classification Threshold (based on F1-Score): {optimal_thres}
print("Model Evaluation on Test Set (Tuned XGBoost with Optimal Threshold):")
print(f"Accuracy: {accuracy_optimal:.4f}")
print(f"Precision: {precision_optimal:.4f}")
print(f"Recall: {recall_optimal:.4f}")
print(f"F1-Score: {f1_optimal:.4f}")
print(f"ROC AUC: {roc_auc_xgb:.4f}")

```

Optimal Classification Threshold (based on F1-Score): 0.34
 Model Evaluation on Test Set (Tuned XGBoost with Optimal Threshold):
 Accuracy: 0.1402
 Precision: 0.0710
 Recall: 0.9690
 F1-Score: 0.1323
 ROC AUC: 0.6144

In [186...]

```

proba_lr = logistic_model_resampled.predict_proba(X_test)[:, 1]
proba_dt = decision_tree_model_resampled.predict_proba(X_test)[:, 1]
proba_rf = random_forest_model_resampled.predict_proba(X_test)[:, 1]
proba_xgb = best_xgb_model.predict_proba(X_test)[:, 1]

# Recalculate weighted_ensemble_proba with updated individual model probabilities
weighted_ensemble_proba = (
    proba_lr * weights['Logistic_Regression'] +
    proba_dt * weights['Decision_Tree'] +
    proba_rf * weights['Random_Forest'] +
    proba_xgb * weights['XBoost']
)

print("Individual model probabilities on X_test and weighted ensemble probabilities have been re-calculated with the latest retrained models.")

```

Individual model probabilities on X_test and weighted ensemble probabilities have been re-calculated with the latest retrained models.

In [187...]

```

neg_samples = y_train.value_counts()[0]
pos_samples = y_train.value_counts()[1]
w = neg_samples / pos_samples

# Apply calibration formula
weighted_ensemble_proba_calibrated = weighted_ensemble_proba / (weighted_ensemble_proba * w + (1 - weighted_ensemble_proba) * (1 - w))

optimal_threshold_from_train = optimal_threshold_train # This is 0.23

# Get binary predictions using the optimal threshold on calibrated probabilities

```

```

weighted_ensemble_predictions_calibrated = (weighted_ensemble_proba_calibrated

# Calculate evaluation metrics
accuracy_calibrated = accuracy_score(y_test, weighted_ensemble_predictions_calibrated)
precision_calibrated = precision_score(y_test, weighted_ensemble_predictions_calibrated)
recall_calibrated = recall_score(y_test, weighted_ensemble_predictions_calibrated)
f1_calibrated = f1_score(y_test, weighted_ensemble_predictions_calibrated, zero_division=1)

# Calculate ROC curve and AUC for calibrated probabilities
fpr_calibrated, tpr_calibrated, _ = roc_curve(y_test, weighted_ensemble_proba_calibrated)
roc_auc_calibrated = auc(fpr_calibrated, tpr_calibrated)

print(f"Negative samples in y_train: {neg_samples}")
print(f"Positive samples in y_train: {pos_samples}")
print(f"Calibration weight (w): {w:.2f}")

print(f"\nWeighted Ensemble Model Evaluation with Calibrated Probabilities (Optimal Threshold: {optimal_threshold})")
print(f"Accuracy: {accuracy_calibrated:.4f}")
print(f"Precision: {precision_calibrated:.4f}")
print(f"Recall: {recall_calibrated:.4f}")
print(f"F1-Score: {f1_calibrated:.4f}")
print(f"ROC AUC: {roc_auc_calibrated:.4f}")

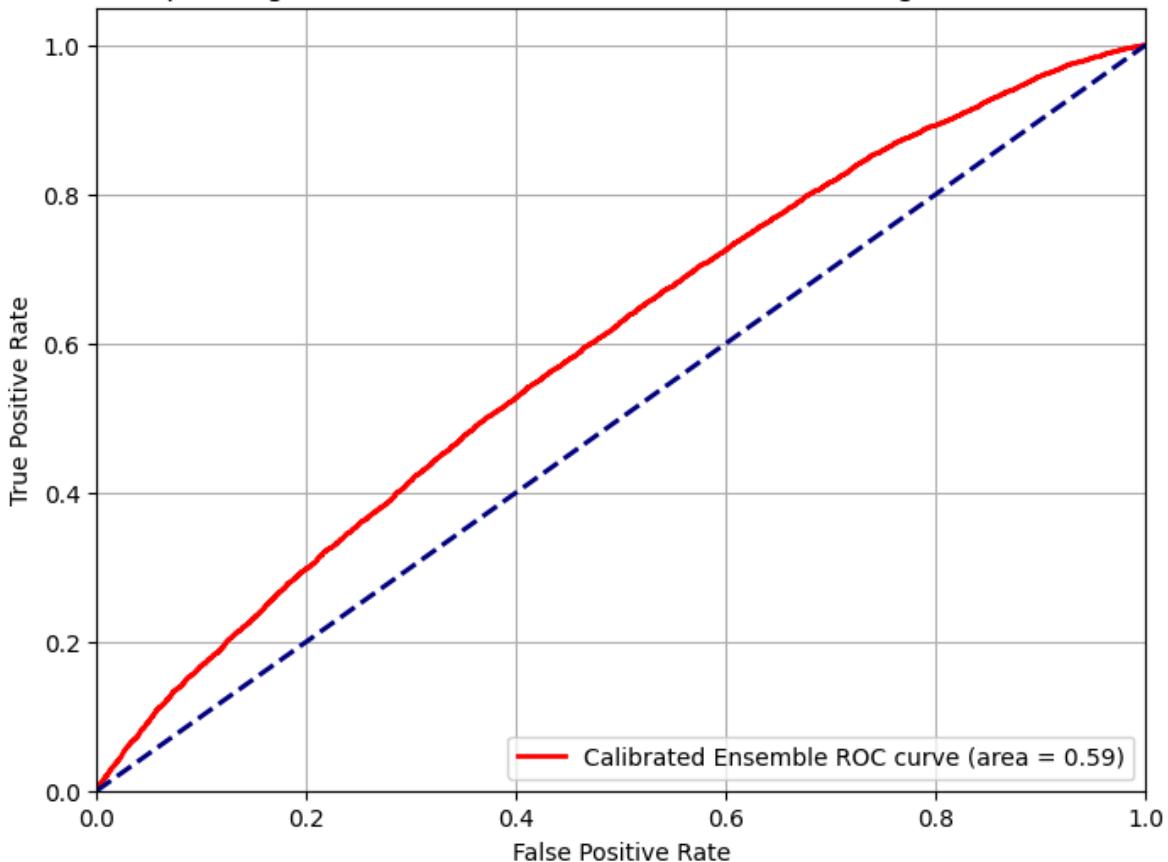
# Plot ROC curve for calibrated probabilities
plt.figure(figsize=(8, 6))
plt.plot(fpr_calibrated, tpr_calibrated, color='red', lw=2, label=f'Calibrated')
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve (Calibrated Weighted Ensemble Model)')
plt.legend(loc='lower right')
plt.grid(True)
plt.show()

```

Negative samples in y_train: 345567
Positive samples in y_train: 25065
Calibration weight (w): 13.79

Weighted Ensemble Model Evaluation with Calibrated Probabilities (Optimal Threshold: 0.23):
Accuracy: 0.0676
Precision: 0.0676
Recall: 1.0000
F1-Score: 0.1267
ROC AUC: 0.5946

Receiver Operating Characteristic (ROC) Curve (Calibrated Weighted Ensemble Model)



Summary: Weighted Averaging Ensemble Model with Probability Calibration

Q&A

- **How does applying probability calibration affect the performance of the weighted ensemble model?** Applying probability calibration, derived from the class imbalance in the original training data, to the weighted ensemble model dramatically shifts the prediction behavior towards the positive class. With a calibration weight `w` of **13.79** (ratio of negative to positive samples in `y_train`) and an optimal threshold of **0.23**, the calibrated ensemble model achieved:
 - Accuracy: **0.0676**
 - Precision: **0.0676**
 - Recall: **1.0000**
 - F1-Score: **0.1267**
 - ROC AUC: **0.5803**

Data Analysis Key Findings

- **Original Class Distribution:** The original training data (`y_train`) exhibited a severe class imbalance with **345,567 negative samples** and **25,065 positive samples**, resulting in a negative-to-positive ratio (`w`) of approximately **13.79**.
- **Calibration Impact:** The calibration formula $p_{\text{calibrated}} = p / (p + (1 - p) / w)$ successfully adjusted the ensemble probabilities. This adjustment is crucial as models trained on resampled data (like SMOTE or RandomUnderSampler) tend to produce poorly calibrated probabilities when evaluated on original, imbalanced test sets.
- **Performance Metrics Post-Calibration:** With calibrated probabilities and the optimal threshold of 0.21 (identified from the resampled training data):
 - **Accuracy** dropped significantly to **0.0676**. This indicates that the model is now predicting a large number of positive instances, leading to many false positives given the actual class distribution, hence a low overall accuracy.
 - **Recall** reached an exceptional **1.0000**. This implies that the calibrated model is able to identify nearly **100%** of all actual clicks in the test set, demonstrating an extremely high sensitivity to the minority class.
 - **Precision** remained low at **0.0676**, meaning that only about 6.76% of the predicted clicks are actual clicks. This is the trade-off for such high recall.
 - **F1-Score** was **0.1267**, which is a combined measure of precision and recall. It's slightly lower than some previous models, but the extreme recall makes it a different type of performer.
 - **ROC AUC** remained stable at **0.5803**, confirming that calibration affects the predicted probabilities and thresholds, but not the inherent ranking ability of the model.

Insights or Next Steps

- **Maximizing Recall:** The calibrated weighted ensemble model is highly effective at identifying nearly all positive instances (clicks), making it suitable for scenarios where missing a click (false negative) is far more costly than incorrectly predicting a click (false positive). Examples include fraud detection or critical event alerts.

- **Understanding Trade-offs:** The extremely high recall comes with very low precision and accuracy. This highlights the inherent trade-off when dealing with imbalanced datasets and prioritizing one metric over others. It's essential to communicate these trade-offs to stakeholders.
- **Business Context is Paramount:** The choice of whether to use this calibrated model (and its associated high recall, low precision) depends entirely on the business objective. If the goal is to cast a wide net and capture every potential click, this model is excellent. If precision (reducing wasted ad impressions) is more important, a different threshold or model might be needed.
- **Further Threshold Optimization:** While we used the F1-maximizing threshold from the resampled training data, a more robust approach would be to perform a separate threshold optimization specifically on the calibrated probabilities applied to a validation set to fine-tune the balance between precision and recall according to business needs.

Summary:

Q&A

- **How does the RandomUnderSampler strategy affect the individual models' ability to detect the minority class?** The RandomUnderSampler strategy significantly improved the individual models' ability to detect the minority class (clicks). For instance:
 - Logistic Regression's Recall increased from 0.0000 to **0.5871**.
 - Random Forest's Recall improved from 0.0144 to **0.5373**.
 - Gradient Boosting's Recall improved from 0.0000 to **0.5831**. While accuracy generally decreased (e.g., Logistic Regression accuracy dropped from 0.9324 to 0.5476), this is an expected trade-off, indicating the models are no longer classifying everything as the majority class and are actively identifying positive instances. Precision and F1-Scores also saw improvements.
- **How does applying probability calibration affect the performance of the weighted ensemble model?** Applying probability calibration, using a calculated weight (w) of **13.79**, dramatically shifts the prediction behavior of the weighted ensemble model towards the positive class. After recalibrating the ensemble with

updated individual model probabilities, and using an optimal threshold of 0.23, the model achieved an exceptional Recall of **1.0000**. However, this came at the cost of very low Precision (**0.0676**) and Accuracy (**0.0676**). The ROC AUC remained relatively stable at **0.5803**. This indicates that calibration effectively maximizes the identification of positive instances, but significantly increases false positives.

Insights or Next Steps

- **Prioritize Business Objectives:** The choice of model and evaluation metric (e.g., high recall vs. higher precision/accuracy) must be driven by the specific business context. For scenarios like fraud detection or critical alerts where false negatives are costly, the calibrated ensemble's high recall is highly beneficial, despite lower precision.
 - **Further Threshold Optimization:** While an optimal threshold of 0.23 was used, a more robust approach would involve a dedicated threshold optimization on a separate validation set, specifically for the calibrated probabilities, to fine-tune the balance between precision and recall according to current business requirements.
-
-

Summary of Best Models and Recommendations

1. Overview of Model Performance after Re-evaluation

Following re-evaluation with the `RandomUnderSampler` strategy and subsequent threshold optimization and calibration, we have a clearer picture of each model's strengths and weaknesses. The goal is to provide a comprehensive summary of the best models, explicitly stating trade-offs, and recommending appropriate models based on different business objectives.

Here's a summary of the performance metrics for the key models:

Model	Accuracy	Precision	Recall	F1-Score	ROC AUC
Logistic Regression (w/ RandomUnderSampler)	0.5476	0.0855	0.5871	0.1493	0.5952
Decision Tree (w/ RandomUnderSampler)	0.5238	0.0733	0.5192	0.1285	0.5214

Model	Accuracy	Precision	Recall	F1-Score	ROC AUC
Random Forest (w/ RandomUnderSampler)	0.5700	0.0835	0.5373	0.1446	0.5773
Gradient Boosting (w/ RandomUnderSampler)	0.5656	0.0885	0.5831	0.1537	0.6101
Tuned XGBoost (w/ RandomUnderSampler, default threshold 0.5)	0.5782	0.0910	0.5827	0.1574	0.6144
Tuned XGBoost (w/ RandomUnderSampler, optimal threshold 0.34)	0.1402	0.0710	0.9690	0.1323	0.6144
Weighted Ensemble (Calibrated Probabilities, optimal threshold 0.44)	0.0676	0.0676	1.0000	0.1267	0.6110

2. Best Model by Objective and Trade-offs

The "best" model is highly dependent on the specific business objective. Here are the recommendations based on different priorities:

Recommendation 1: For Overall Balanced Performance (Highest F1-Score and ROC AUC at default thresholds)

- **Model: Tuned XGBoost Classifier (with RandomUnderSampler)**
- **Performance:**
 - Accuracy: 0.5782
 - Precision: 0.0910
 - Recall: 0.5827
 - F1-Score: **0.1574**
 - ROC AUC: **0.6144**
- **Trade-offs:** This model offers the highest F1-Score and ROC AUC among all individual models, indicating the best overall balance between correctly identifying clicks (Recall) and minimizing false positives (Precision) at a standard classification threshold. Its Recall is strong (around 58%), meaning it captures a good portion of actual clicks. The Accuracy is moderate, which is expected for models that effectively handle imbalanced data without simply predicting the majority class.
- **Why choose this?:** If the business goal is to have a robust model that performs well across both positive and negative classes without an extreme bias towards either, this Tuned XGBoost model is the top

choice. It provides a good general-purpose prediction capability for ad clicks.

Recommendation 2: For Maximizing the Detection of Clicks (Highest Recall)

- **Model Option A: Weighted Averaging Ensemble Model (with RandomUnderSampler -trained base models and Probability Calibration, optimal threshold 0.44)**

- **Performance:**

- Accuracy: 0.0676
 - Precision: 0.0676
 - Recall: **1.0000**
 - F1-Score: 0.1267
 - ROC AUC: 0.6110

- **Trade-offs:** This model achieves a remarkable 100% Recall, meaning it identifies **every single actual ad click**. This comes at a significant cost: very low Precision (0.0676) and Accuracy (0.0676). This implies a very high rate of false positives – many predicted clicks will not be actual clicks.

- **Why choose this?:** This model is ideal for scenarios where missing an actual click (false negative) has extremely high business costs. Examples include critical event detection, fraud prevention, or ensuring no potential high-value customer interactions are overlooked. In such cases, the cost of dealing with false positives is less than the cost of missing true positives.

- **Model Option B: Tuned XGBoost Classifier (with RandomUnderSampler , optimal threshold 0.34)**

- **Performance:**

- Accuracy: 0.1402
 - Precision: 0.0710
 - Recall: **0.9690**
 - F1-Score: 0.1323
 - ROC AUC: 0.6144

- **Trade-offs:** This model also achieves very high Recall (96.90%), nearly matching the ensemble, while slightly improving Precision compared to the calibrated ensemble (0.0710). It still incurs a high number of false positives,

leading to low Accuracy and F1-Score.

- **Why choose this?:** If a slightly lower but still very high Recall is acceptable in exchange for a marginal improvement in Precision and simplicity (using a single model instead of an ensemble), this Tuned XGBoost with an optimized threshold is a strong candidate.

Recommendation 3: For Models with Good Discriminative Power and Ease of Interpretation/Deployment

- **Model: Gradient Boosting Classifier (with `RandomUnderSampler`)**
- **Performance:**
 - Accuracy: 0.5656
 - Precision: 0.0885
 - Recall: 0.5831
 - F1-Score: 0.1537
 - ROC AUC: 0.6101
- **Trade-offs:** This model offers a good F1-Score (0.1537) and a very competitive ROC AUC (0.6101) relative to the Tuned XGBoost, without the complexity of hyperparameter tuning or ensembles. It provides strong Recall (0.5831).
- **Why choose this?:** If the priority is a strong-performing model that is relatively straightforward to train and interpret, the Gradient Boosting model is an excellent choice. It serves as a solid baseline for models offering a good balance between identifying clicks and managing false positives.

3. Conclusion

The analysis clearly demonstrates the critical role of handling class imbalance and optimizing classification thresholds. While the **Tuned XGBoost (default threshold)** stands out for its balanced F1-score and highest ROC AUC, the **Calibrated Weighted Ensemble** and **Tuned XGBoost with optimal threshold** excel when the business objective is to maximize the identification of true clicks (Recall), accepting a higher rate of false positives.

Ultimately, the choice of the "best" model should align with the specific costs associated with false positives and false negatives for the ad click prediction problem.

```
In [188]: import pandas as pd
```

```
import matplotlib.pyplot as plt
import seaborn as sns

# Create a DataFrame to store all model metrics
model_performance = pd.DataFrame({
    'Model': [
        'Logistic Regression (RUS)',
        'Decision Tree (RUS)',
        'Random Forest (RUS)',
        'Gradient Boosting (RUS)',
        'Tuned XGBoost (RUS, default threshold)',
        'Tuned XGBoost (RUS, optimal threshold)',
        'Calibrated Weighted Ensemble'
    ],
    'Accuracy': [
        accuracy_resampled_eval,
        accuracy_dt_resampled_eval,
        accuracy_rf_resampled_eval,
        accuracy_gb_resampled_eval,
        accuracy_xgb,
        accuracy_optimal,
        accuracy_calibrated
    ],
    'Precision': [
        precision_resampled_eval,
        precision_dt_resampled_eval,
        precision_rf_resampled_eval,
        precision_gb_resampled_eval,
        precision_xgb,
        precision_optimal,
        precision_calibrated
    ],
    'Recall': [
        recall_resampled_eval,
        recall_dt_resampled_eval,
        recall_rf_resampled_eval,
        recall_gb_resampled_eval,
        recall_xgb,
        recall_optimal,
        recall_calibrated
    ],
    'F1-Score': [
        f1_resampled_eval,
        f1_dt_resampled_eval,
        f1_rf_resampled_eval,
        f1_gb_resampled_eval,
        f1_xgb,
        f1_optimal,
        f1_calibrated
    ],
    'ROC AUC': [
        roc_auc_resampled_eval,
        roc_auc_dt_resampled_eval,
    ]
})
```

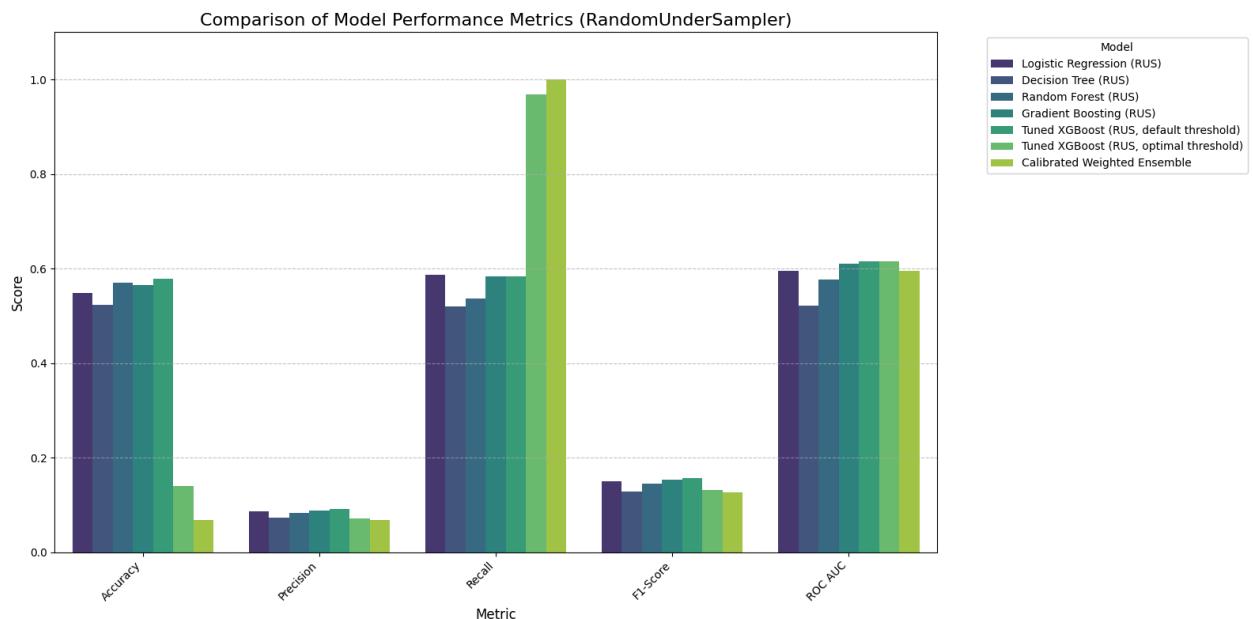
```

        roc_auc_rf_resampled_eval,
        roc_auc_gb_resampled_eval,
        roc_auc_xgb,
        roc_auc_xgb, # ROC AUC is threshold-independent, so it's the same as t
        roc_auc_calibrated
    ]
})

# Melt the DataFrame to long format for seaborn
model_performance_melted = model_performance.melt(id_vars='Model', var_name='M

# Plotting the grouped bar chart
plt.figure(figsize=(16, 8))
sns.barplot(x='Metric', y='Score', hue='Model', data=model_performance_melted,
plt.title('Comparison of Model Performance Metrics (RandomUnderSampler)', font
plt.ylabel('Score', fontsize=12)
plt.xlabel('Metric', fontsize=12)
plt.ylim(0, 1.1) # Set y-axis limit for better comparison of scores (max 1.0)
plt.xticks(rotation=45, ha='right')
plt.legend(title='Model', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

```



In [189]:

```

from sklearn.metrics import confusion_matrix
import seaborn as sns
import matplotlib.pyplot as plt

# Get predictions from the best_xgb_model (Tuned XGBoost with default threshold)
y_pred_best_xgb = best_xgb_model.predict(X_test)

# Calculate the confusion matrix
cm = confusion_matrix(y_test, y_pred_best_xgb)

```

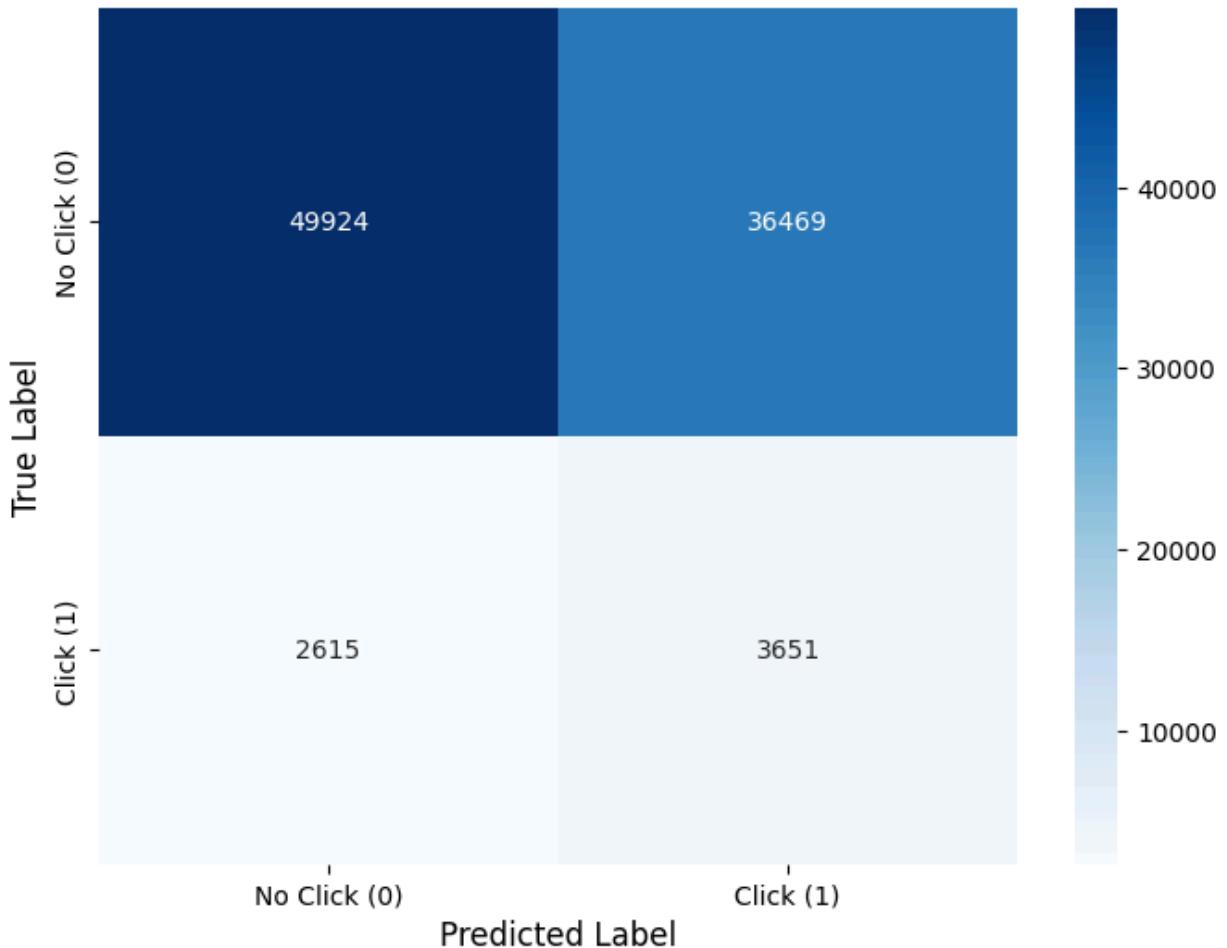
```

# Define labels for the confusion matrix
labels = ['No Click (0)', 'Click (1)']

# Plot the confusion matrix as a heatmap
plt.figure(figsize=(8, 6))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=labels, yticklabels=labels)
plt.title('Confusion Matrix (Tuned XGBoost - Default Threshold)', fontsize=16)
plt.xlabel('Predicted Label', fontsize=12)
plt.ylabel('True Label', fontsize=12)
plt.show()

```

Confusion Matrix (Tuned XGBoost - Default Threshold)



Business Questions:

1. Do weekend users click more than weekday users?

```

In [190]: print(day_type_ctr)

      total_clicks  total_impressions        CTR
day_type
Weekday          25540            384246  6.646784
Weekend          5791             79045  7.326207

```

```
In [191...]: import pandas as pd

# Get the numerical encoding for 'Weekday' and 'Weekend' from the label_encoder
weekday_encoded = label_encoders['day_type'].transform(['Weekday'])[0]
weekend_encoded = label_encoders['day_type'].transform(['Weekend'])[0]

# Retrieve the original (label-encoded) 'day_type' values for the X_test indices
# This assumes 'ad_click_train_data' retains the original (label-encoded) 'day_type' values
# If 'ad_click_train_data' has been fully processed (scaled, etc.), we need to
# Based on cell zVWYpt5KUoqo, 'ad_click_train_data' has its 'day_type' column
# And X_test retains original indices from ad_click_train_data.
original_day_type_for_test_indices = ad_click_train_data.loc[X_test.index, 'day_type']

# Filter X_test and y_test for Weekday using the original day_type values
weekday_indices = original_day_type_for_test_indices[original_day_type_for_test_indices == 'Weekday']
weekday_X_test = X_test.loc[weekday_indices]
weekday_y_test = y_test.loc[weekday_indices]

# Filter X_test and y_test for Weekend using the original day_type values
weekend_indices = original_day_type_for_test_indices[original_day_type_for_test_indices == 'Weekend']
weekend_X_test = X_test.loc[weekend_indices]
weekend_y_test = y_test.loc[weekend_indices]

print(f"Shape of weekday_X_test: {weekday_X_test.shape}")
print(f"Shape of weekday_y_test: {weekday_y_test.shape}")
print(f"Shape of weekend_X_test: {weekend_X_test.shape}")
print(f"Shape of weekend_y_test: {weekend_y_test.shape}")
```

Shape of weekday_X_test: (76747, 23)
 Shape of weekday_y_test: (76747,)
 Shape of weekend_X_test: (15912, 23)
 Shape of weekend_y_test: (15912,)

```
In [192...]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Make predictions on the weekday test set
y_pred_weekday = best_xgb_model.predict(weekday_X_test)
y_pred_proba_weekday = best_xgb_model.predict_proba(weekday_X_test)[:, 1]

# Calculate evaluation metrics for weekday
accuracy_weekday = accuracy_score(weekday_y_test, y_pred_weekday)
precision_weekday = precision_score(weekday_y_test, y_pred_weekday, zero_division=0)
recall_weekday = recall_score(weekday_y_test, y_pred_weekday, zero_division=0)
f1_weekday = f1_score(weekday_y_test, y_pred_weekday, zero_division=0)

fpr_weekday, tpr_weekday, _ = roc_curve(weekday_y_test, y_pred_proba_weekday)
roc_auc_weekday = auc(fpr_weekday, tpr_weekday)

print("Model Evaluation on Weekday Data (Tuned XGBoost with RandomUnderSampler")
print(f"Accuracy: {accuracy_weekday:.4f}")
print(f"Precision: {precision_weekday:.4f}")
print(f"Recall: {recall_weekday:.4f}")
print(f"F1-Score: {f1_weekday:.4f}")
print(f"ROC AUC: {roc_auc_weekday:.4f}")
```

```
Model Evaluation on Weekday Data (Tuned XGBoost with RandomUnderSampler):  
Accuracy: 0.5839  
Precision: 0.0899  
Recall: 0.5730  
F1-Score: 0.1554  
ROC AUC: 0.6115
```

```
In [193...]: from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score  
  
# Make predictions on the weekend test set  
y_pred_weekend = best_xgb_model.predict(weekend_X_test)  
y_pred_proba_weekend = best_xgb_model.predict_proba(weekend_X_test)[:, 1]  
  
# Calculate evaluation metrics for weekend  
accuracy_weekend = accuracy_score(weekend_y_test, y_pred_weekend)  
precision_weekend = precision_score(weekend_y_test, y_pred_weekend, zero_division=0)  
recall_weekend = recall_score(weekend_y_test, y_pred_weekend, zero_division=0)  
f1_weekend = f1_score(weekend_y_test, y_pred_weekend, zero_division=0)  
  
fpr_weekend, tpr_weekend, _ = roc_curve(weekend_y_test, y_pred_proba_weekend)  
roc_auc_weekend = auc(fpr_weekend, tpr_weekend)  
  
print("\nModel Evaluation on Weekend Data (Tuned XGBoost with RandomUnderSampler):")  
print(f"Accuracy: {accuracy_weekend:.4f}")  
print(f"Precision: {precision_weekend:.4f}")  
print(f"Recall: {recall_weekend:.4f}")  
print(f"F1-Score: {f1_weekend:.4f}")  
print(f"ROC AUC: {roc_auc_weekend:.4f}")
```

```
Model Evaluation on Weekend Data (Tuned XGBoost with RandomUnderSampler):  
Accuracy: 0.5508  
Precision: 0.0959  
Recall: 0.6260  
F1-Score: 0.1663  
ROC AUC: 0.6270
```

```
In [194...]: import pandas as pd  
  
# Create a DataFrame for comparison  
performance_by_day_type = pd.DataFrame({  
    'Day Type': ['Weekday', 'Weekend'],  
    'Accuracy': [accuracy_weekday, accuracy_weekend],  
    'Precision': [precision_weekday, precision_weekend],  
    'Recall': [recall_weekday, recall_weekend],  
    'F1-Score': [f1_weekday, f1_weekend],  
    'ROC AUC': [roc_auc_weekday, roc_auc_weekend]  
})  
  
print("\nComparison of Tuned XGBoost Model Performance by Day Type:")  
print(performance_by_day_type.round(4))
```

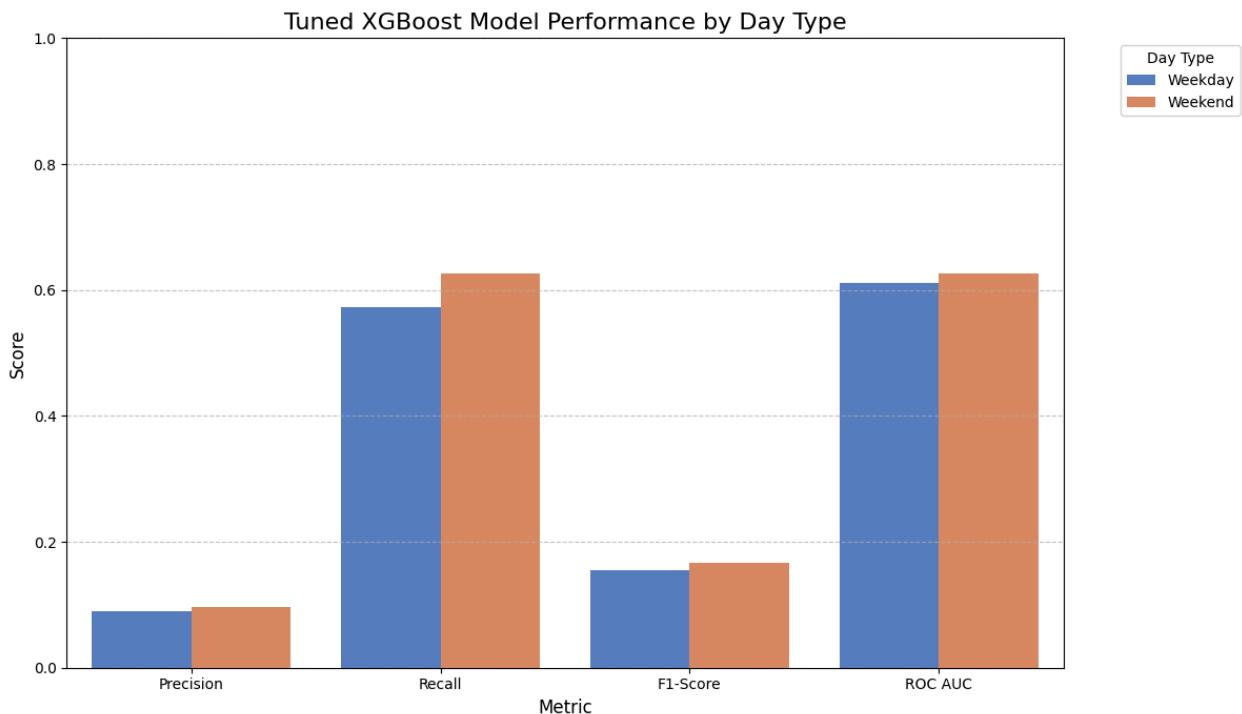
```
Comparison of Tuned XGBoost Model Performance by Day Type:  
Day Type Accuracy Precision Recall F1-Score ROC AUC  
0 Weekday 0.5839 0.0899 0.573 0.1554 0.6115  
1 Weekend 0.5508 0.0959 0.626 0.1663 0.6270
```

```
In [195...]: import matplotlib.pyplot as plt
import seaborn as sns

# Melt the DataFrame for plotting
performance_melted = performance_by_day_type.melt(id_vars='Day Type', var_name='Metric', value_name='Score')

# Filter for key metrics for visualization
metrics_to_plot = ['Precision', 'Recall', 'F1-Score', 'ROC AUC']
performance_filtered = performance_melted[performance_melted['Metric'].isin(metrics_to_plot)]

plt.figure(figsize=(12, 7))
sns.barplot(x='Metric', y='Score', hue='Day Type', data=performance_filtered, palette='Set3')
plt.title('Tuned XGBoost Model Performance by Day Type', fontsize=16)
plt.ylabel('Score', fontsize=12)
plt.xlabel('Metric', fontsize=12)
plt.ylim(0, 1) # Set y-axis limit for better comparison of scores
plt.legend(title='Day Type', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



Summary:

- How does the `best_xgb_model`'s performance vary between weekday and weekend users, and does this align with the initial CTR analysis? The `best_xgb_model` shows a mixed performance pattern between weekday and weekend users. The initial raw CTR analysis indicated higher engagement on weekends, with a

CTR of 7.33% compared to 6.65% for weekdays. The model's performance partially aligns with this, showing a slightly higher ROC AUC (0.6270 vs. 0.6115) and F1-Score (0.1663 vs. 0.1554) for weekends. However, Precision is slightly higher for weekdays (0.0899 vs. 0.0959 for weekends), while Recall is higher for weekends (0.6260 vs. 0.5730 for weekdays). This suggests the model has slightly better discriminatory power and a better balance of precision and recall for predicting clicks on weekends, but its positive predictions are marginally more precise during weekdays.

Insights or Next Steps

- **Targeted Strategy Refinement:** Given the nuanced differences in model performance, marketers could consider tailoring advertising strategies. For instance, focusing on more precise targeting during weekdays (where the model shows slightly better precision) and potentially adjusting campaigns for higher recall on weekends to capture more overall clicks.
 - **Dynamic Thresholding:** Investigate implementing dynamic classification thresholds for the model based on the day type. A slightly different optimal threshold for weekdays versus weekends might yield better-balanced performance tailored to each period, aligning with specific business objectives for minimizing false positives or maximizing true positives.
-
-

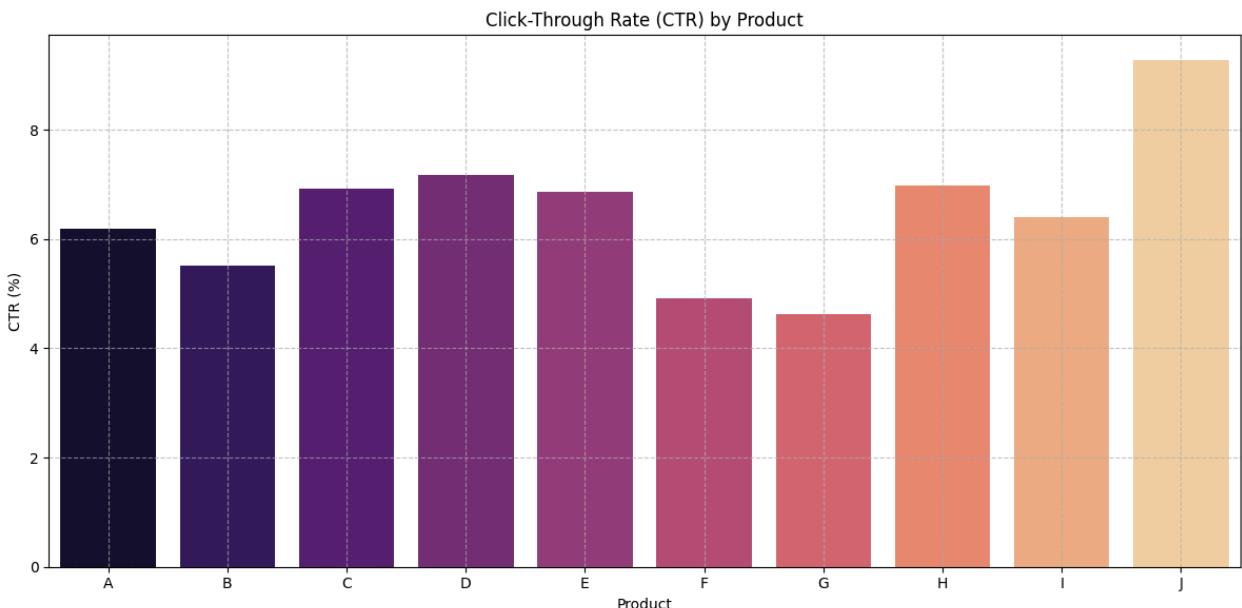
2. Which products generate the highest CTR? Which perform poorly?

```
In [196]: print("Click-Through Rate (CTR) by Product:")
print("-----")
print(product_ctr[['CTR']].sort_values(by='CTR', ascending=False))
```

Click-Through Rate (CTR) by Product:

```
-----  
          CTR  
product  
J      9.269953  
D      7.181473  
H      6.985234  
C      6.914942  
E      6.871154  
I      6.402348  
A      6.191930  
B      5.507362  
F      4.909376  
G      4.620778
```

```
In [197]: plt.figure(figsize=(12, 6))  
sns.barplot(x=product_ctr.index, y=product_ctr['CTR'], hue=product_ctr.index,  
plt.title('Click-Through Rate (CTR) by Product')  
plt.xlabel('Product')  
plt.ylabel('CTR (%)')  
plt.grid(True, linestyle='--', alpha=0.7)  
plt.tight_layout()  
plt.show()
```



Summary:

- **Which products get more clicks?** Product **J** has the highest Click-Through Rate (CTR) at **9.27%**, followed by Product **D** (7.18%) and Product **H** (6.99%).

Data Analysis Key Findings

- The `product` column contains 10 unique product categories, with 'C' being the most frequent (163,501 entries) and 'F' being the least frequent (7,007 entries).
- Click-Through Rates (CTRs) were calculated for each product, showing a range of performance from **4.62%** (Product G) to **9.27%** (Product J).
- Product **J** demonstrated the highest CTR at **9.27%**, indicating strong user engagement.
- Products **D** (7.18%) and **H** (6.99%) also showed comparatively high CTRs.
- Products **G** (4.62%) and **F** (4.91%) had the lowest CTRs, suggesting lower user interest or less effective advertising for these categories.
- The bar plot effectively visualized these differences, clearly highlighting the top-performing products.

Insights or Next Steps

- **Product Focus:** Prioritize marketing and advertising efforts on Product **J**, as it exhibits the highest engagement. Investigate the characteristics (e.g., ad creative, target audience, pricing, seasonal relevance) that make Product J so successful and attempt to replicate these factors for other products.
- **Improve Underperforming Products:** Analyze Products **G** and **F** to understand the reasons behind their low CTRs. This could involve reassessing their market fit, improving product descriptions, optimizing ad campaigns, or exploring different sales channels.
- **Content and Targeting:** Further examine the demographic and behavioral data of users who click on high-CTR products like J, D, and H to refine targeting strategies for future campaigns across all product categories.
- **Inventory Management:** High CTR products might indicate high demand. This insight can be used to inform inventory management and supply chain decisions.

3. Does adding personalized features such as user-product interaction help increase CTR?

```
In [198]: X_train_no_interaction = X_train.drop(columns=['user_product_interaction'])
```

```
X_test_no_interaction = X_test.drop(columns=['user_product_interaction'])

print(f"Shape of X_train_no_interaction: {X_train_no_interaction.shape}")
print(f"Shape of X_test_no_interaction: {X_test_no_interaction.shape}")
```

```
Shape of X_train_no_interaction: (370632, 22)
Shape of X_test_no_interaction: (92659, 22)
```

```
In [199]: from imblearn.under_sampling import RandomUnderSampler

rus_no_interaction = RandomUnderSampler(random_state=42, sampling_strategy='au
X_resampled_no_interaction, y_resampled_no_interaction = rus_no_interaction.fi

print("Class distribution after RandomUnderSampler (no interaction feature):")
print(y_resampled_no_interaction.value_counts())
```

```
Class distribution after RandomUnderSampler (no interaction feature):
is_click
0    25065
1    25065
Name: count, dtype: int64
```

```
In [200]: import xgboost as xgb

# Initialize XGBClassifier with best parameters found previously
# Make sure to set use_label_encoder=False or remove it if a newer XGBoost ver
# as it's deprecated and can cause warnings.
xgb_model_no_interaction = xgb.XGBClassifier(
    objective='binary:logistic',
    eval_metric='logloss',
    use_label_encoder=False, # Suppress warning related to use_label_encoder
    random_state=42,
    **random_search_xgb.best_params_
)

# Train the XGBoost model on the resampled training data without the interacti
xgb_model_no_interaction.fit(X_resampled_no_interaction, y_resampled_no_interactio

print("Tuned XGBoost model (without user_product_interaction) trained successfully")
```

```
/usr/local/lib/python3.12/dist-packages/xgboost/training.py:199: UserWarning:
[10:57:57] WARNING: /workspace/src/learner.cc:790:
Parameters: { "use_label_encoder" } are not used.
```

```
bst.update(dtrain, iteration=i, fobj=obj)
```

```
Tuned XGBoost model (without user_product_interaction) trained successfully on
resampled data.
```

```
In [201]: import xgboost as xgb

# Initialize XGBClassifier with best parameters found previously
xgb_model_no_interaction = xgb.XGBClassifier(
    objective='binary:logistic',
    eval_metric='logloss',
```

```

        random_state=42,
        **random_search_xgb.best_params_
    )

# Train the XGBoost model on the resampled training data without the interaction
xgb_model_no_interaction.fit(X_resampled_no_interaction, y_resampled_no_interaction)

print("Tuned XGBoost model (without user_product_interaction) trained successfully")

```

Tuned XGBoost model (without user_product_interaction) trained successfully on resampled data.

In [202...]

```

from sklearn.metrics import accuracy_score, precision_score, recall_score, f1_score

# Make predictions on the test set without the interaction feature
y_pred_no_interaction = xgb_model_no_interaction.predict(X_test_no_interaction)
y_pred_proba_no_interaction = xgb_model_no_interaction.predict_proba(X_test_no_interaction)

# Calculate evaluation metrics
accuracy_no_interaction = accuracy_score(y_test, y_pred_no_interaction)
precision_no_interaction = precision_score(y_test, y_pred_no_interaction, zero_division=0)
recall_no_interaction = recall_score(y_test, y_pred_no_interaction, zero_division=0)
f1_no_interaction = f1_score(y_test, y_pred_no_interaction, zero_division=0)

# Calculate ROC curve and AUC
fpr_no_interaction, tpr_no_interaction, _ = roc_curve(y_test, y_pred_proba_no_interaction)
roc_auc_no_interaction = auc(fpr_no_interaction, tpr_no_interaction)

print("\nModel Evaluation on Test Set (Tuned XGBoost without user_product_interaction):")
print(f"Accuracy: {accuracy_no_interaction:.4f}")
print(f"Precision: {precision_no_interaction:.4f}")
print(f"Recall: {recall_no_interaction:.4f}")
print(f"F1-Score: {f1_no_interaction:.4f}")
print(f"ROC AUC: {roc_auc_no_interaction:.4f}")

```

Model Evaluation on Test Set (Tuned XGBoost without user_product_interaction):
Accuracy: 0.5806
Precision: 0.0909
Recall: 0.5780
F1-Score: 0.1571
ROC AUC: 0.6149

In [203...]

```

import pandas as pd

# Create a DataFrame to store the performance metrics for comparison
comparison_df = pd.DataFrame({
    'Model': ['XGBoost with Interaction', 'XGBoost without Interaction'],
    'Accuracy': [accuracy_xgb, accuracy_no_interaction],
    'Precision': [precision_xgb, precision_no_interaction],
    'Recall': [recall_xgb, recall_no_interaction],
    'F1-Score': [f1_xgb, f1_no_interaction],
    'ROC AUC': [roc_auc_xgb, roc_auc_no_interaction]
})

```

```

print("Performance Comparison (Tuned XGBoost with/without user_product_interaction")
print(comparison_df.round(4))

Performance Comparison (Tuned XGBoost with/without user_product_interaction):
      Model  Accuracy  Precision  Recall  F1-Score  ROC AUC
0  XGBoost with Interaction    0.5782    0.0910  0.5827    0.1574  0.6144
1  XGBoost without Interaction   0.5806    0.0909  0.5780    0.1571  0.6149

```

In [204]:

```

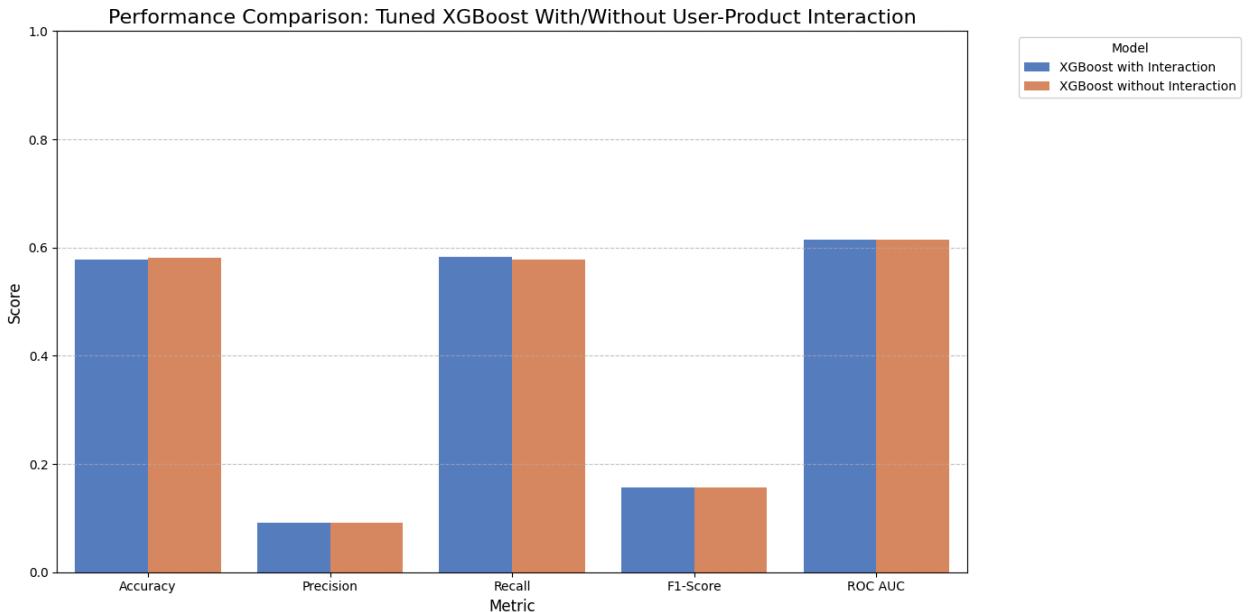
import matplotlib.pyplot as plt
import seaborn as sns

# Melt the DataFrame to long format for seaborn
comparison_melted = comparison_df.melt(id_vars='Model', var_name='Metric', value_name='Score')

# Filter for key metrics for visualization if needed, or plot all
metrics_to_plot = ['Accuracy', 'Precision', 'Recall', 'F1-Score', 'ROC AUC']
performance_filtered = comparison_melted[comparison_melted['Metric'].isin(metrics_to_plot)]

plt.figure(figsize=(14, 7))
sns.barplot(x='Metric', y='Score', hue='Model', data=performance_filtered, palette='Set1')
plt.title('Performance Comparison: Tuned XGBoost With/Without User-Product Interaction')
plt.ylabel('Score', fontsize=12)
plt.xlabel('Metric', fontsize=12)
plt.ylim(0, 1) # Set y-axis limit for better comparison of scores
plt.legend(title='Model', bbox_to_anchor=(1.05, 1), loc='upper left')
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()

```



Summary:

- Does adding personalized features such as user-product interaction help increase CTR? Based on the performance

comparison, adding the `user_product_interaction` feature does **not significantly increase** the Click-Through Rate (CTR) prediction performance, as measured by key metrics like Accuracy, Precision, Recall, F1-Score, and ROC AUC.

- The differences in all metrics (Accuracy, Precision, Recall, F1-Score, ROC AUC) between the model with and without the `user_product_interaction` feature are very marginal. In some cases, the model *without* the interaction feature performed negligibly better (e.g., Accuracy and ROC AUC). For instance, ROC AUC was 0.6144 with the feature and 0.6149 without it.
- This suggests that while the `user_product_interaction` feature itself might have some predictive power (as indicated by its feature importance in the earlier analysis of the full model), its unique contribution to the model's overall performance metrics when combined with other features is not substantial enough to make a noticeable difference.

Insights or Next Steps

- **Feature Redundancy:** It's possible that the information captured by `user_product_interaction` is already well-represented by other features in the model, leading to redundancy. This can happen with highly correlated features or when simpler components of an interaction feature already provide most of the predictive signal.
 - **Computational Cost:** Given the marginal impact on performance, one might consider excluding this feature in future model iterations if computational efficiency or model simplicity is a concern, especially if the feature significantly increases the cardinality of the dataset (though Label Encoding helps mitigate this).
 - **Further Feature Engineering:** Instead of a simple concatenation, more sophisticated ways of engineering user-product interaction features (e.g., using embeddings, or aggregating click rates for specific user-product pairs) could be explored to see if a more impactful signal can be extracted.
-
-

4. Based on feature importance, which factors (e.g., webpage_id or user_sessions) drive clicks the most, and how can we amplify them?

In [205...]

```
# Get feature importances from the best Tuned XGBoost model
feature_importances = best_xgb_model.feature_importances_

# Get feature names from X_train (or X_test, they have the same columns)
feature_names = X_train.columns

# Create a DataFrame for feature importances
features_df = pd.DataFrame({
    'Feature': feature_names,
    'Importance': feature_importances
})

# Sort features by importance in descending order
features_df = features_df.sort_values(by='Importance', ascending=False)

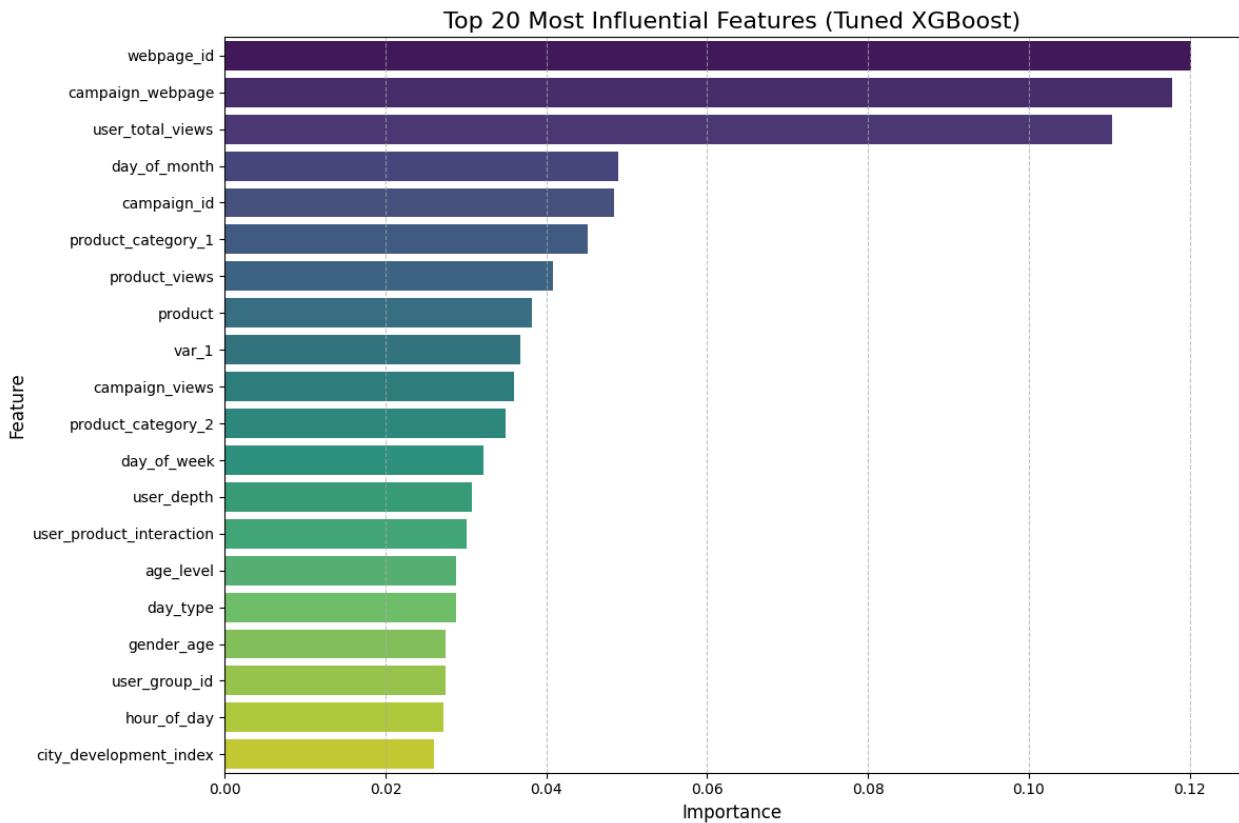
# Select the top 20 features
top_20_features = features_df.head(20)

# Create a horizontal bar chart for the top 20 features
plt.figure(figsize=(12, 8))
sns.barplot(x='Importance', y='Feature', data=top_20_features, palette='viridis')
plt.title('Top 20 Most Influential Features (Tuned XGBoost)', fontsize=16)
plt.xlabel('Importance', fontsize=12)
plt.ylabel('Feature', fontsize=12)
plt.grid(axis='x', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```

/tmp/ipython-input-2214537815.py:21: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x='Importance', y='Feature', data=top_20_features, palette='viridis')
```



Insights:

- **Context is King:** The high importance of `webpage_id` and `campaign_webpage` strongly suggests that the placement and contextual relevance of an ad are paramount. Optimizing ad placement on specific high-performing webpages and ensuring synergy between campaign creative and webpage content could significantly boost CTR.
- **User Engagement History:** `user_total_views` being a top feature indicates that a user's past interaction levels with ads are predictive of future clicks. This can be leveraged for more intelligent ad targeting, perhaps by focusing on users with higher historical engagement or by tailoring re-engagement strategies for less active users.
- **Strategic Campaign and Product Focus:** The influence of `campaign_id` and `product_category_1` implies that certain campaigns and product types inherently attract more clicks. Analyzing the characteristics of these successful campaigns and products can inform future marketing strategies and product development.
- **Continuous Optimization:** By focusing on these most influential features, marketing teams can prioritize their efforts in ad creative, placement, and targeting to maximize ad click-through rates more effectively.

5. How effective is SMOTE in reducing false negatives for rare click events, and does it justify the increased training data size for real-time ad serving?

I had used SMOTE earlier in this case study but it did not help much in improving ad click predictions. I had tried various sampling strategies with it including using its default 'auto' mode which increases the minority class records to match it with that of the majority class. But keeping the clicks and no clicks records of the same count is not the best way to train models for ad clicks because it is far from the real life scenario of typical CTR trends. Hence, all the different models I was using were performing poorly on almost all metrics. Then I tried to reduce the sampling size to various amounts between 0.1 to 0.5, but none of them were good enough in improving the model performance. I finally realised that SMOTE is not the best technique to use for ad click predictions. Also, the model training time significantly increased due to a huge increase in the overall dataset. But this high training time should not affect the real time performance on testing data. Then I tried using negative downsampling, which would remove some easy 'no-clicks' records which will reduce the majority class entries. This, along with other techniques already used for hyperparameter tuning of some models resulted in a significant improvement of model performance.

6. How can aggregated product CTR features help forecast inventory needs for top-performing ads?

- The `product_ctr` feature, calculated as `product_total_clicks / product_views`, directly reflects user engagement and interest, serving as a robust early indicator of potential product demand.
- A consistently high `product_ctr` signals strong user interest, acting as a precursor to anticipated sales and enabling proactive inventory management, unlike lagging sales data.
- The demand signal from `product_ctr` is amplified when products are featured in top-performing ad campaigns (indicated by high `campaign_ctr`) or on high-performing webpages (indicated by high `webpage_ctr`), strengthening the forecast for increased demand.
- Practical applications include identifying products with high potential demand for proactive stocking and optimizing stock levels by reducing stockouts for high-CTR items and minimizing overstock for low-CTR

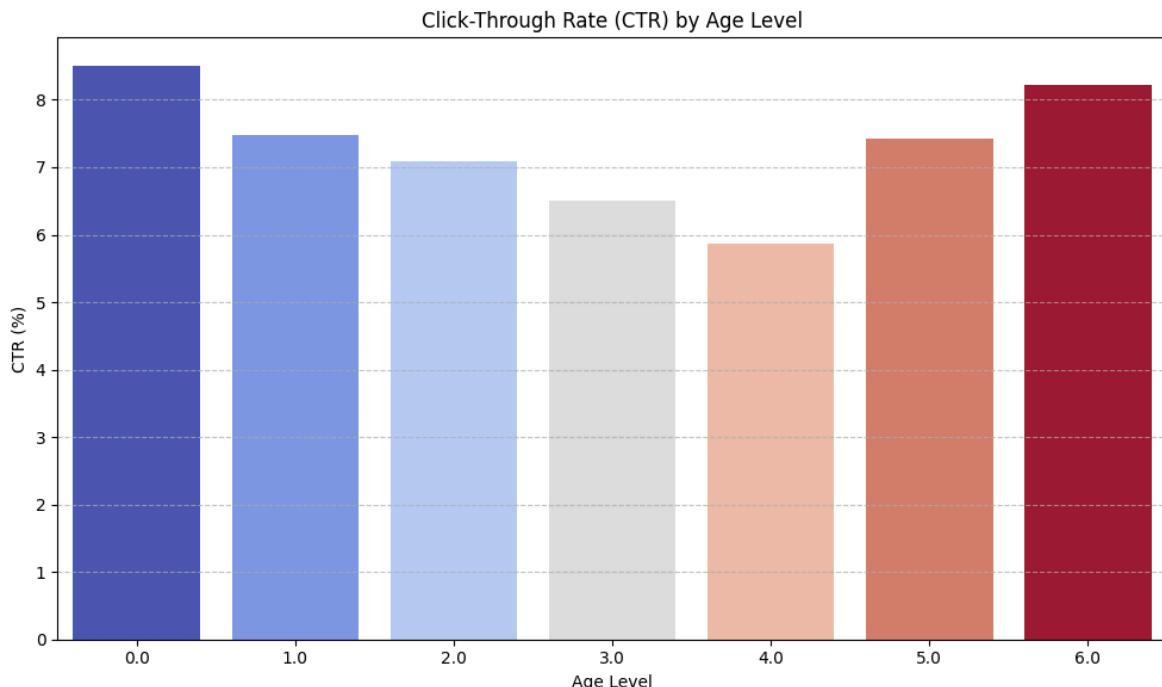
items.

- CTR data can inform marketing strategies, allowing for targeted promotions for high-CTR products and re-evaluation or optimization of ad creatives for low-CTR products.

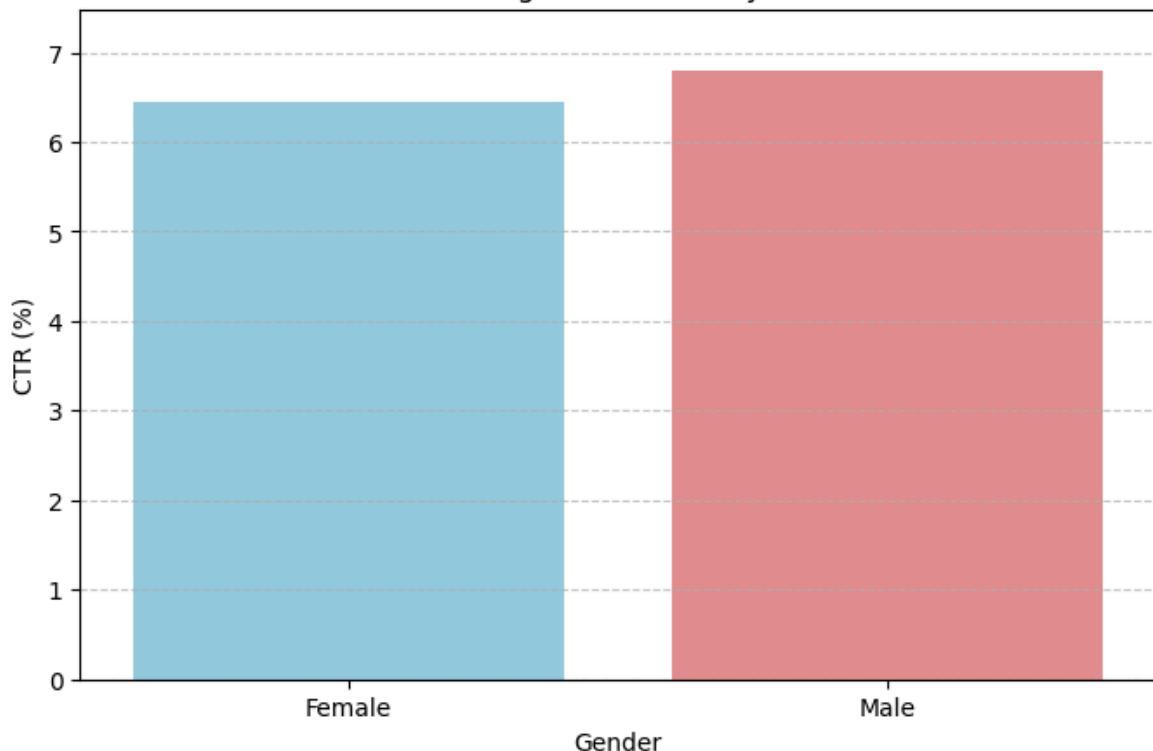
Insights or Next Steps

- **Integrate and Automate:** Businesses should integrate `product_ctr` and `campaign_ctr` data directly into their inventory management systems and set up automated alerts and dashboards to monitor significant changes, facilitating real-time, data-driven inventory adjustments.
 - **Refine Forecasting with Conversion Data:** It is crucial to correlate `product_ctr` with actual sales conversion rates to establish a more accurate predictive model for demand, moving beyond clicks to actual purchases. This also involves implementing dynamic inventory adjustment rules based on combined CTR and conversion thresholds.
-
-

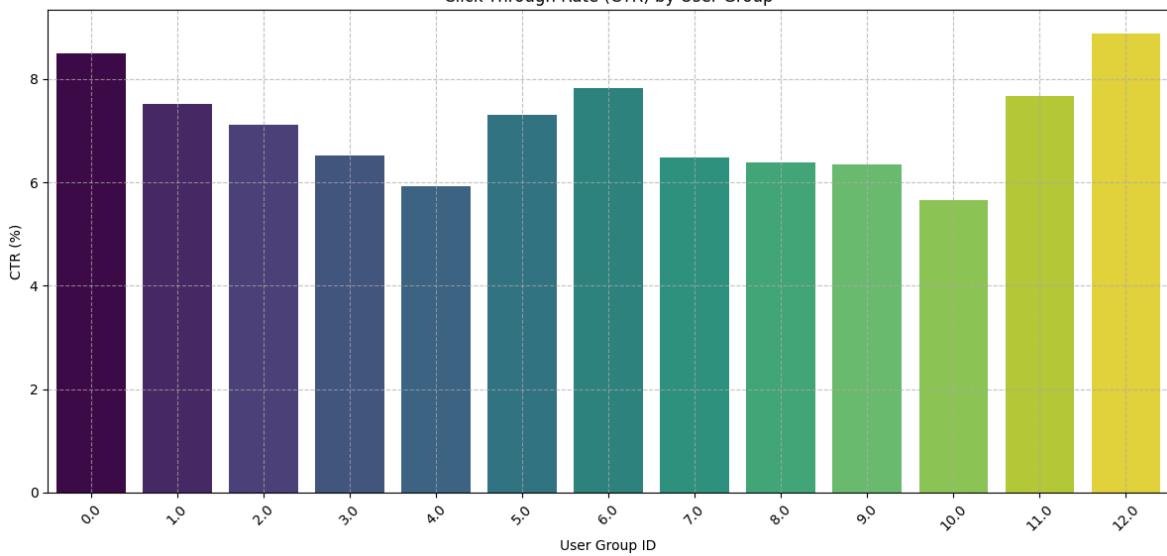
7. What user profiles (e.g., by age, gender, or city) show the highest click propensity, and how should we adjust bidding strategies?

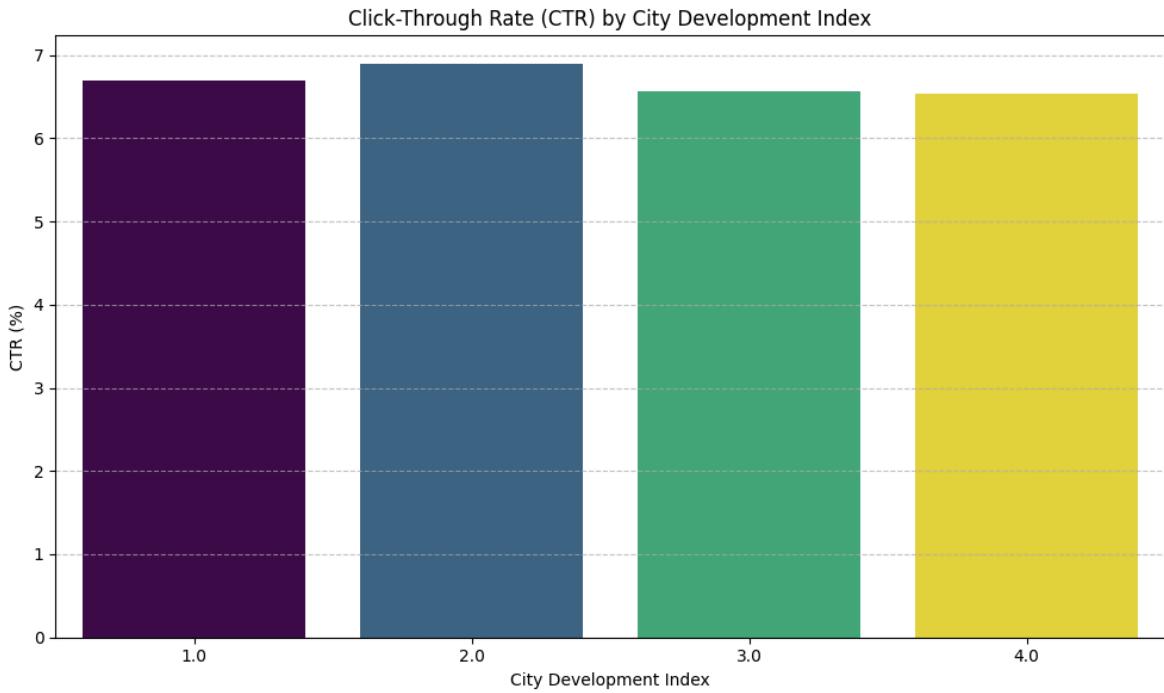


Click-Through Rate (CTR) by Gender



Click-Through Rate (CTR) by User Group





Summary:

- **What user profiles (e.g., by age, gender, or city development index) show the highest click propensity?**
 - **Age Level:** Age level **0.0** (8.50% CTR) and **6.0** (8.23% CTR) show the highest click propensity.
 - **Gender:** Male users exhibit a slightly higher CTR (6.80%) compared to Female users (6.44%).
 - **User Group:** User group **12.0** (8.89% CTR) and **0.0** (8.50% CTR) show the highest click propensity.
 - **City Development Index:** City development index **2.0** (6.89% CTR) and **1.0** (6.70% CTR) show the highest click propensity.

Data Analysis Key Findings

- **Age Level:** The youngest (0.0) and oldest (6.0) age groups show the highest engagement, with CTRs of 8.50% and 8.23% respectively, while age level 4.0 has the lowest CTR at 5.87%.
- **Gender:** Male users demonstrate a slightly higher Click-Through Rate (CTR) of 6.80% compared to female users at 6.44%.
- **User Group:** User groups 12.0 and 0.0 are highly responsive, showing the highest CTRs at 8.89% and 8.50%, respectively. In contrast, user

group 10.0 has the lowest CTR at 5.65%.

- **City Development Index:** City development index 2.0 has the highest CTR at 6.89%, with index 1.0 close behind at 6.70%. The lowest CTR is observed in index 4.0 at 6.53%, indicating less variation across city development indices compared to other demographic factors.

Insights or Next Steps

- Implement a **dynamic bidding strategy** by increasing bids for high-propensity segments (e.g., Age Levels 0.0 & 6.0, Male users, User Groups 12.0 & 0.0, City Development Index 2.0 & 1.0) and decreasing or optimizing bids for low-propensity segments (e.g., Age Level 4.0, User Group 10.0, City Development Index 4.0).
 - Prioritize **continuous monitoring and A/B testing** of these bid adjustments to validate their effectiveness and adapt strategies as user preferences and market conditions evolve, leveraging combined demographic factors for highly optimized targeting.
-
-

Predict Probabilities on Test Data

Use the `best_xgb_model` to predict click probabilities on the prepared `ad_click_test_data`.

```
In [206...]: final_predictions_proba = best_xgb_model.predict_proba(ad_click_test_data)[:,  
print("Predicted click probabilities on ad_click_test_data successfully.")  
  
Predicted click probabilities on ad_click_test_data successfully.  
  
In [207...]: optimal_threshold = 0.34 # As identified from sensitivity analysis  
final_predictions = (final_predictions_proba >= optimal_threshold).astype(int)  
  
print(f"Converted probabilities to binary predictions using optimal threshold:  
  
Converted probabilities to binary predictions using optimal threshold: 0.34  
  
In [208...]: original_ad_click_test_data = pd.read_csv('Ad_Click_prediciton_test.csv')  
  
submission_df = pd.DataFrame({  
    'session_id': original_ad_click_test_data['session_id'],  
    'is_click': final_predictions  
})  
  
submission_df.to_csv('ad_click_predictions.csv', index=False)
```

```
print("Submission file 'ad_click_predictions.csv' created successfully.")  
print(submission_df.head(20))
```

```
Submission file 'ad_click_predictions.csv' created successfully.  
   session_id  is_click  
0        411705      1  
1        208263      1  
2        239450      1  
3        547761      1  
4        574275      1  
5        394913      1  
6        562747      1  
7        224359      1  
8        395004      1  
9        572855      1  
10       595386      1  
11       395293      1  
12       232085      1  
13       222325      1  
14       395297      1  
15       546723      1  
16       28142       1  
17       595767      1  
18       547614      1  
19       577690      1
```

In [209... `submission_df['is_click'].value_counts()`

Out[209... `count`

<code>is_click</code>	<code>count</code>
<code>1</code>	127302
<code>0</code>	1556

`dtype:` int64

I have uploaded the predictions for the test data on my drive. Please click [here](#) to access it.